

CSED311 Computer Architecture – Lecture 5


Single-Cycle CPU

Eunhyeok Park

Department of Computer Science and Engineering
POSTECH

Disclaimer: Slides developed in part by Profs. Austin, Brehob, Falsafi, Hill, Hoe, Lipasti, Martin, Roth, Shen, Smith, Sohi, Tyson, Vijaykumar, and Wenisch @ Carnegie Mellon University, University of Michigan, Purdue University, University of Pennsylvania, University of Wisconsin and POSTECH.

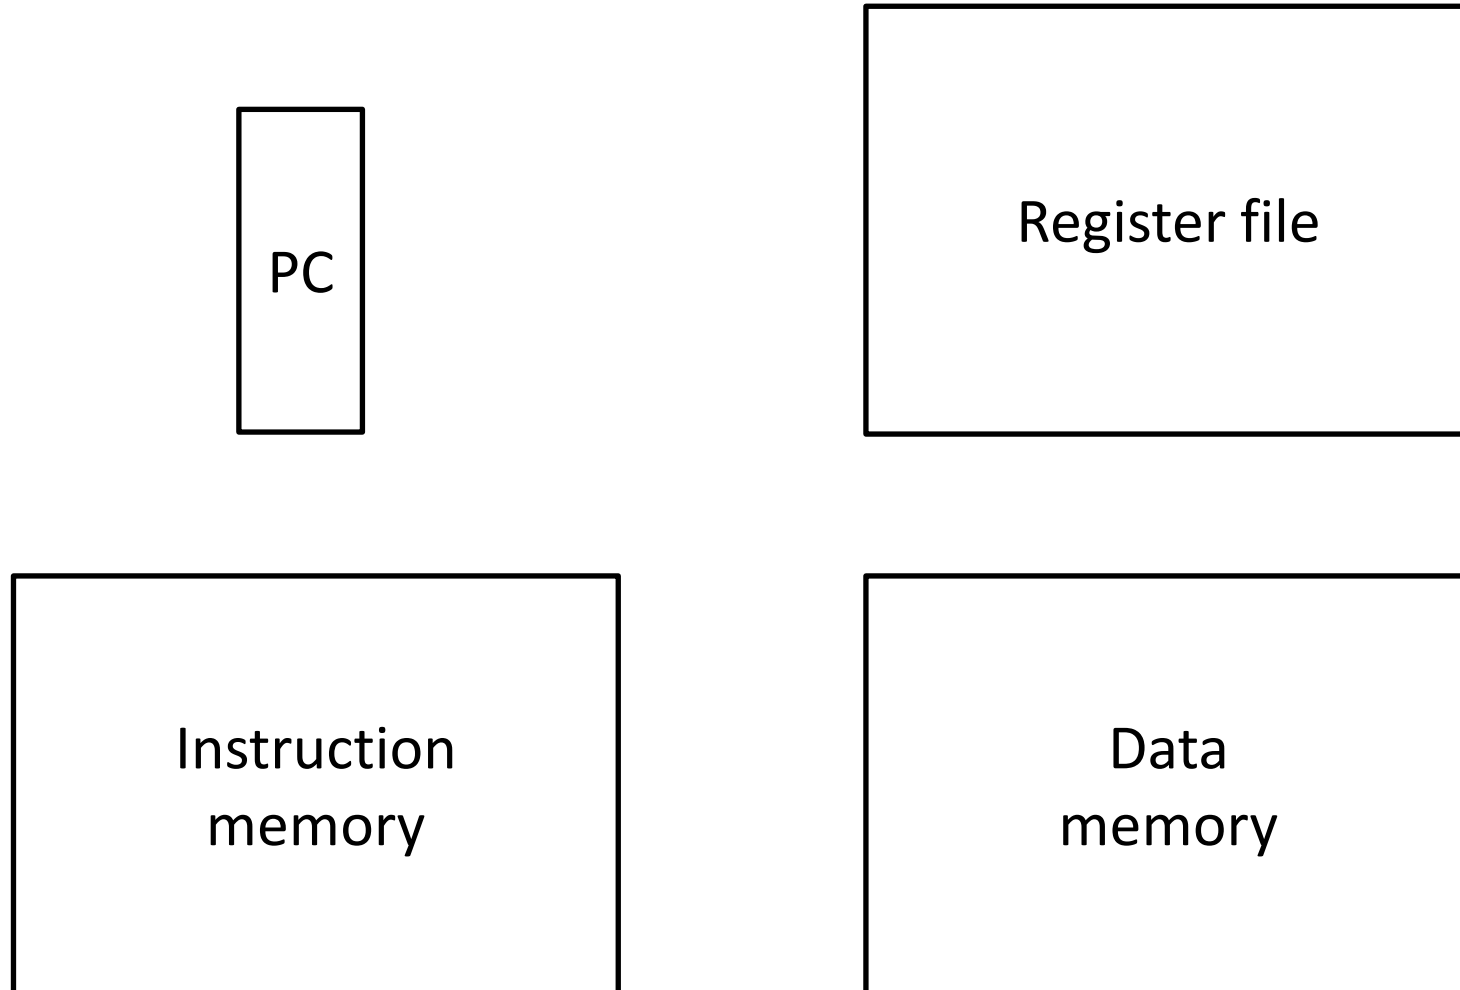
CPU: Instruction Processing FSM



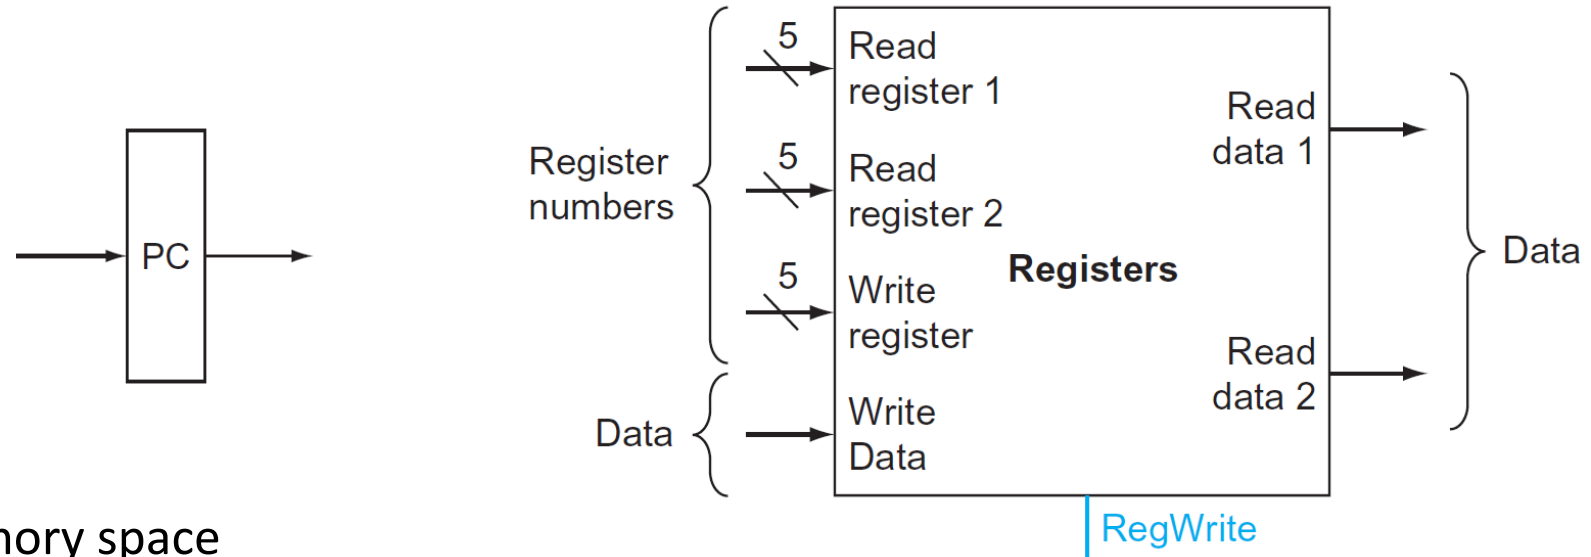
- An ISA describes an abstract finite-state machine (FSM)
 - State = program visible state
 - Next-state logic = instruction execution
- Nice ISAs have atomic instruction semantics
 - One state transition per instruction in abstract FSM
- Implementation of FSMs can vary

Programmer-Visible (Architectural) State


- What are the inputs and outputs for each component?



Programmer-Visible (Architectural) State



Single memory space



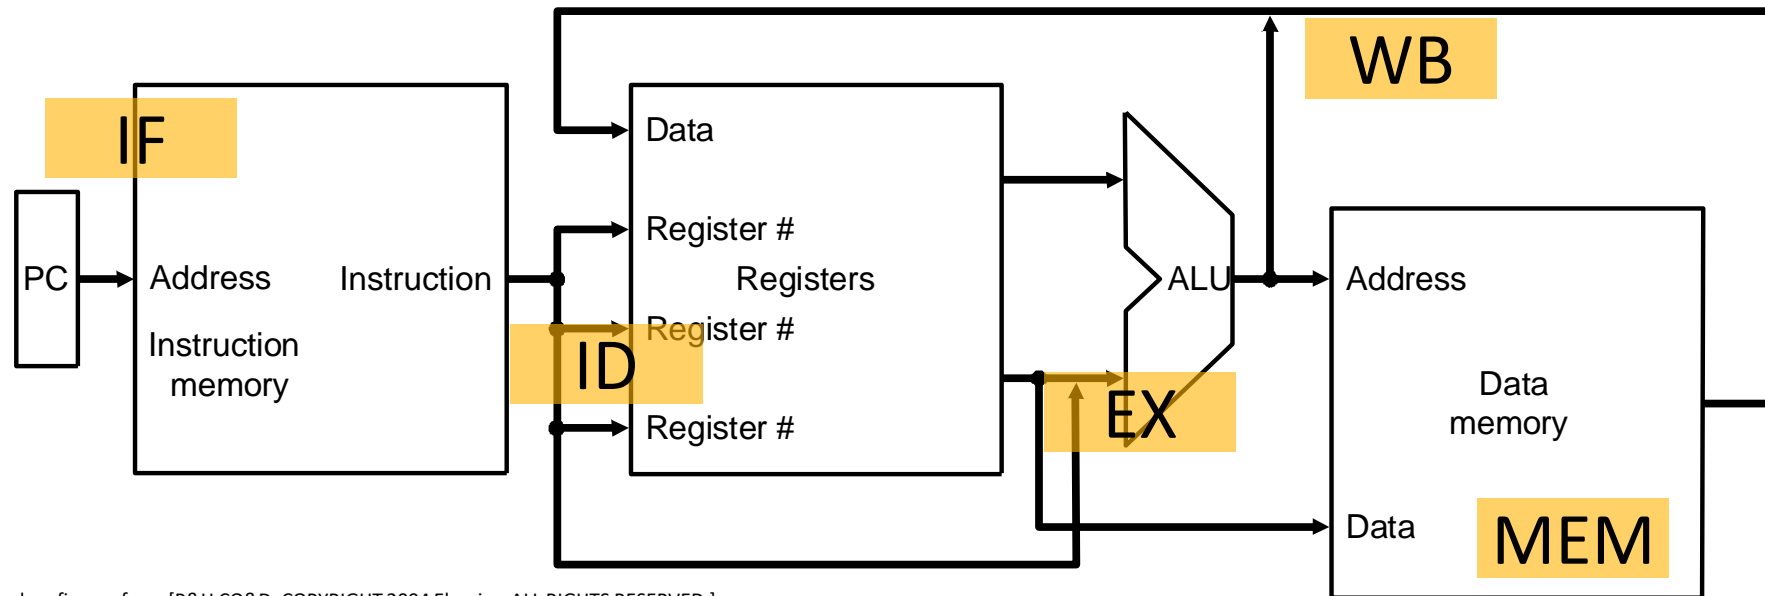
“Magic” Register File and Memory

- Combinational read
 - The output of the read data port is a combinational function of the register file contents and the corresponding read select port
 - Similarly, instruction/data memory work like combinational logic
- Synchronous write
 - The selected register (or memory location) is updated **only** on the positive-edge clock transition when write enable is asserted

Cannot affect read output in between clock edges
- The instruction memory also needs to be written to when the program is loaded, but let's consider it read-only for simplicity for now
- We will learn about more realistic design of memory hierarchy later

Instruction Processing


- 5 generic steps
 - **IF**: Instruction fetch
 - **ID**: Instruction decode and operand fetch
 - **EX**: ALU/execute
 - **MEM**: (Data) Memory access (only for load & store instructions)
 - **WB**: Write-back



**Based on figures from [P&H CO&D, COPYRIGHT 2004 Elsevier. ALL RIGHTS RESERVED.]

Contents

- Datapath for ALU instructions
- Datapath for LD/ST instructions
- Datapath for control flow instructions
- Control



Today's lecture is about *microarchitecture*!

The microarchitecture *implements* the RISC-V *architecture*

Single-Cycle Datapath for Arithmetic and Logical Instructions

Reminder: R-Type Instruction

- Assembly example: `ADD rd, rs1, rs2`

■ Encoding:

31	25 24	20 19	15 14	12 11	7 6	0
funct7	rs2	rs1	funct3	rd	opcode	
7	5	5	3	5	7	
0000000	src2	src1	ADD/SLT/SLTU	dest	OP	} = 0110011
0000000	src2	src1	AND/OR/XOR	dest	OP	
0000000	src2	src1	SLL/SRL	dest	OP	
0100000	src2	src1	SUB/SRA	dest	OP	

- Semantics

- $\text{GPR}[\text{rd}] \leftarrow \text{GPR}[\text{rs1}] \text{ (OP) } \text{GPR}[\text{rs2}]$ $[(\text{OP}): +, -, |, \&, \ll, \gg, \dots]$
- $\text{PC} \leftarrow \text{PC} + 4$


- Variations

- Arithmetic: ADD, SUB
- Compare (Set if Less Than): SLT (signed), SLTU (unsigned)
- Logical: AND, OR, XOR
- Shift: SLL (left), SRL (right-logical), SRA (right-arithmetic)

$\text{GPR}[\text{rd}] \leftarrow 1 \text{ if } \text{GPR}[\text{rs1}] < \text{GPR}[\text{rs2}]$
 $\text{GPR}[\text{rd}] \leftarrow 0 \text{ otherwise}$

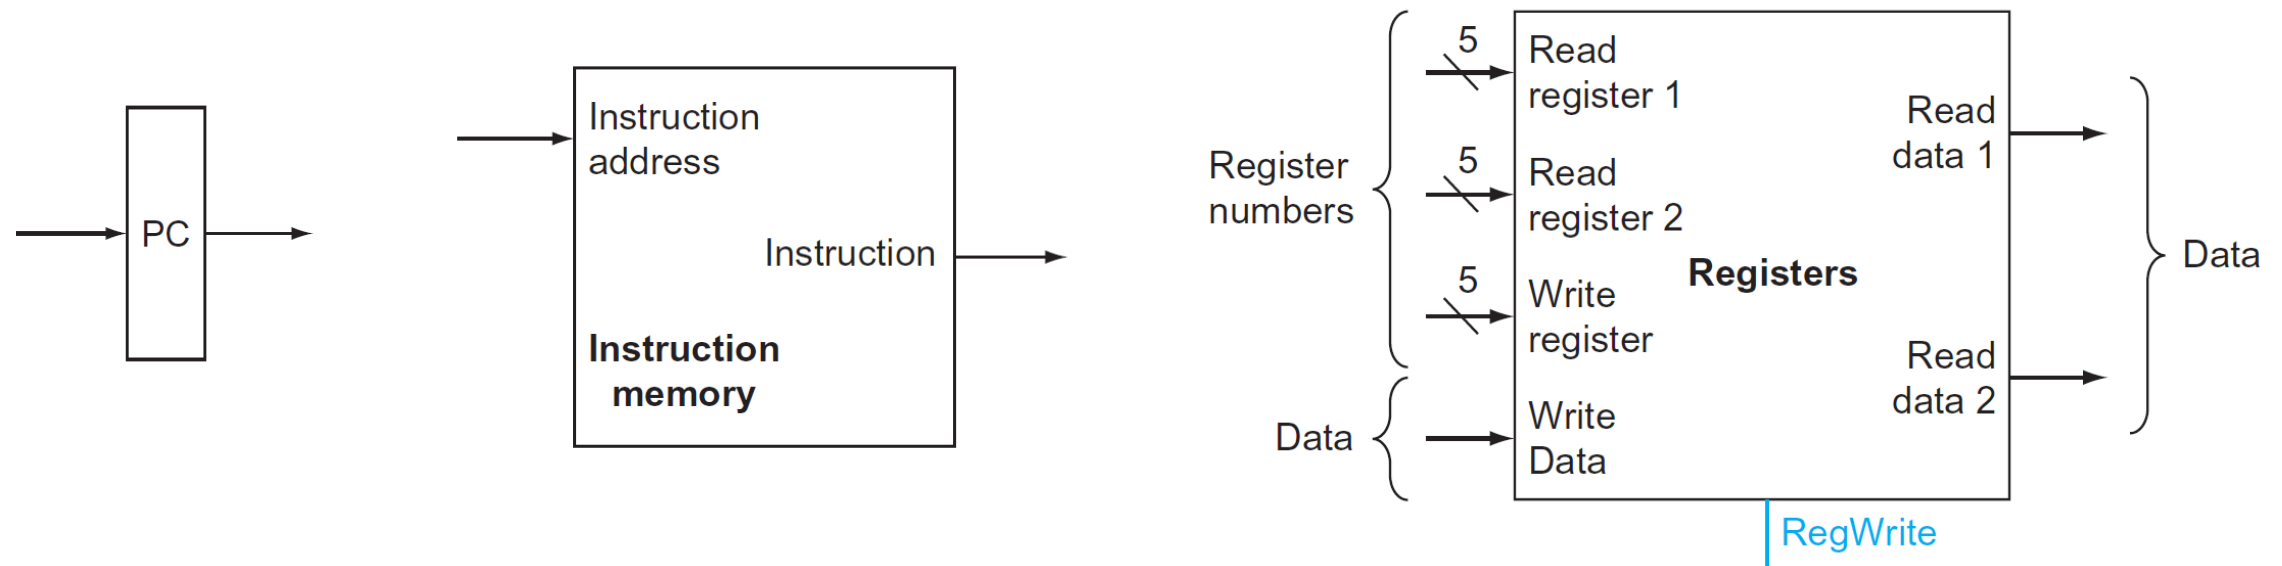
- Exception: none (ignore carry and overflow) – see page 13 of the ISA manual₉

R-Type Instruction Datapath




◆ Semantics (ADD)

- $GPR[rd] \leftarrow GPR[rs1] + GPR[rs2]$
- $PC \leftarrow PC + 4$



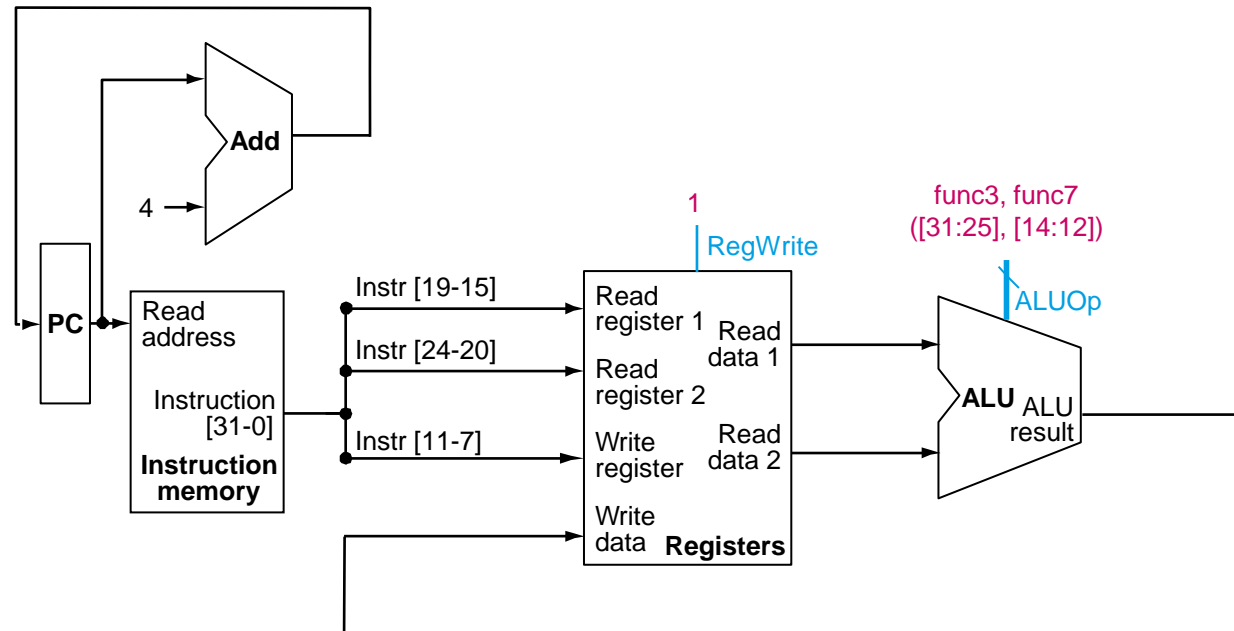
First, we are going to build the datapath and identify the control signals that need to be generated. Then, we can work on the control signals at the end.

R-Type Instruction Datapath



◆ Semantics (ADD)

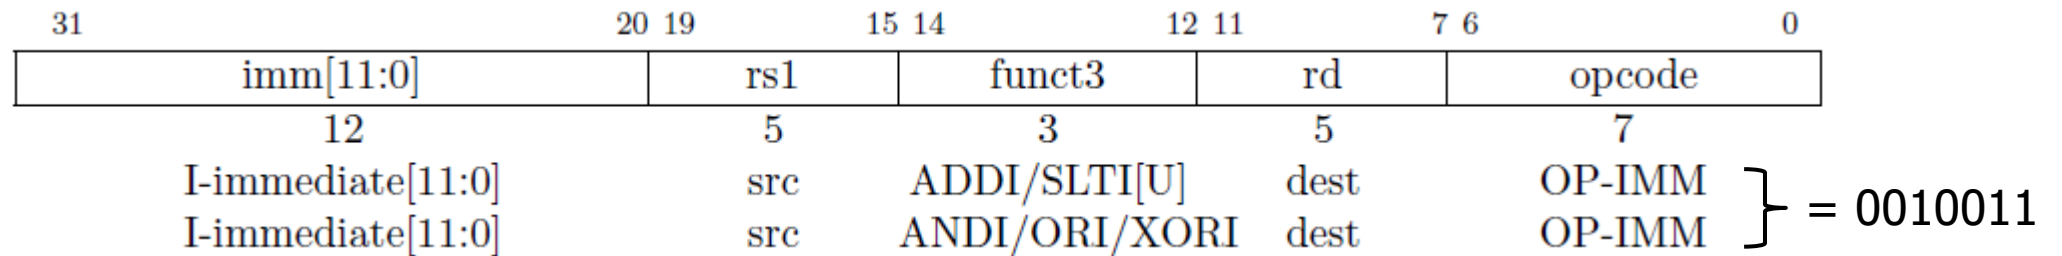
- $GPR[rd] \leftarrow GPR[rs1] + GPR[rs2]$
- $PC \leftarrow PC + 4$



Reminder: I-Type ALU Instruction

■ Assembly example: `ADDI rd, rs1, imm12`

■ Encoding:



■ Semantics

- $GPR[rd] \leftarrow GPR[rs1] \text{ (OP) sign-extend (immediate)}$ $[(OP): +, |, \&, \dots]$
- $PC \leftarrow PC + 4$

■ Variations:

- Arithmetic: ADDI, ~~SUBI~~ – why no SUBI?
- Logical: ANDI, ORI, XORI
- Compare (Set if Less Than Immediate): SLTI (signed), SLTIU (unsigned)

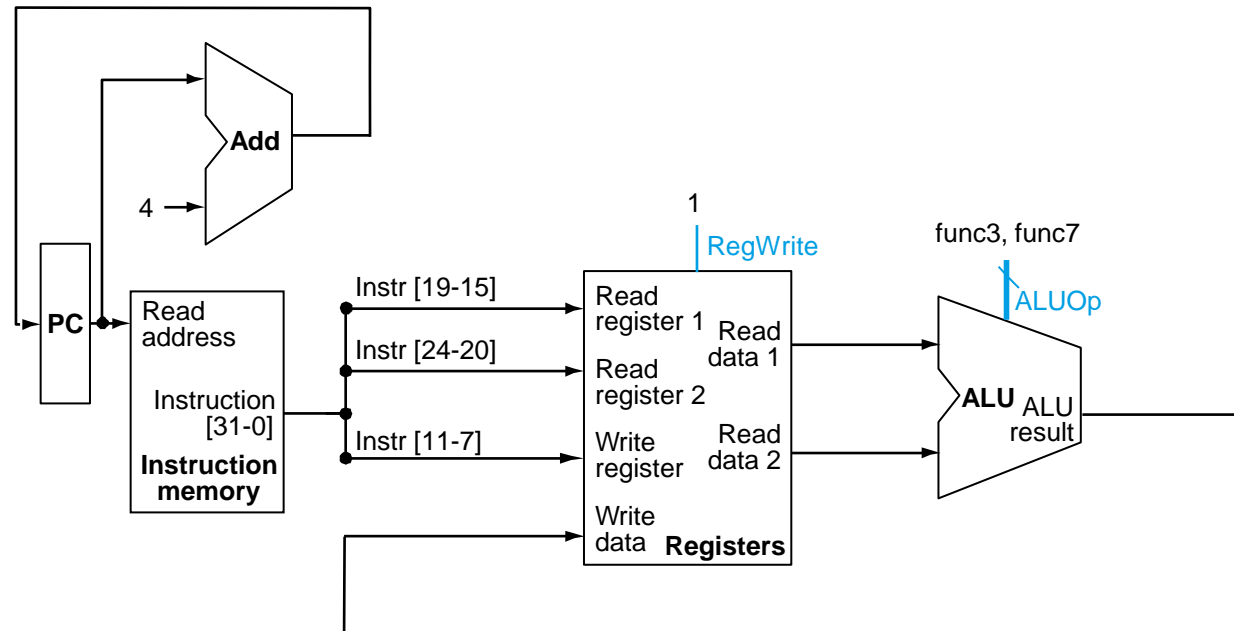

■ Exception: none (ignore carry and overflow)

I-Type ALU Instruction Datapath


■ How should this change?

◆ Semantics

- $\text{GPR}[\text{rd}] \leftarrow \text{GPR}[\text{rs1}] + \text{sign-extend}(\text{imm})$
- $\text{PC} \leftarrow \text{PC} + 4$

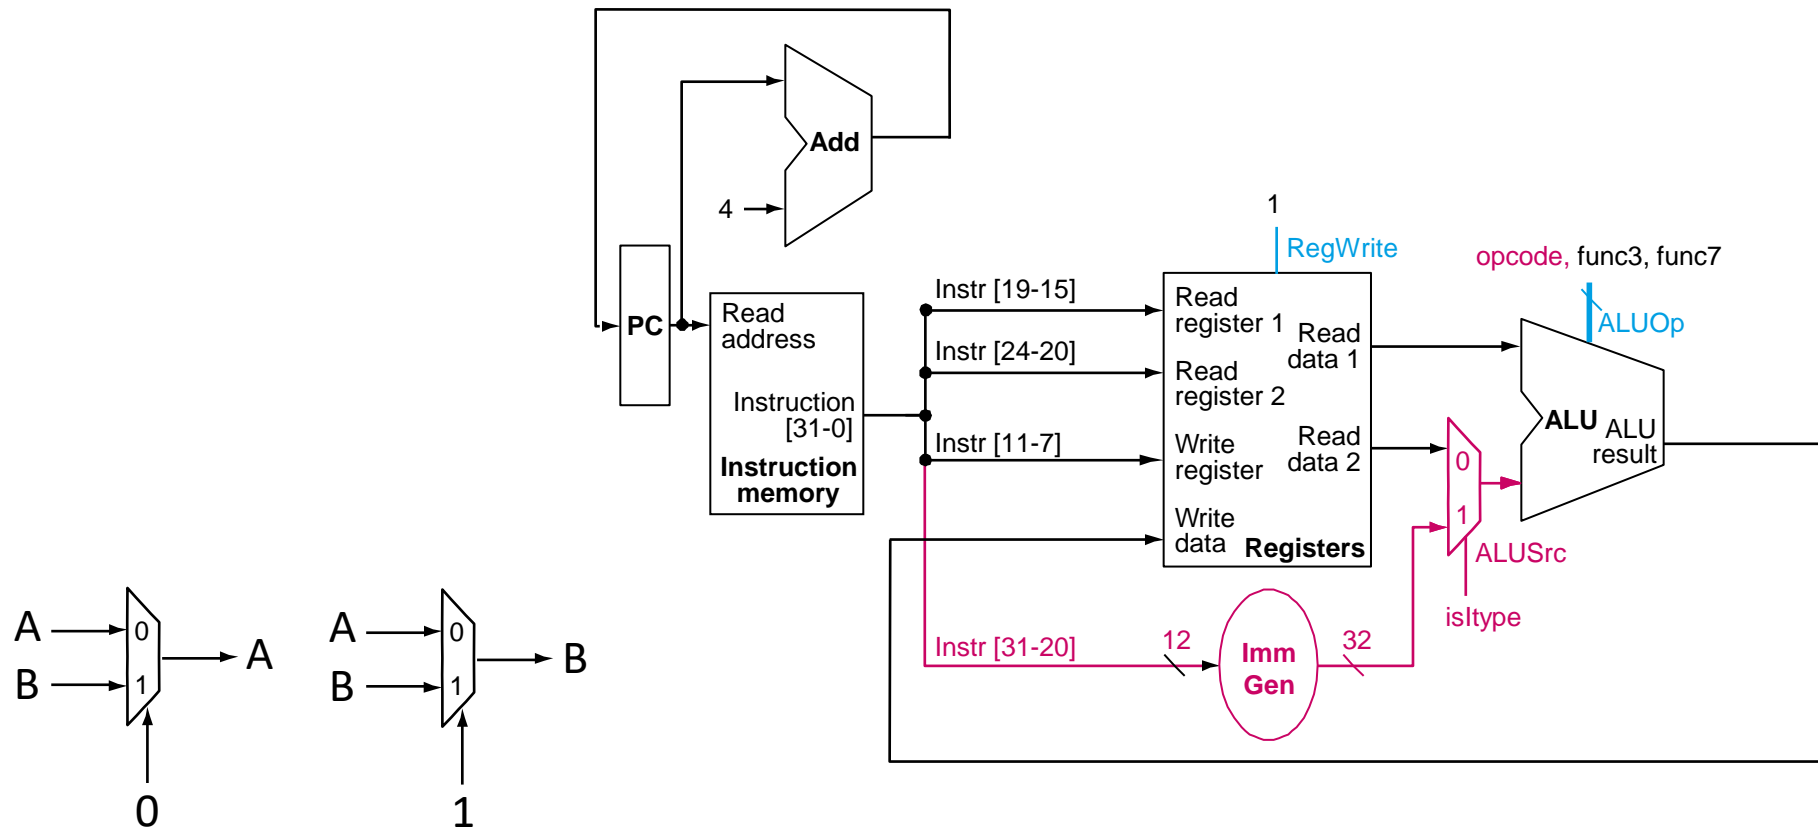


I-Type ALU Instruction Datapath



◆ Semantics

- $\text{GPR}[\text{rd}] \leftarrow \text{GPR}[\text{rs1}] + \text{sign-extend}(\text{imm})$
- $\text{PC} \leftarrow \text{PC} + 4$




Single-Cycle Datapath for Data Movement Instructions

Reminder: Load Instructions (I-type)

- Assembly example: `LW rd, offset12(base)`


- Encoding:



- Semantics:
 - $\text{byte_address}_{32} = \text{sign-extend}(\text{offset}_{12}) + \text{GPR}[\text{base}]$
 - $\text{GPR}[\text{rd}] \leftarrow \text{MEM}_{32}[\text{byte_address}_{32}]$
 - $\text{PC} \leftarrow \text{PC} + 4$
- Exceptions: will be discussed later
- Variations:
 - LW (word), LH (halfword), LB (byte): sign-extend
 - LHU, LBU: zero-extend


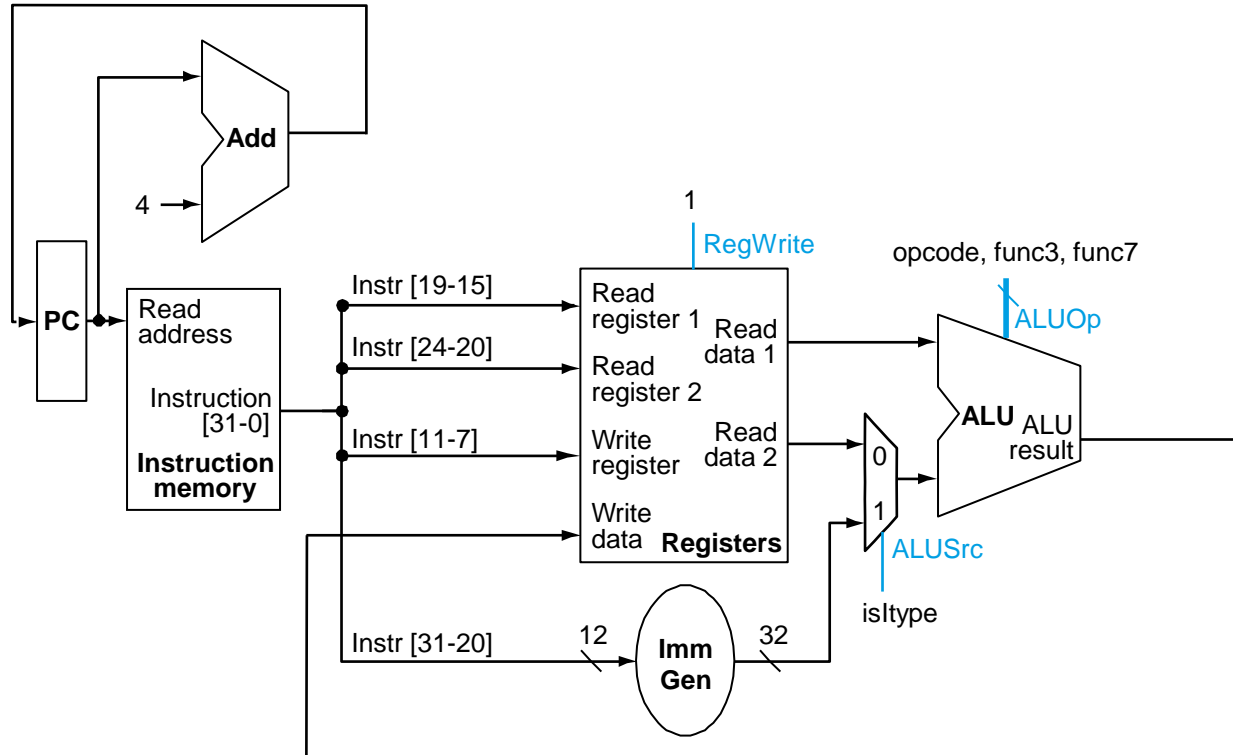
Load Datapath

■ How should this change for load?




◆ Semantics:

- $\text{byte_addr}_{32} = \text{sign-extend}(\text{offset}_{12}) + \text{GPR}[\text{base}]$
- $\text{GPR}[\text{rd}] \leftarrow \text{MEM}_{32}[\text{byte_addr}_{32}]$
- $\text{PC} \leftarrow \text{PC} + 4$

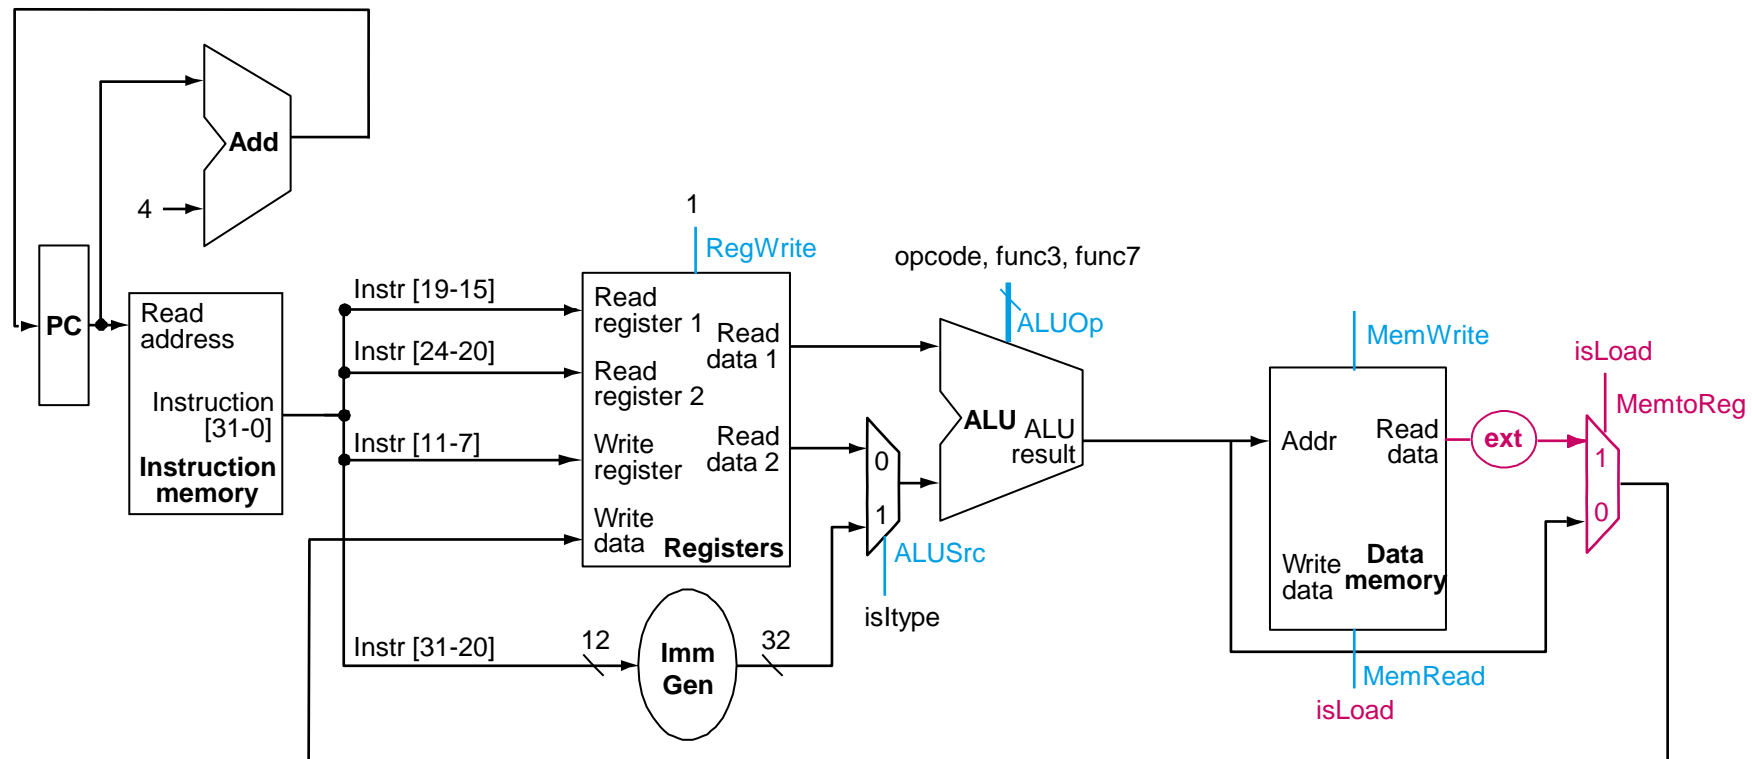


Load Datapath



◆ Semantics:

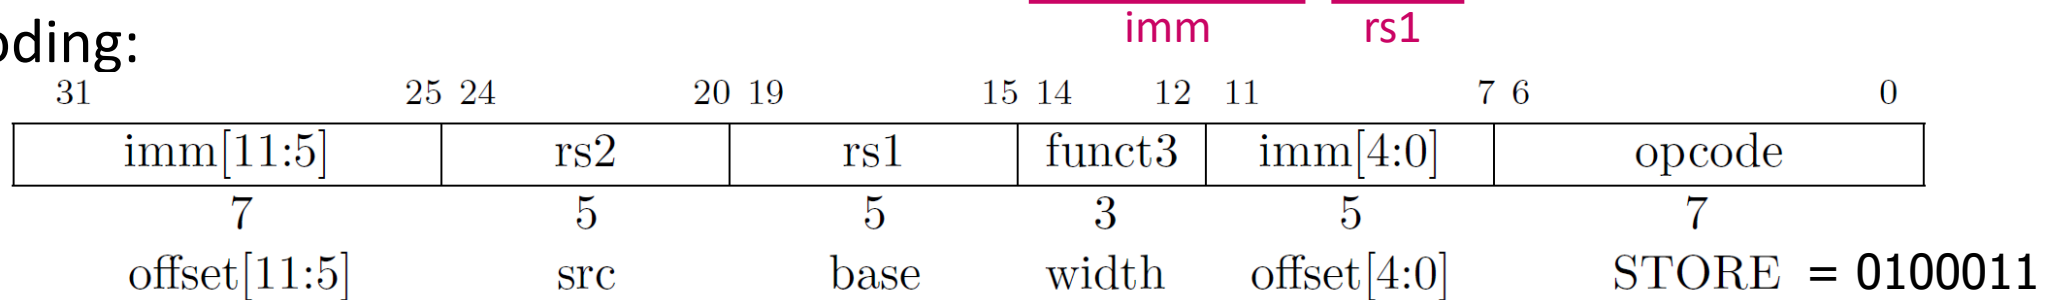
- $\text{byte_addr}_{32} = \text{sign-extend}(\text{offset}_{12}) + \text{GPR}[\text{base}]$
- $\text{GPR}[\text{rd}] \leftarrow \text{MEM}_{32}[\text{byte_addr}_{32}]$
- $\text{PC} \leftarrow \text{PC} + 4$



Reminder: Store Instructions (S-Type)

- Assembly example (Store Word): `SW rs2, offset12(base)`

- Encoding:



- Semantics:

- $\text{byte_address}_{32} = \text{sign-extend}(\text{offset}_{12}) + \text{GPR}[\text{base}]$
- $\text{MEM}_{32}[\text{byte_address}_{32}] \leftarrow \text{GPR}[\text{rs2}]$
- $\text{PC} \leftarrow \text{PC} + 4$

- Exceptions: will be discussed later


- Variations:

- SW (word), SH (halfword), SB (byte)

- Why split immediate?

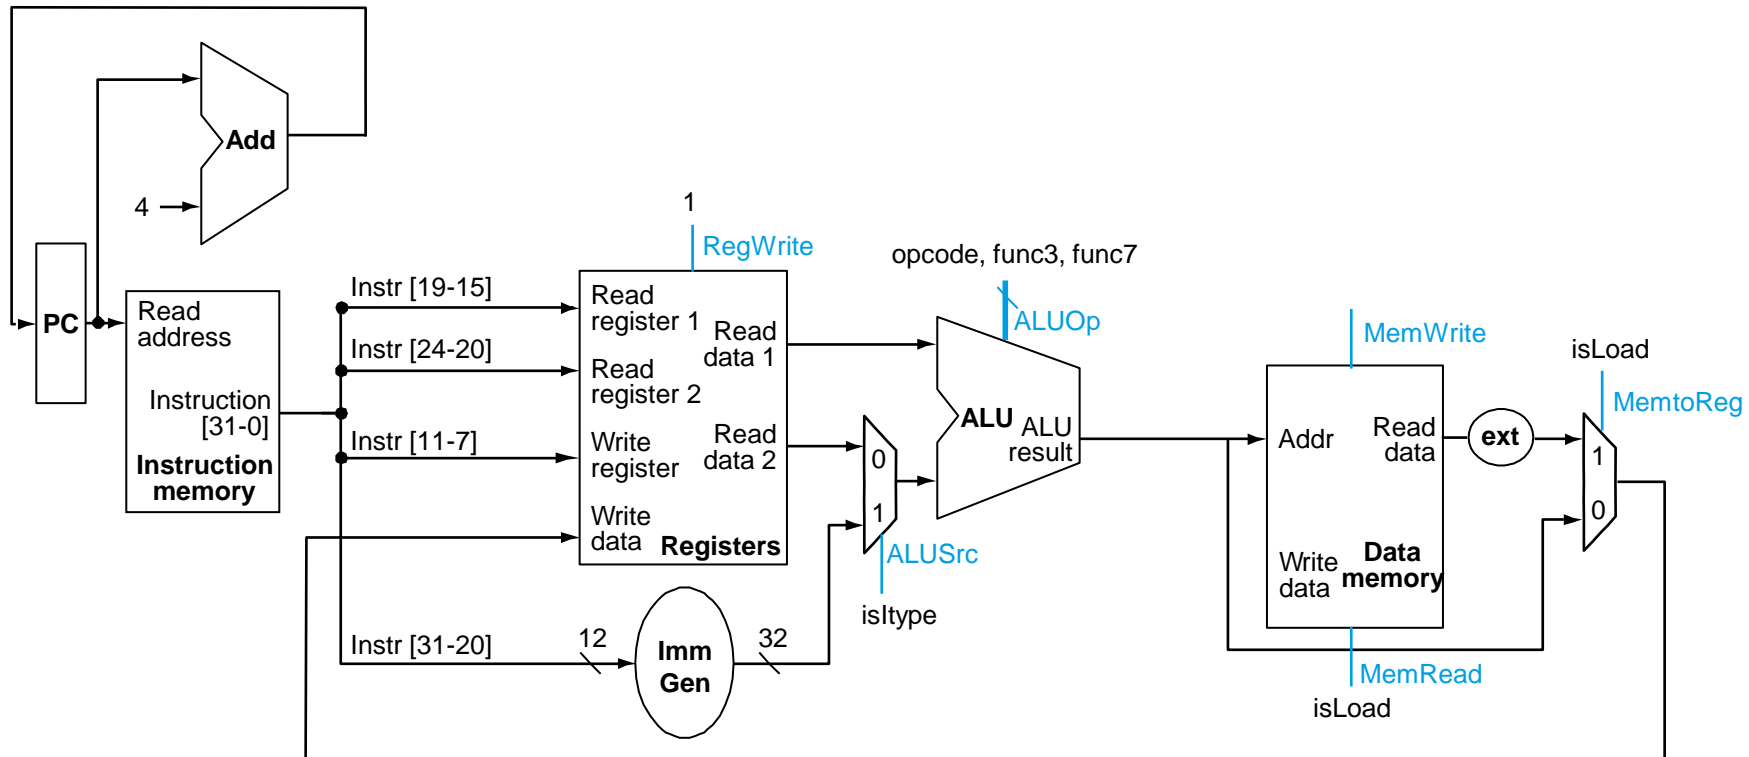
Store Datapath

■ How should this change for store?




- ◆ Semantics:

- $\text{byte_address}_{32} = \text{sign-extend}(\text{offset}_{12}) + \text{GPR}[\text{base}]$
- $\text{MEM}_{32}[\text{byte_address}_{32}] \leftarrow \text{GPR}[\text{rs2}]$
- $\text{PC} \leftarrow \text{PC} + 4$

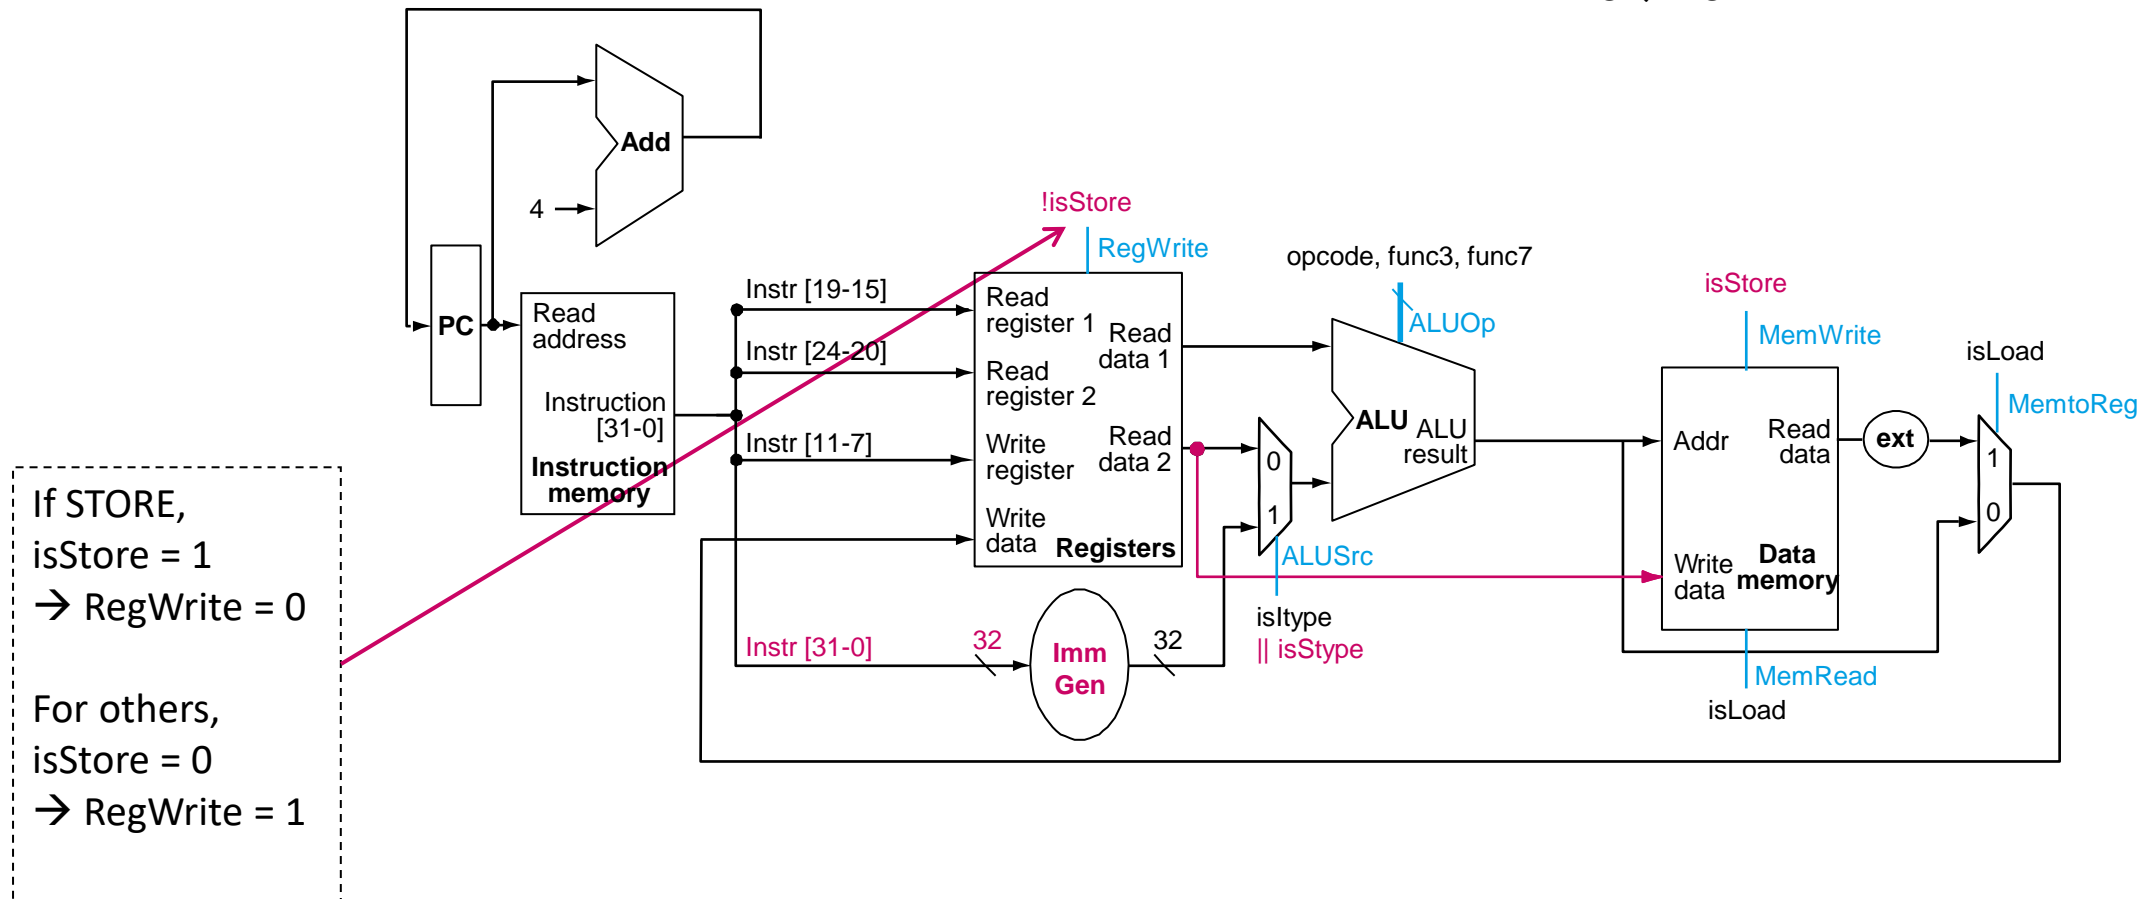


Load/Store Datapath

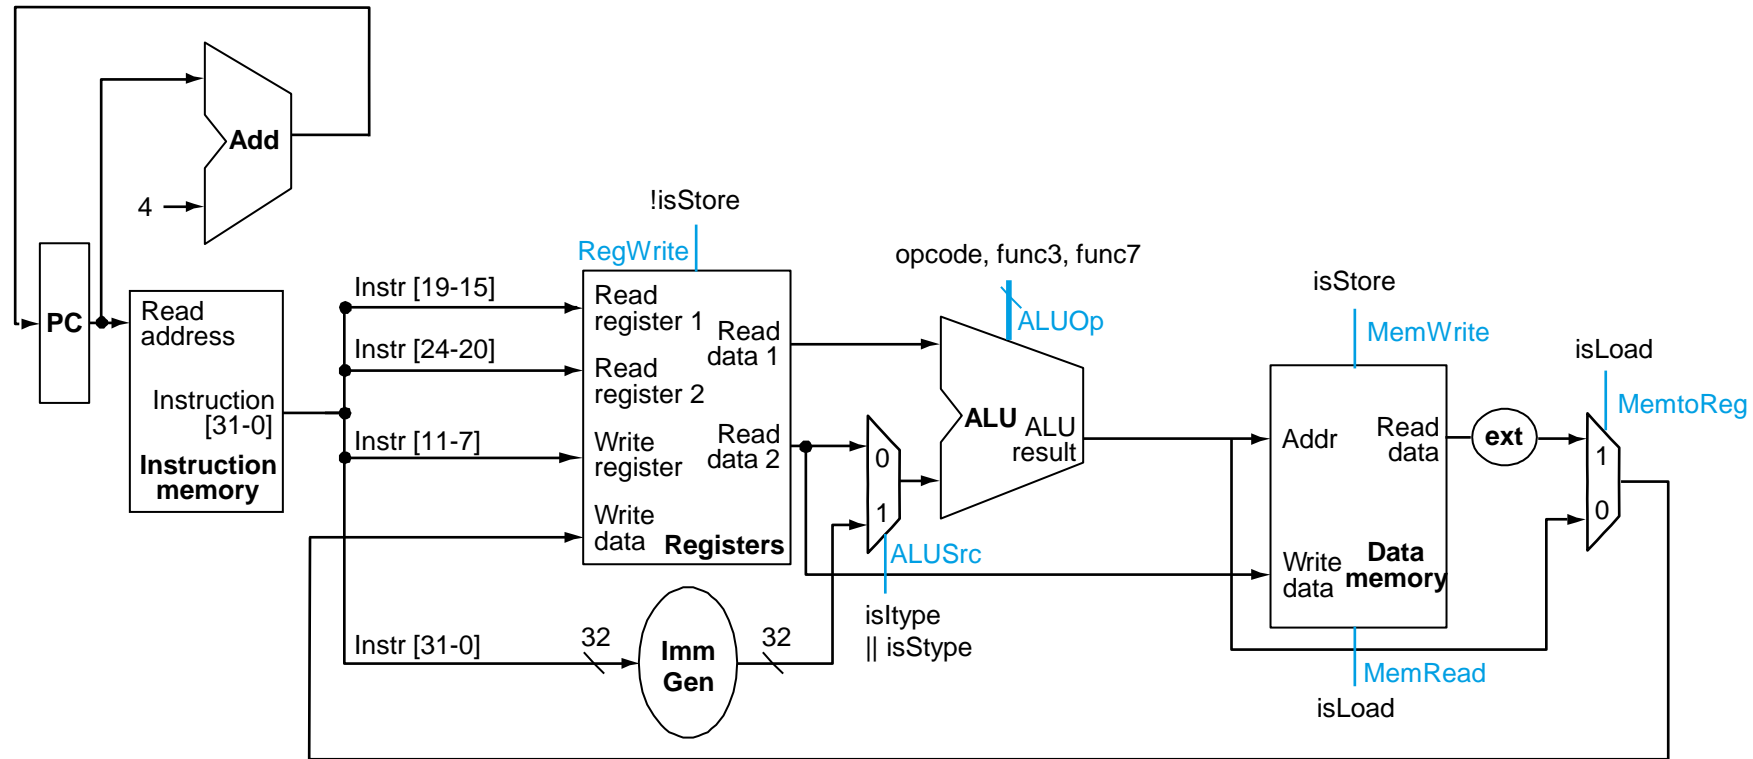


◆ Semantics:

- $\text{byte_address}_{32} = \text{sign-extend}(\text{offset}_{12}) + \text{GPR}[\text{base}]$
- $\text{MEM}_{32}[\text{byte_address}_{32}] \leftarrow \text{GPR}[\text{rs2}]$
- $\text{PC} \leftarrow \text{PC} + 4$




Datapath for Non-Control Flow Instructions



Single-Cycle Datapath for Control Flow Instructions

Reminder: Jump Instruction (UJ-Type)


- Jump and Link assembly: JAL rd, imm₂₁
- Encoding:



- Note: imm[0] == 0 (only branch to even addresses)
- Semantics:
 - target = PC + sign-extend(imm₂₁) // PC-relative addressing
 - GPR[rd] ← PC + 4 // store return address in a register (cf. BEQ/BNE/...)
 - PC ← target // always jump to the same target
- Exception: misaligned target (4-byte)

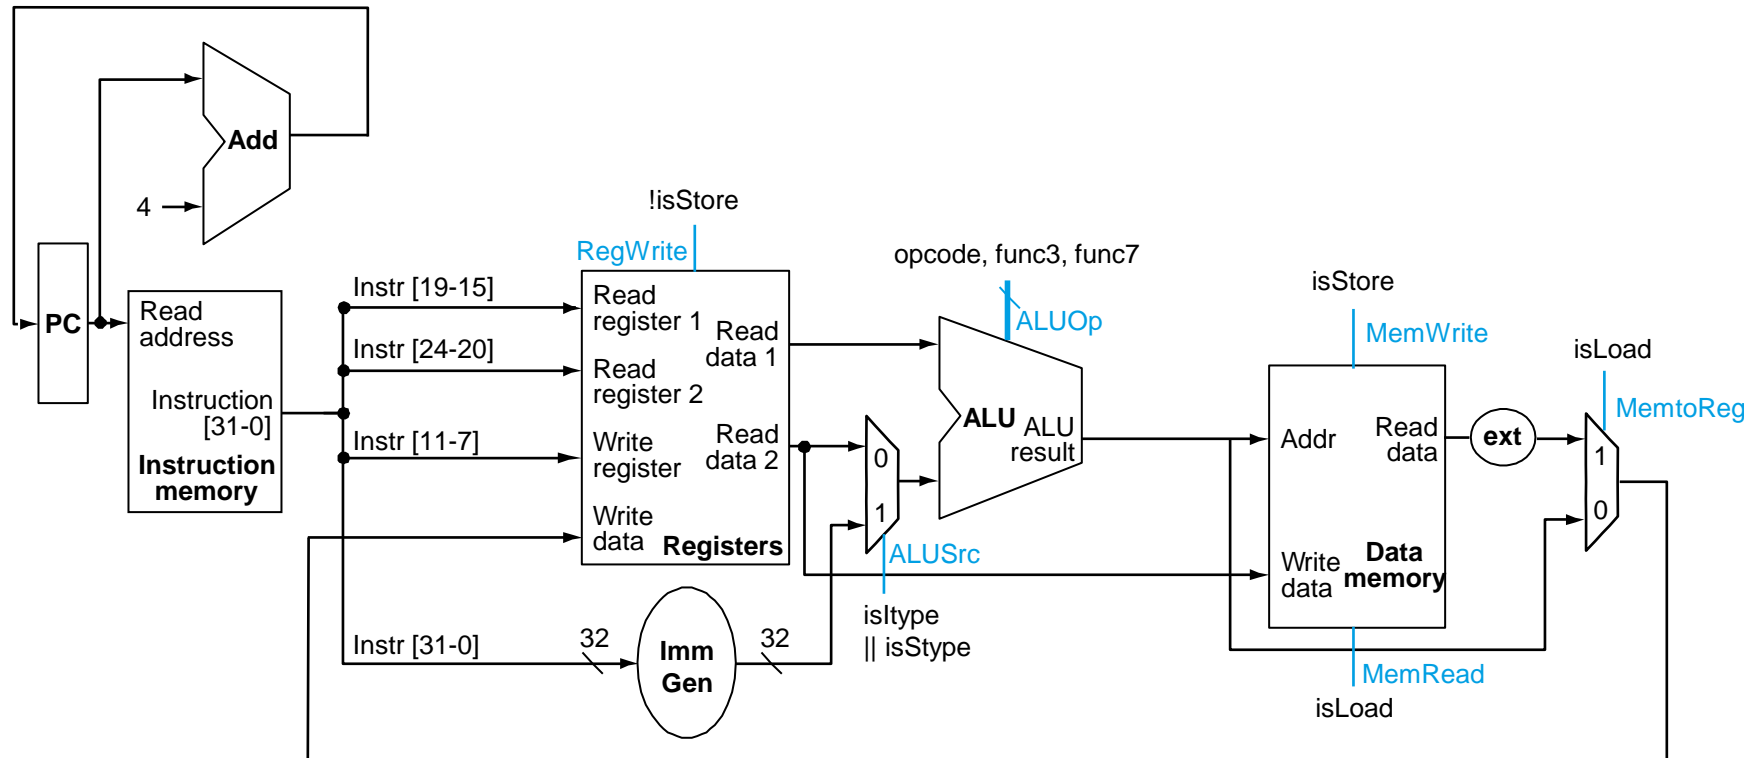
Unconditional Jump Datapath (JAL)

■ How should this change for JAL?




◆ Semantics:

- $\text{target} = \text{PC} + \text{sign-extend}(\text{imm}_{21})$
- $\text{GPR}[\text{rd}] \leftarrow \text{PC} + 4$
- $\text{PC} \leftarrow \text{target}$



Unconditional Jump Datapath (JAL)



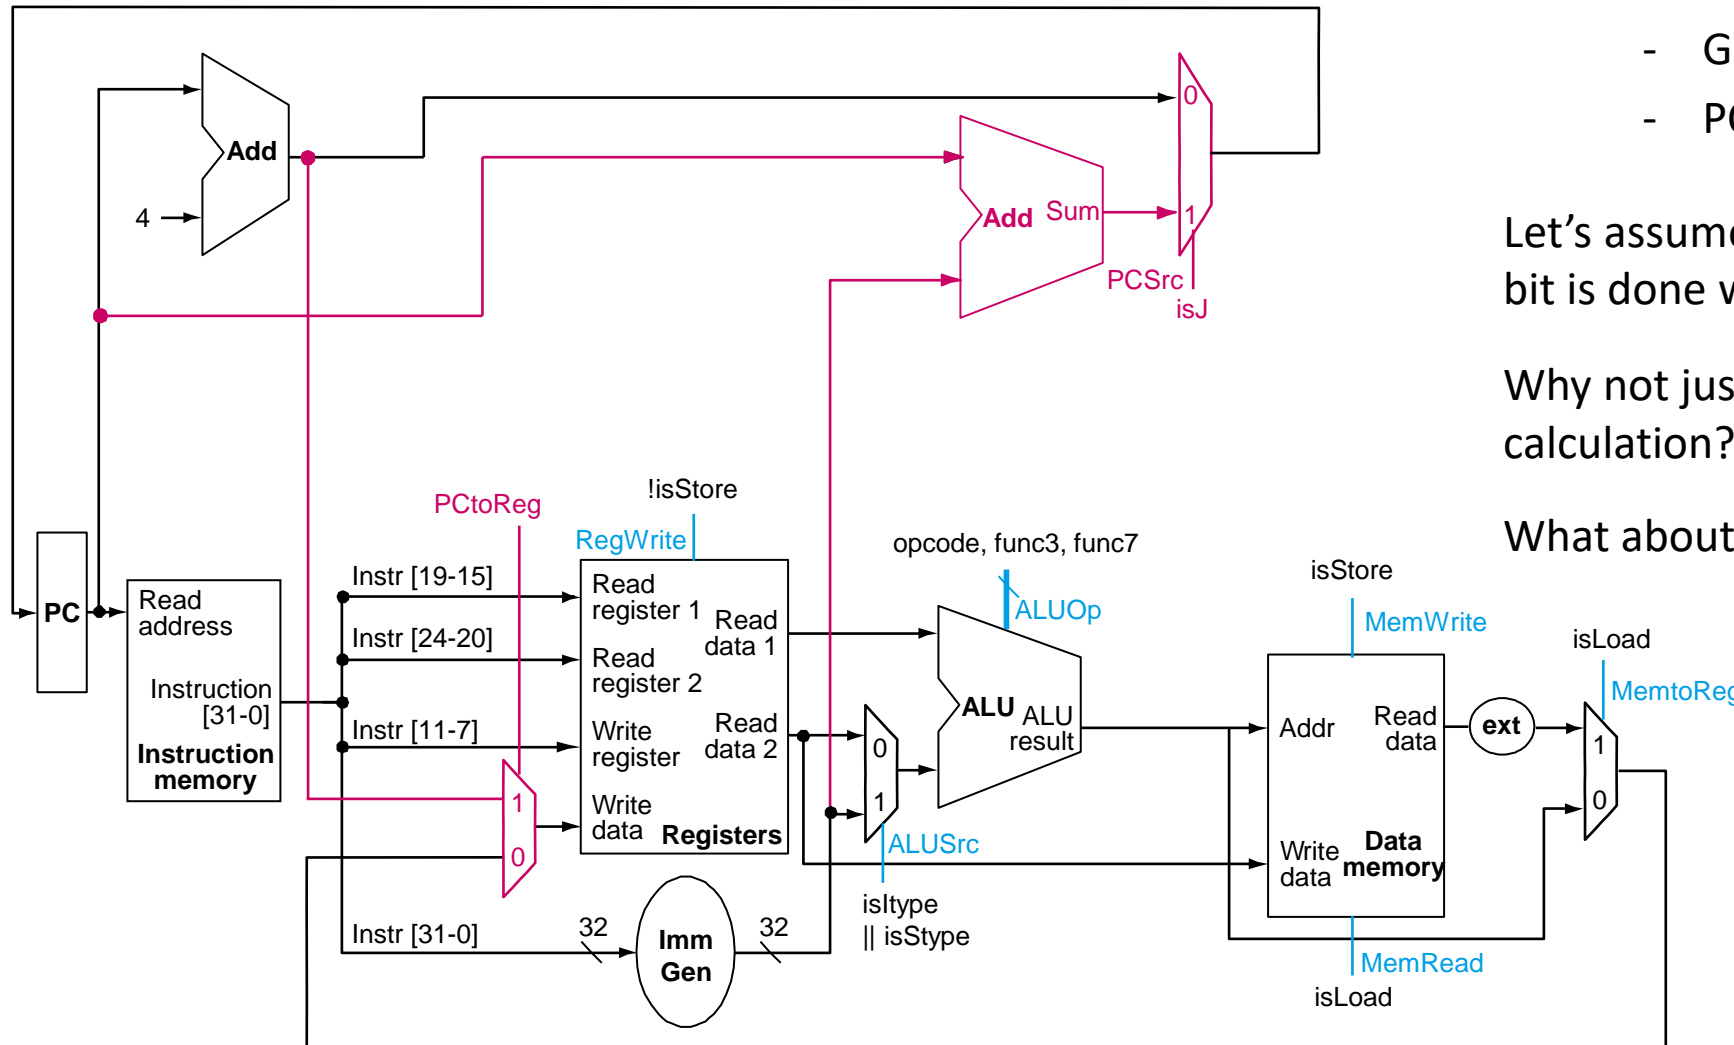
◆ Semantics:

- $\text{target} = \text{PC} + \text{sign-extend}(\text{imm}_{21})$
- $\text{GPR}[\text{rd}] \leftarrow \text{PC} + 4$
- $\text{PC} \leftarrow \text{target}$

Let's assume shifting immediate by one bit is done within "Imm. gen."

Why not just use existing ALU for target calculation?


What about JALR?



Reminder: Jump Indirect Instruction (I-Type)

■ Jump and Link Register assembly: `JALR rd, imm12(rs1)`

■ Encoding:



■ Semantics:

- $\text{target} = \text{GPR}[\text{rs1}] + \text{sign-extend}(\text{imm}_{12})$ // can jump to different targets
- $\text{target} \&= 0\text{FFFFFFFE}$ // set the lowest bit to 0 (but let's ignore this for now)
- $\text{GPR}[\text{rd}] \leftarrow \text{PC} + 4$
- $\text{PC} \leftarrow \text{target}$

■ Usages: - Procedure return: `jalr x0, 0(x1)` // return to GPR[x1]. x0 can't change
- Computed jumps (e.g., case/switch statements)

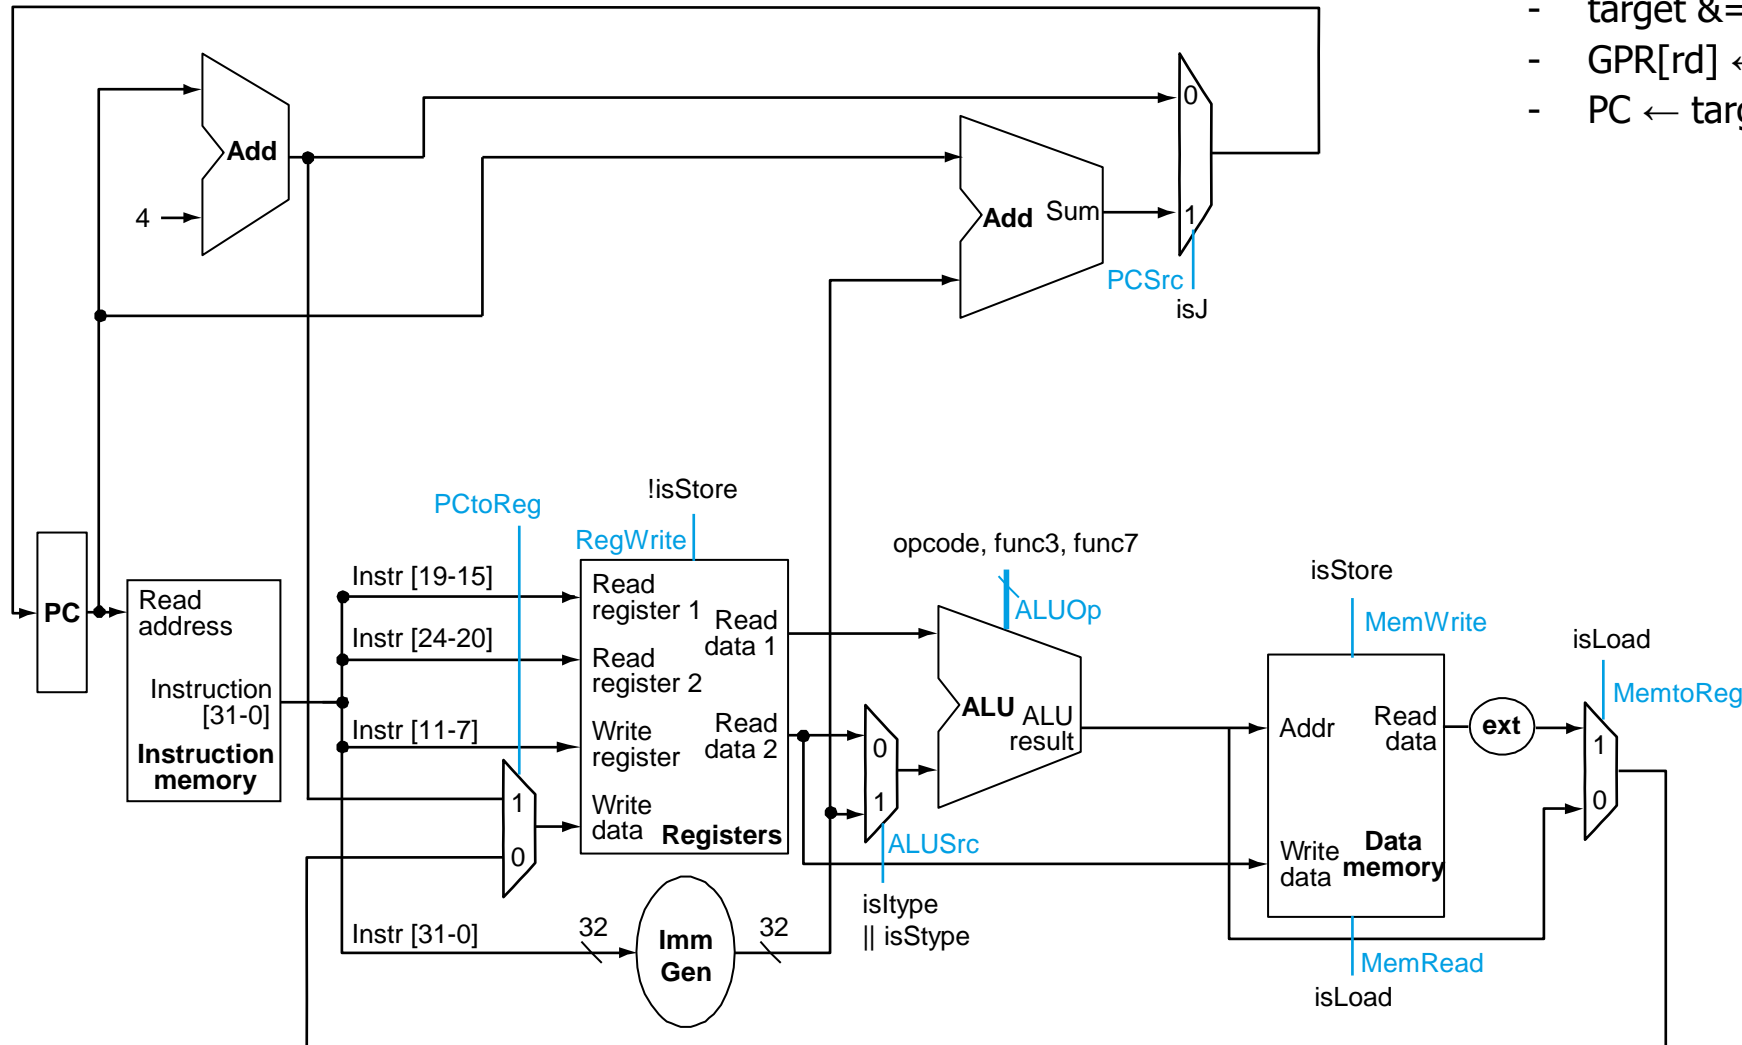
■ Exception: misaligned target (4-byte)

Unconditional Jump Datapath (JALR)

■ How should this change for JALR?

◆ Semantics:

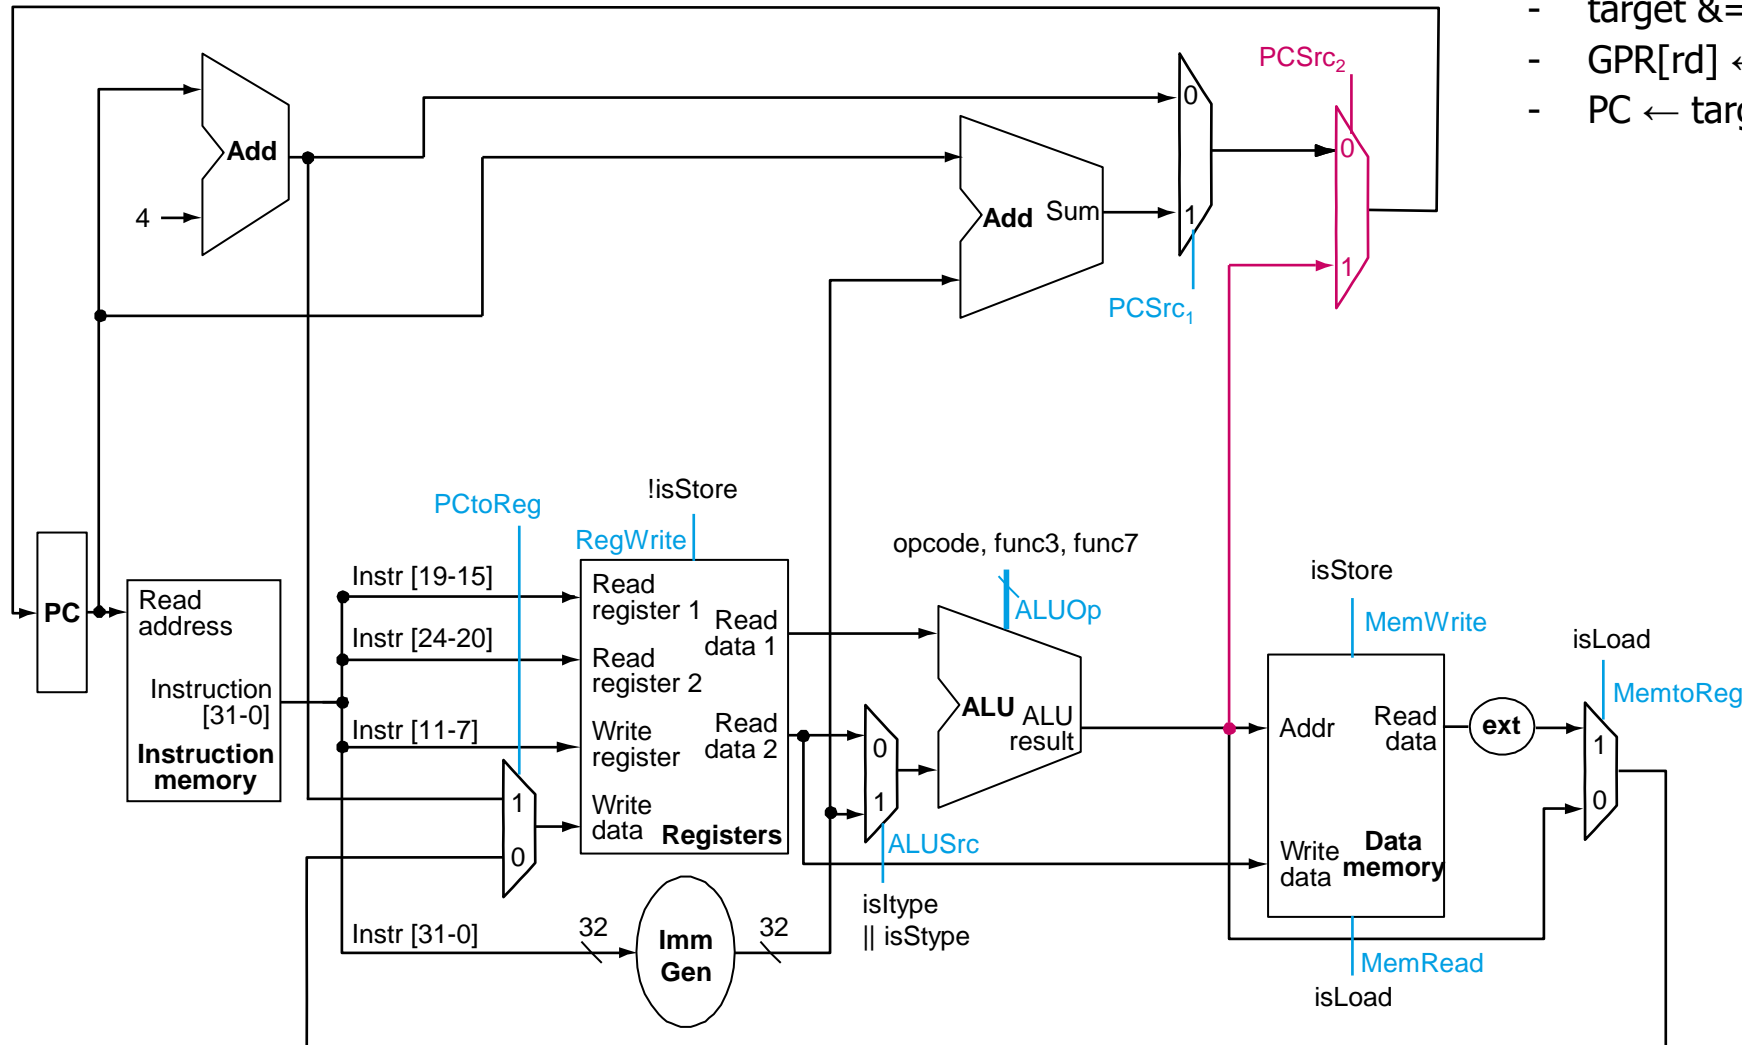
- $\text{target} = \text{GPR}[\text{rs1}] + \text{sign-extend}(\text{imm}_{12})$
- $\text{target} \&= 0\text{FFFFFFFE}$ //ignored for simplicity
- $\text{GPR}[\text{rd}] \leftarrow \text{PC} + 4$
- $\text{PC} \leftarrow \text{target}$



Unconditional Jump Datapath (JALR)

◆ Semantics:


- $\text{target} = \text{GPR}[\text{rs1}] + \text{sign-extend}(\text{imm}_{12})$
- $\text{target} \&= 0\text{FFFFFFFE}$ //ignored for simplicity
- $\text{GPR}[\text{rd}] \leftarrow \text{PC} + 4$
- $\text{PC} \leftarrow \text{target}$



Reminder: Cond. Branch Instructions (SB-Type)

- Assembly example (Branch if Equal): `BEQ rs1, rs2, imm13`

- Encoding:



— Note: `imm[0] == 0` (only branch to even addresses) – is this okay?

- Semantics:

— `Target = PC + sign-extend(imm13)` // PC-relative addressing -- how far can you jump?

— If `GPR[rs1] == GPR[rs2]` then `PC ← target`
else `PC ← PC + 4`

- Exception: misaligned target (4-byte) if taken

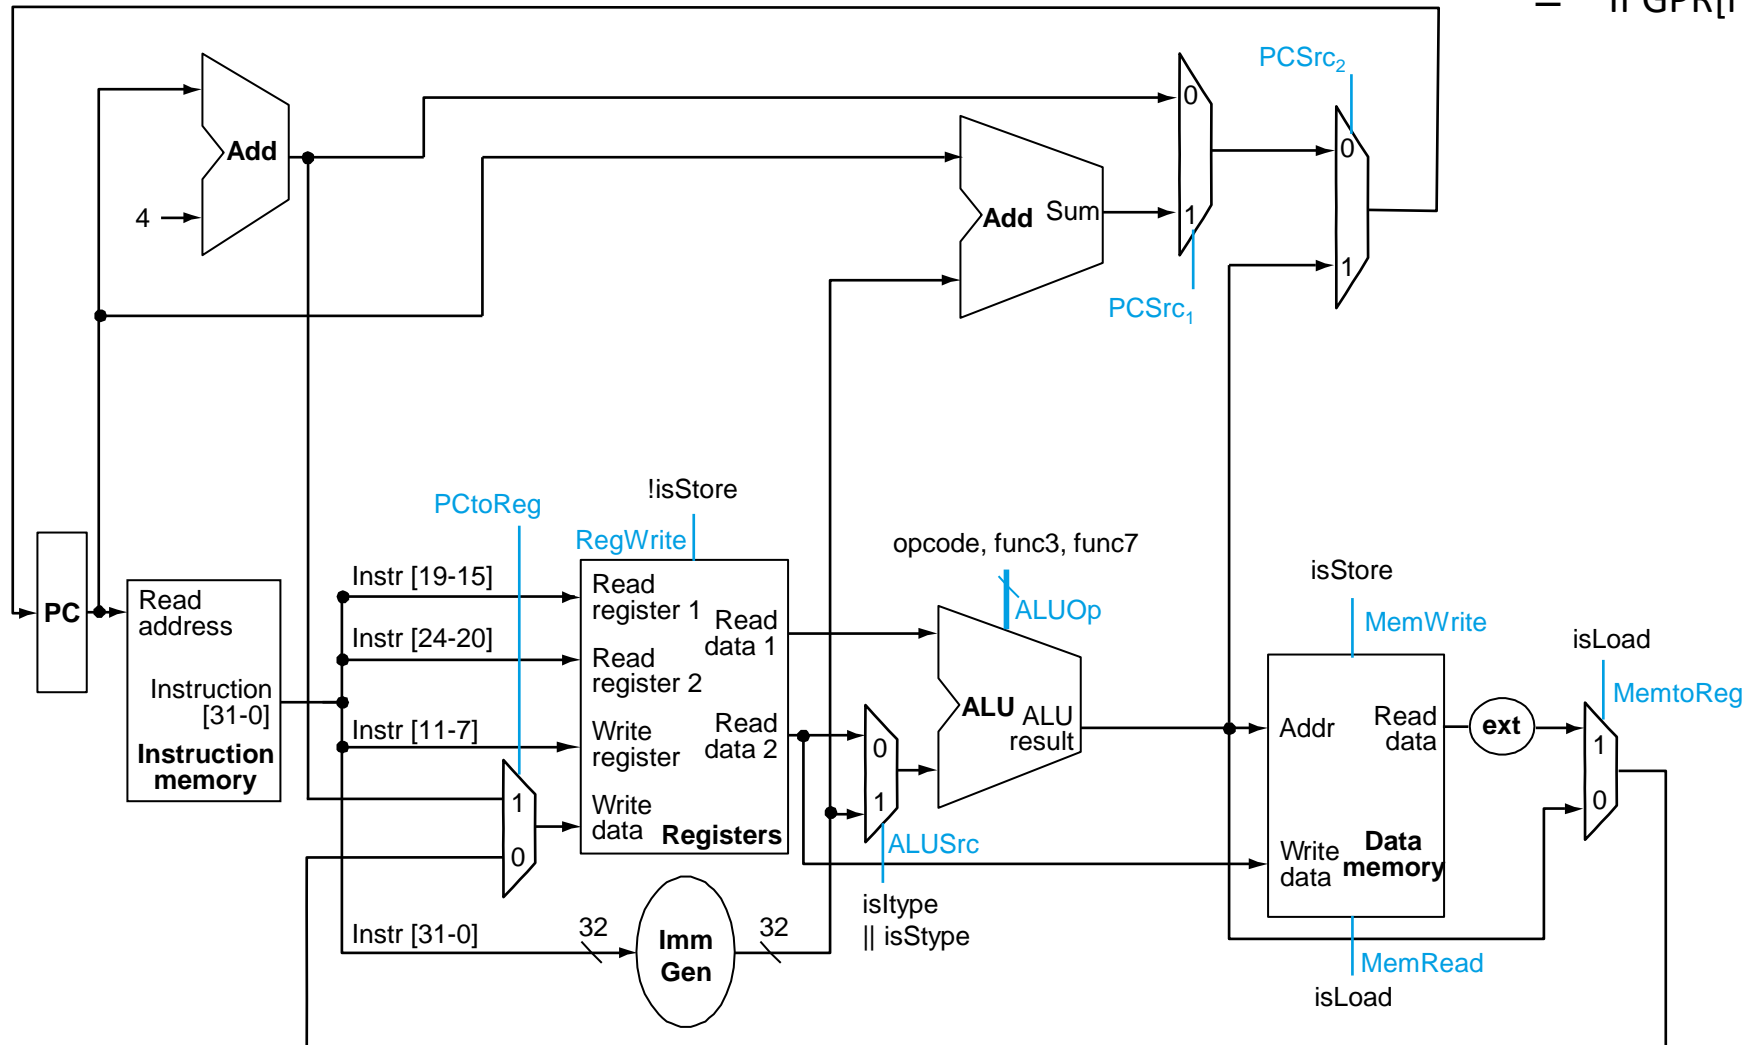
- Variations

— `BEQ, BNE, BLT, BGE` (signed variations) – why no `BGT` (Greater Than)?
— `BLTU, BGEU` (unsigned variations)

Conditional Branch Datapath

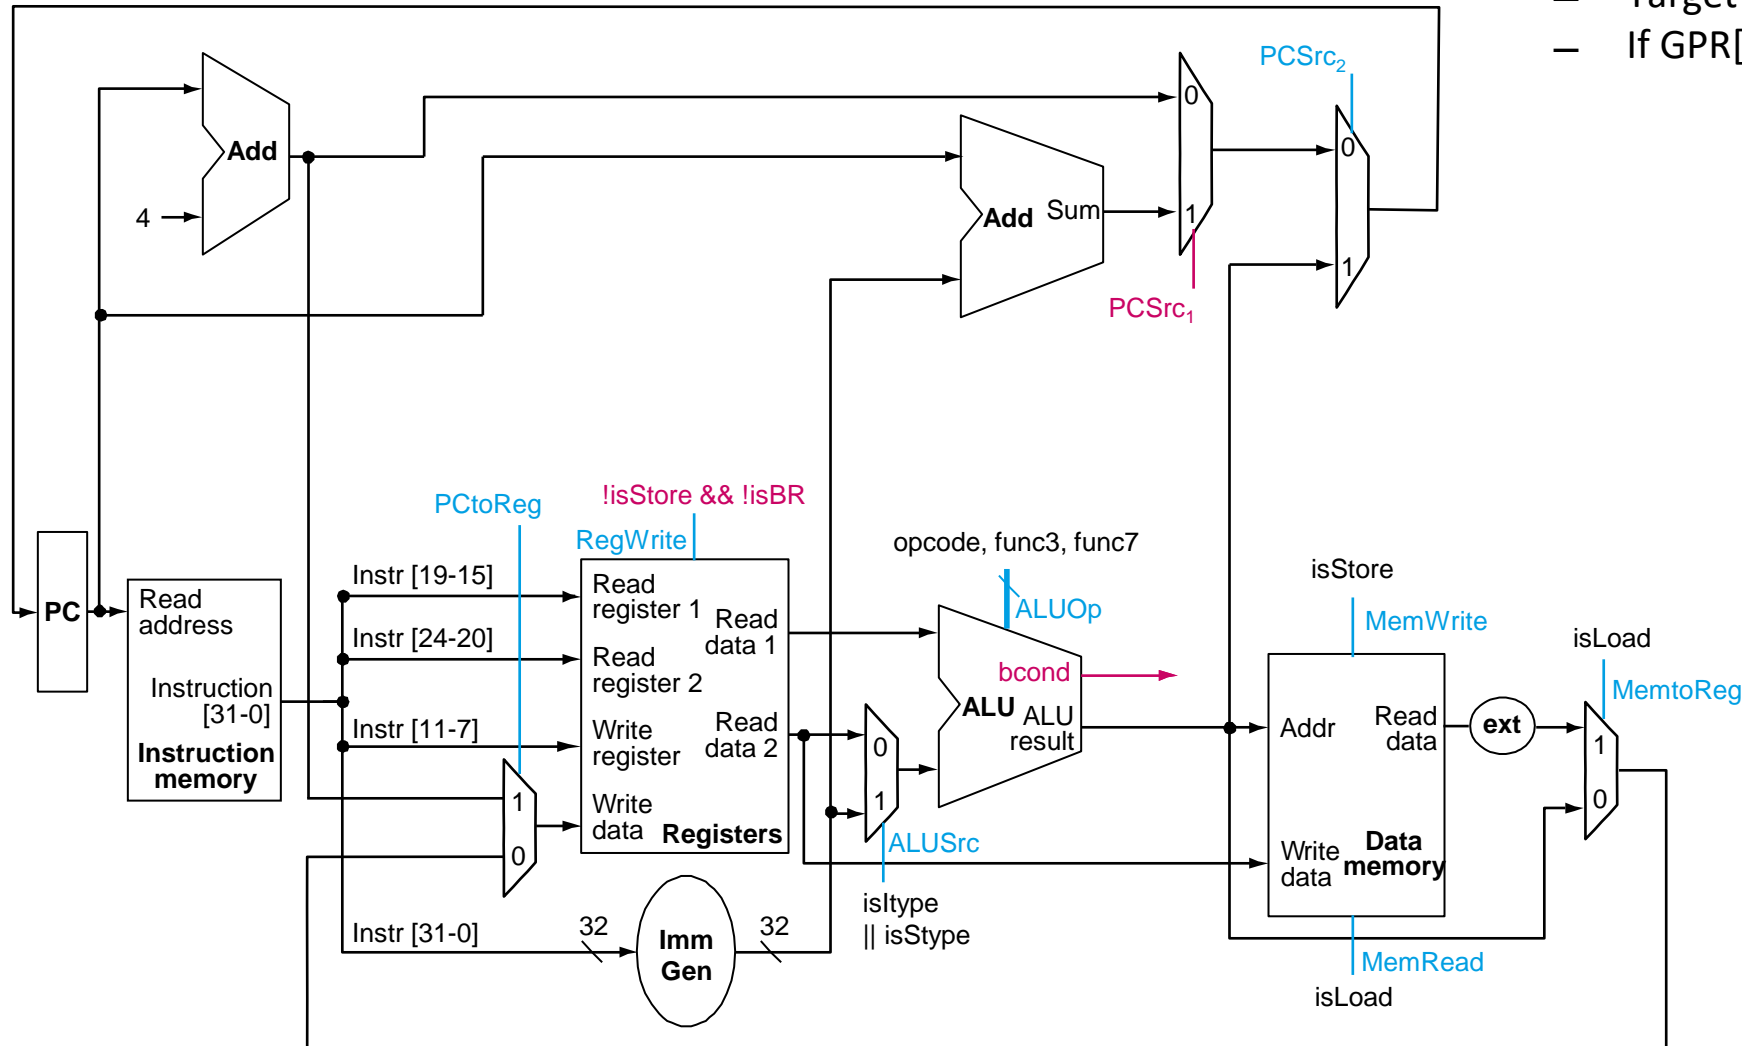
- How should this change for conditional branches?

- Semantics (BEQ):
 - Target = PC + sign-extend(imm₁₃)
 - If GPR[rs1] == GPR[rs2] then PC ← target
else PC ← PC + 4



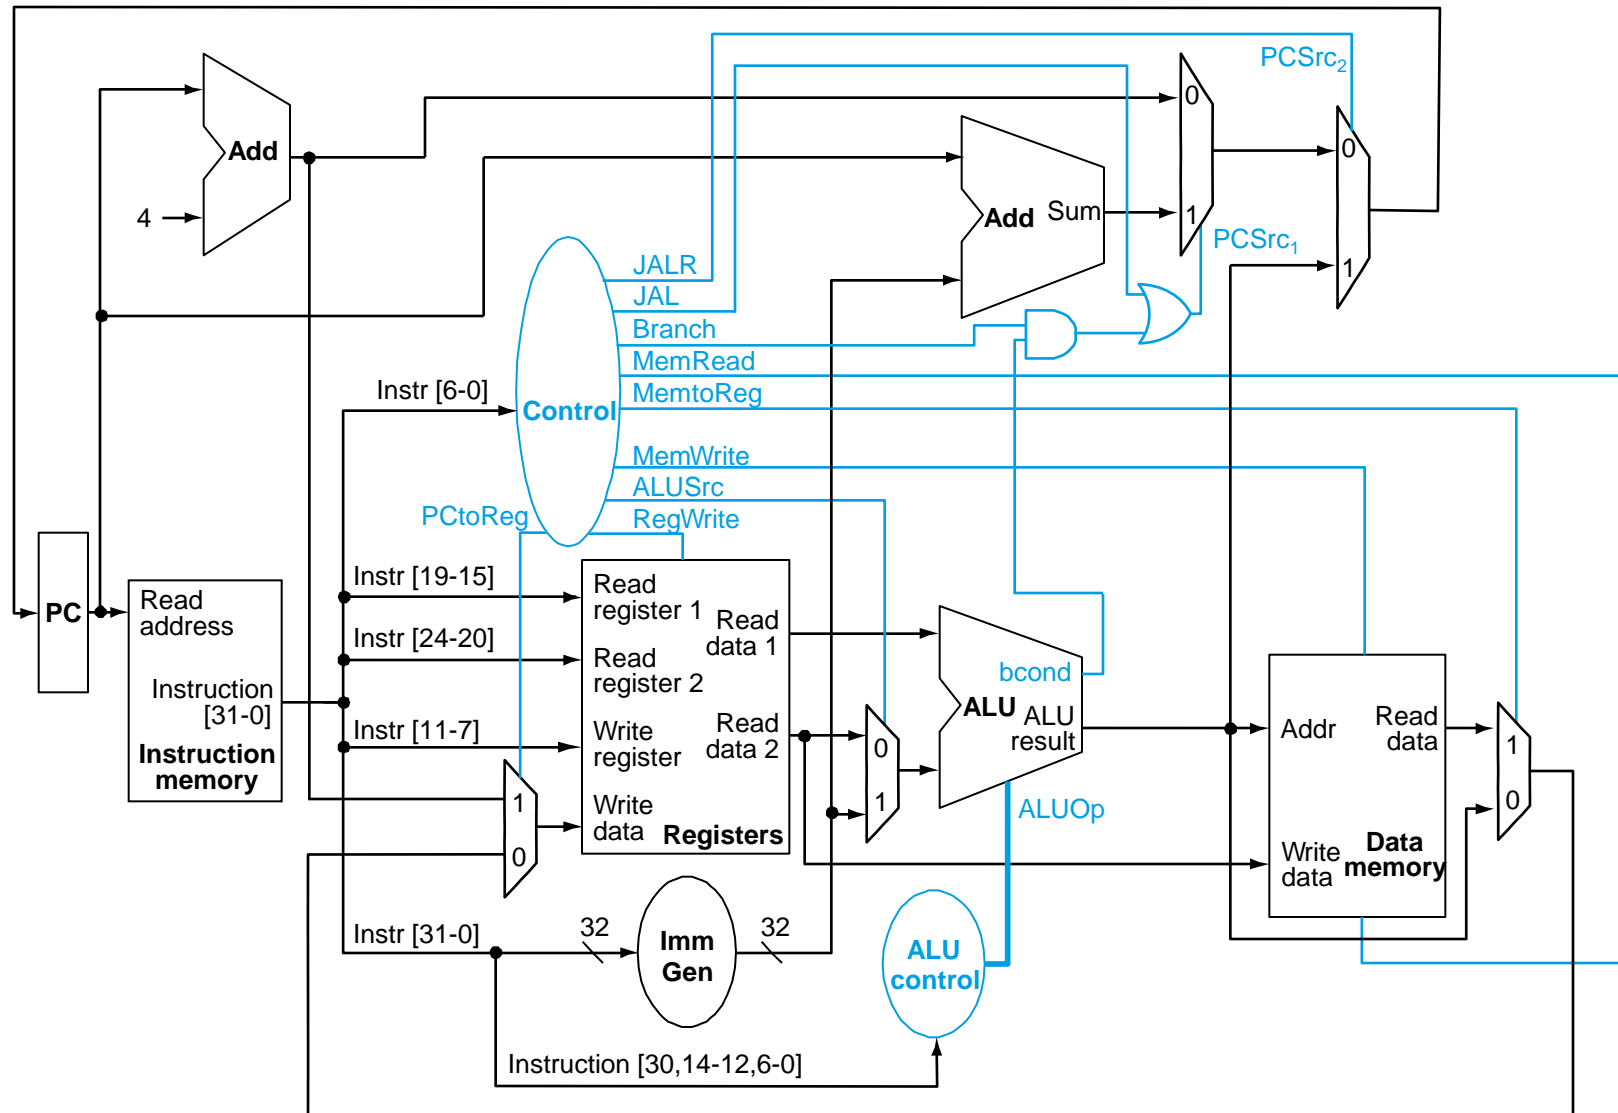
Conditional Branch Datapath

- Semantics (BEQ):
 - Target = PC + sign-extend(imm_{13})
 - If $\text{GPR}[\text{rs1}] == \text{GPR}[\text{rs2}]$ then $\text{PC} \leftarrow \text{target}$
else $\text{PC} \leftarrow \text{PC} + 4$



Control

Datapath with Control



Single-bit Control Values

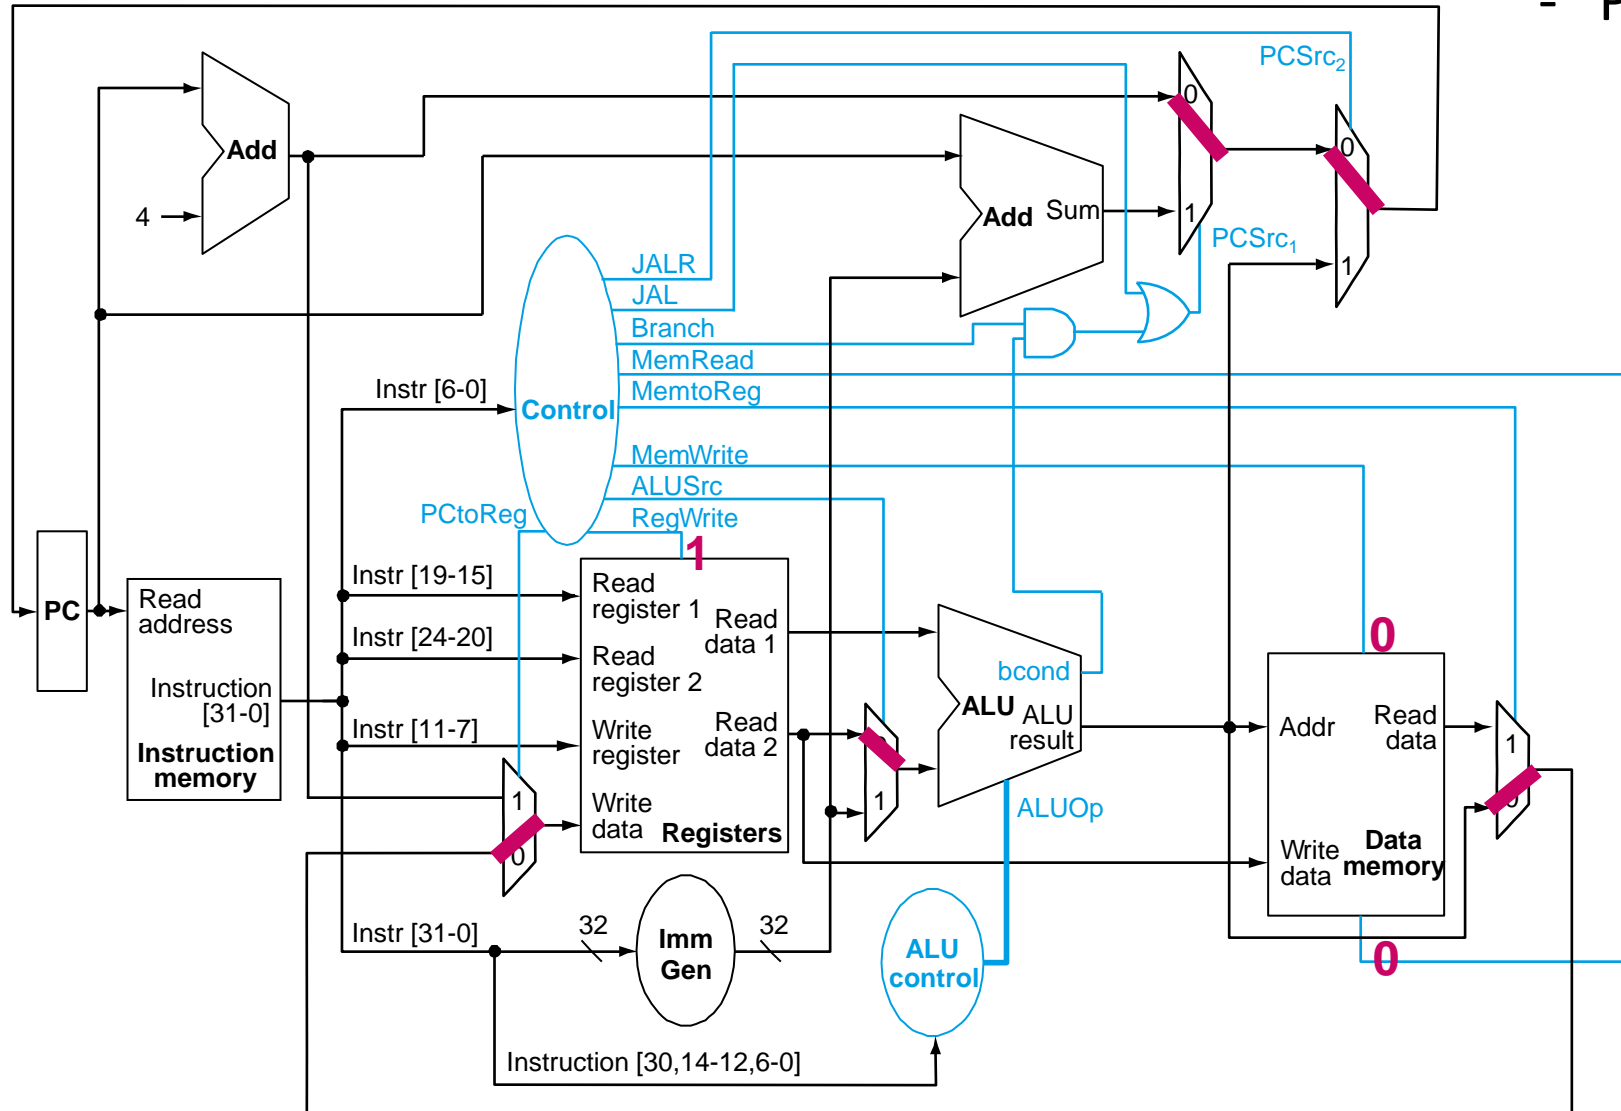
	When de-asserted	When asserted	Equation
RegWrite	GPR write disabled	GPR write enabled	$(\text{opcode} \neq \text{SW/SH/SB}) \ \&\& \ (\text{opcode} \neq \text{Bxx})$
ALUSrc	2 nd ALU input from 2 nd GPR read port	2 nd ALU input from immediate	$(\text{opcode} \neq \text{isRtype}) \ \&\& \ (\text{opcode} \neq \text{isSBtype})$
MemRead	Memory read disabled	Memory read port return load value	$\text{opcode} == \text{LW/LH/LB}$
MemWrite	Memory write disabled	Memory write enabled	$\text{opcode} == \text{SW/SH/SB}$
MemtoReg	Steer ALU result to GPR write port	Steer memory load to GPR write port	$\text{opcode} == \text{LW/LH/LB}$
PCtoReg	Steer above result to GPR write port	Steer PC+4 to GPR write port	$\text{opcode} == \text{JAL/JALR}$
PCSrc₁	Next PC = PC + 4	Next PC = PC + immediate	$\text{opcode} == \text{JAL} \ \ (\text{opcode} == \text{isSBtype} \ \&\& \ \text{bcond})$
PCSrc₂	Next PC is determined by PCSrc₁	Next PC = GPR + immediate	$\text{opcode} == \text{JALR}$

Multi-bit Control Values

	Options	Equation
ALUOp	<ul style="list-style-type: none">• ADD, SUB, AND, OR, XOR, NOR, LT, and Shift• btype: EQ, NE, GE, LT	case opcode RTypeALU : according to funct3 , funct7[5] ITypeALU : according to funct3 only (except shift) LW/SW/JALR : Add Bxx : Subtract and select btype
ImmGen	I-type, I-type(unsigned), S-type, SB-type, U-type, UJ-type	<ul style="list-style-type: none">• Select based on instruction format type

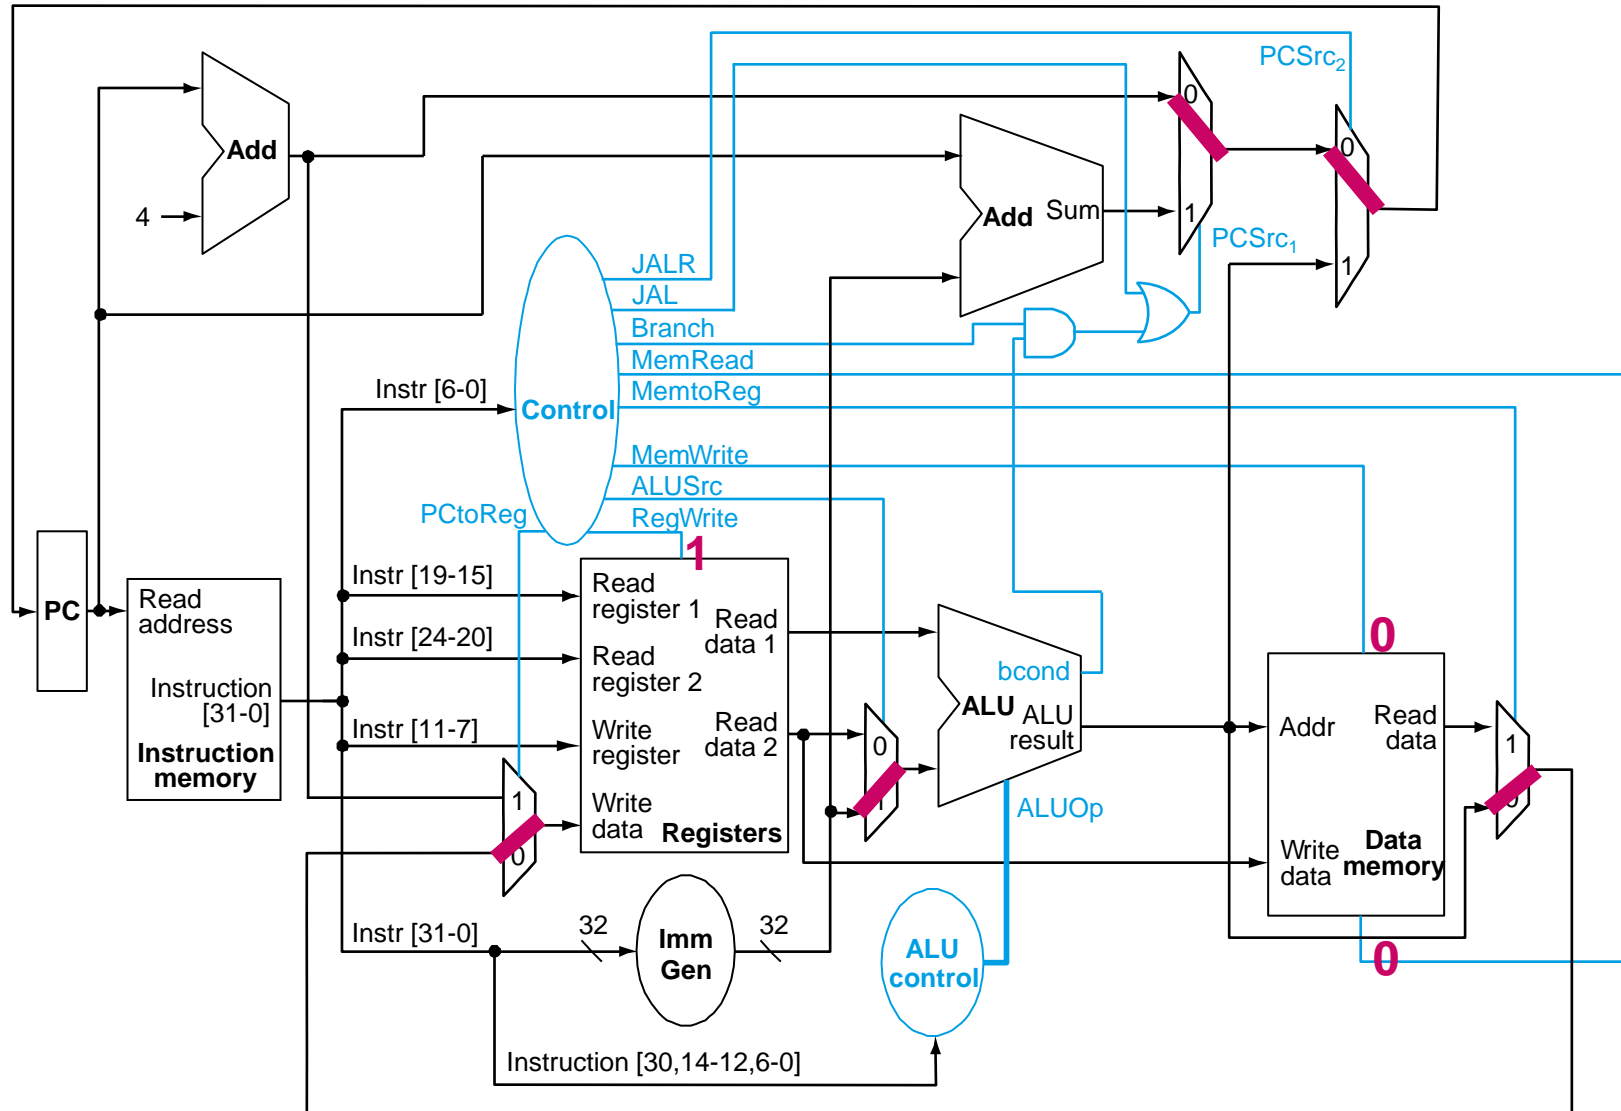
R-Type ALU

- $\text{GPR}[\text{rd}] \leftarrow \text{GPR}[\text{rs1}] (\text{op}) \text{GPR}[\text{rs2}]$
- $\text{PC} \leftarrow \text{PC} + 4$



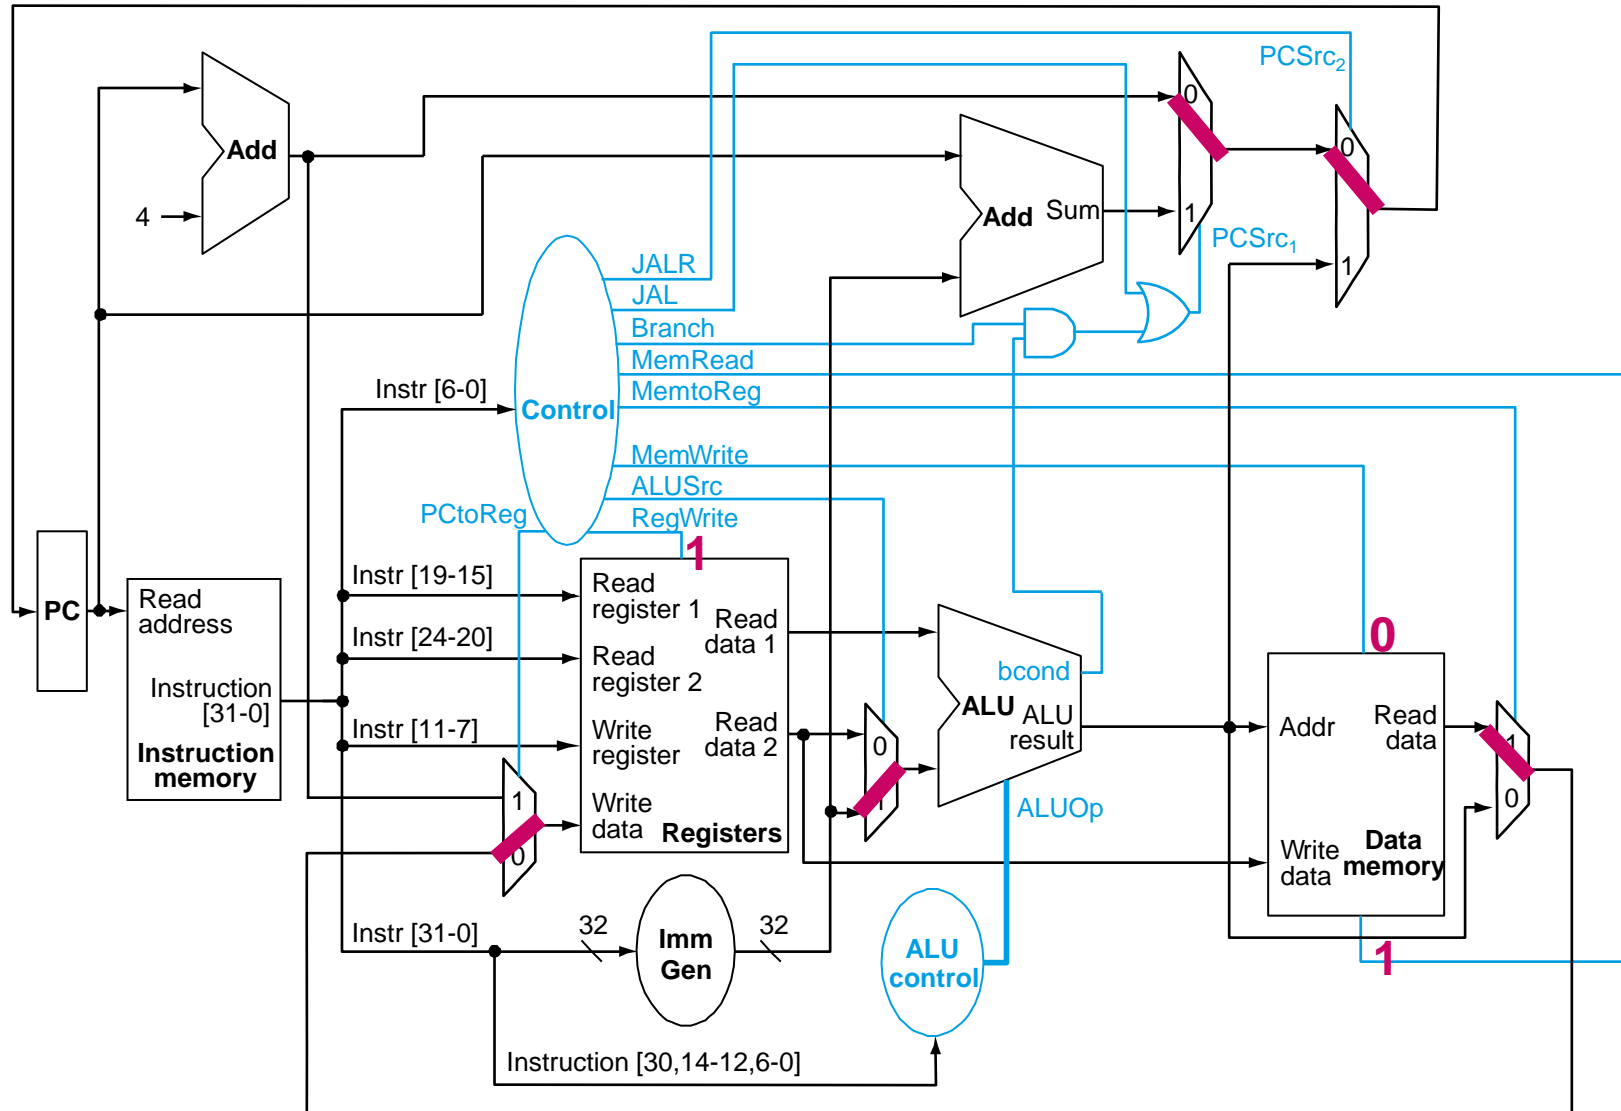
I-Type ALU

- $\text{GPR}[\text{rd}] \leftarrow \text{GPR}[\text{rs1}] \text{ (op) sign-extend}(\text{imm})$
- $\text{PC} \leftarrow \text{PC} + 4$



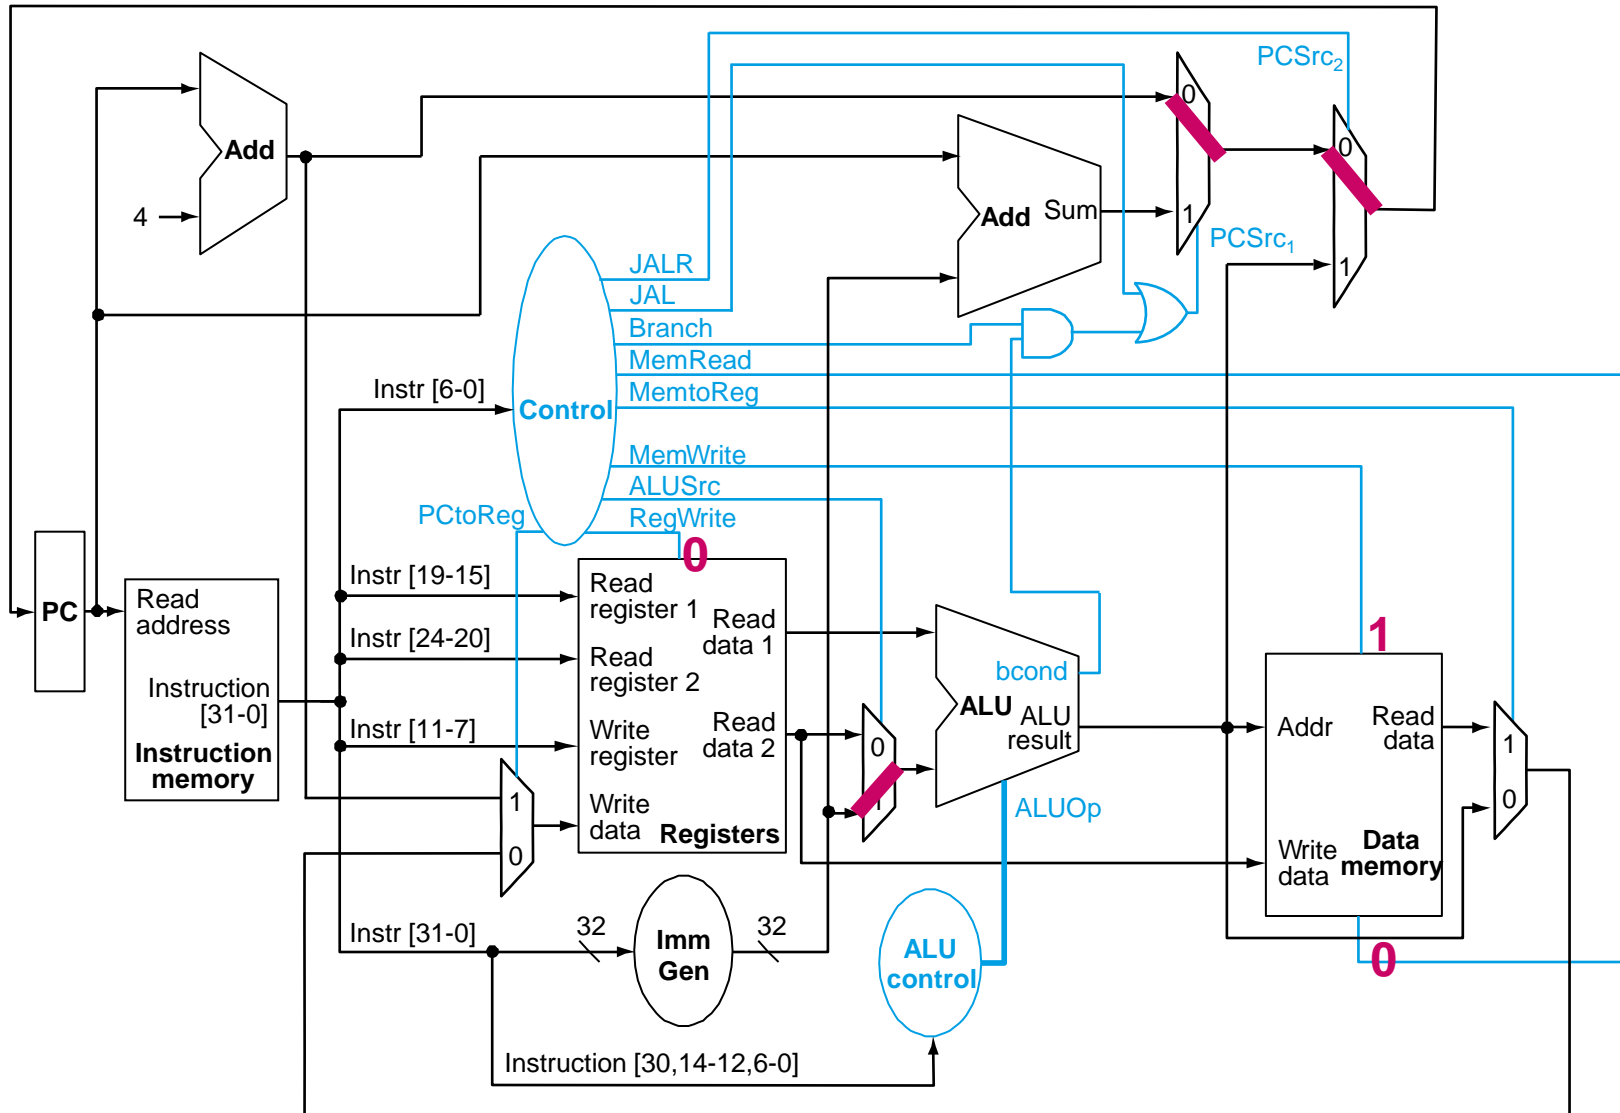
LD

- $\text{byte_addr}_{32} = \text{sign-extend}(\text{offset}_{12}) + \text{GPR}[\text{base}]$
- $\text{GPR}[\text{rd}] \leftarrow \text{MEM}_{32}[\text{byte_addr}_{32}]$
- $\text{PC} \leftarrow \text{PC} + 4$



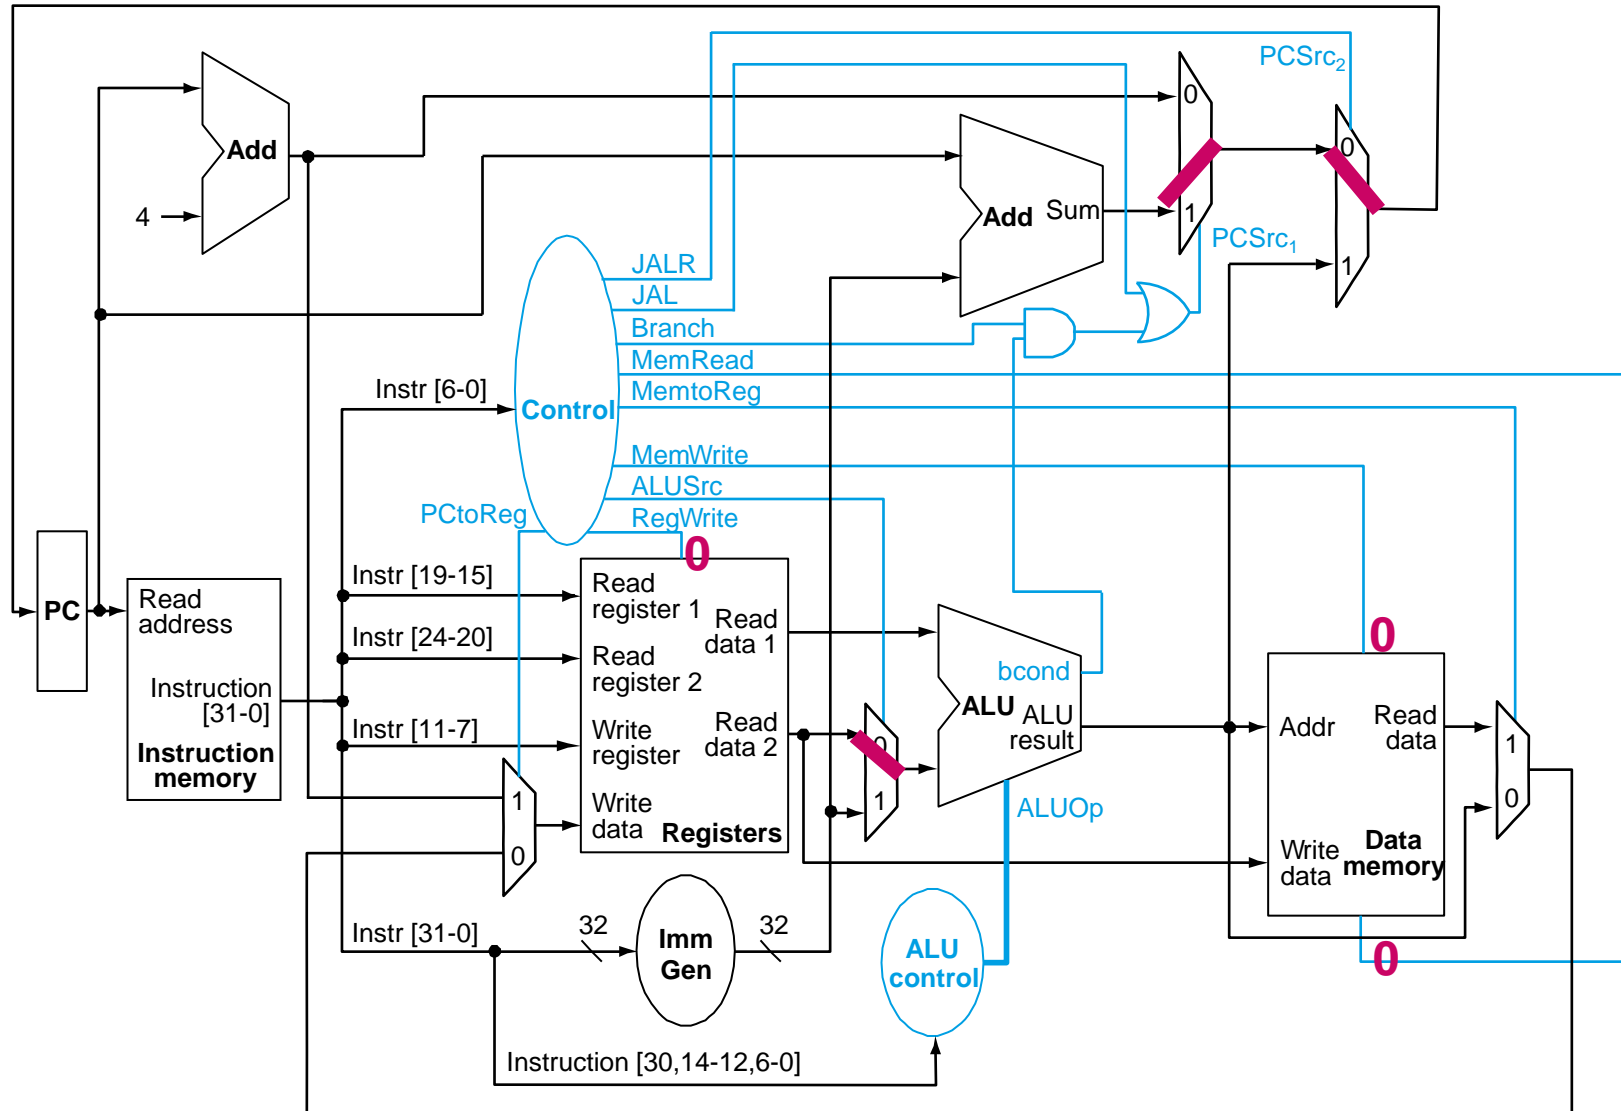
SD

- $\text{byte_address}_{32} = \text{sign-extend}(\text{offset}_{12}) + \text{GPR}[\text{base}]$
- $\text{MEM}_{32}[\text{byte_address}_{32}] \leftarrow \text{GPR}[\text{rs2}]$
- $\text{PC} \leftarrow \text{PC} + 4$



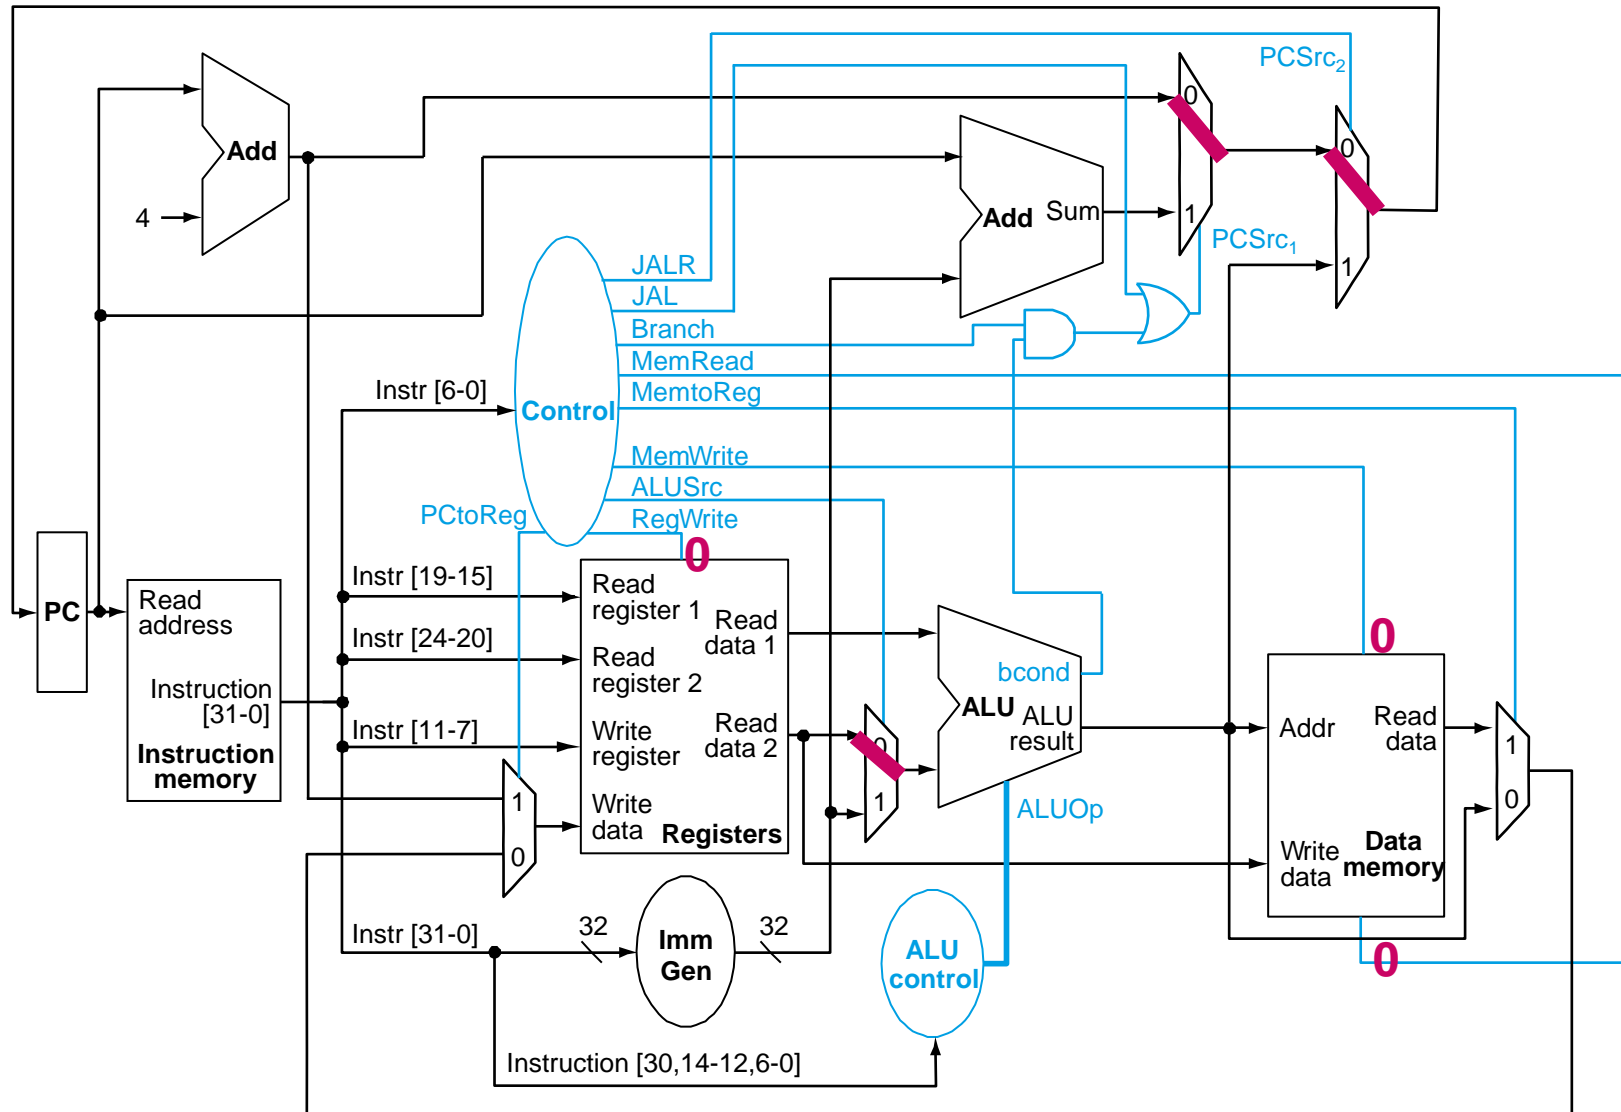
Branch Taken

- Target = PC + sign-extend(imm₁₃)
- If GPR[rs1] == GPR[rs2] then PC ← target
else PC ← PC + 4



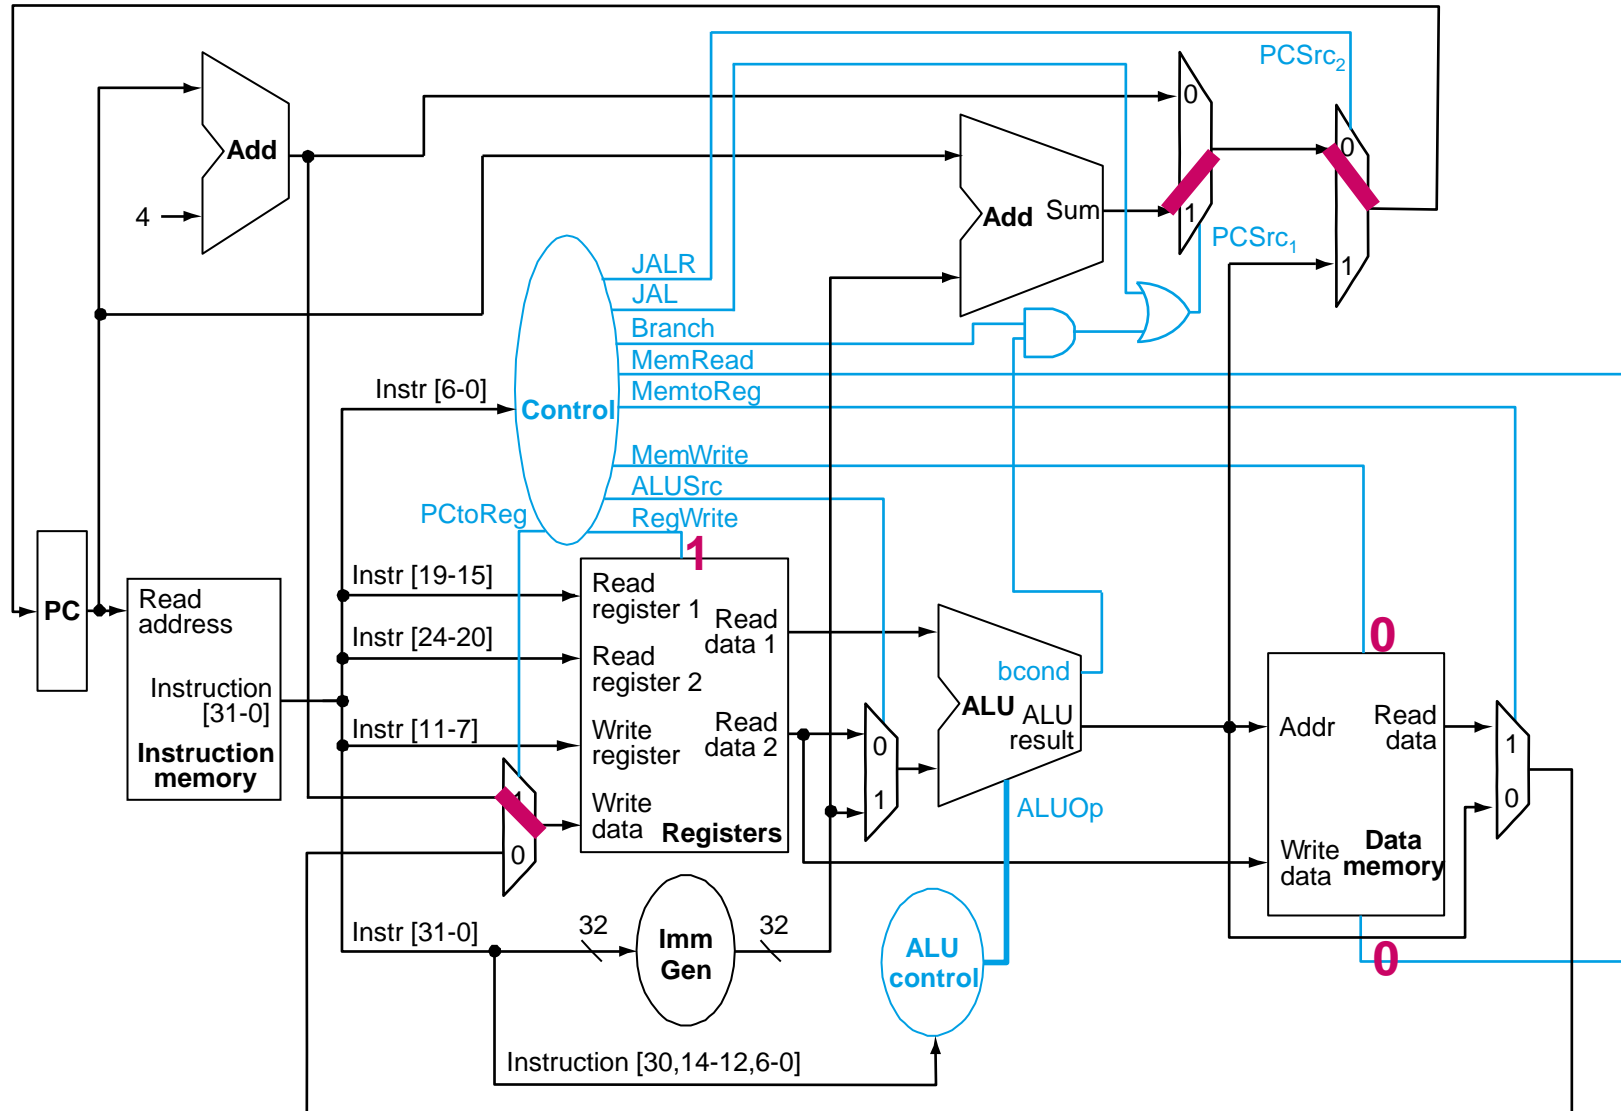
Branch Not Taken

- Target = PC + sign-extend(imm₁₃)
- If GPR[rs1] == GPR[rs2] ~~then PC ← target~~
else PC ← PC + 4



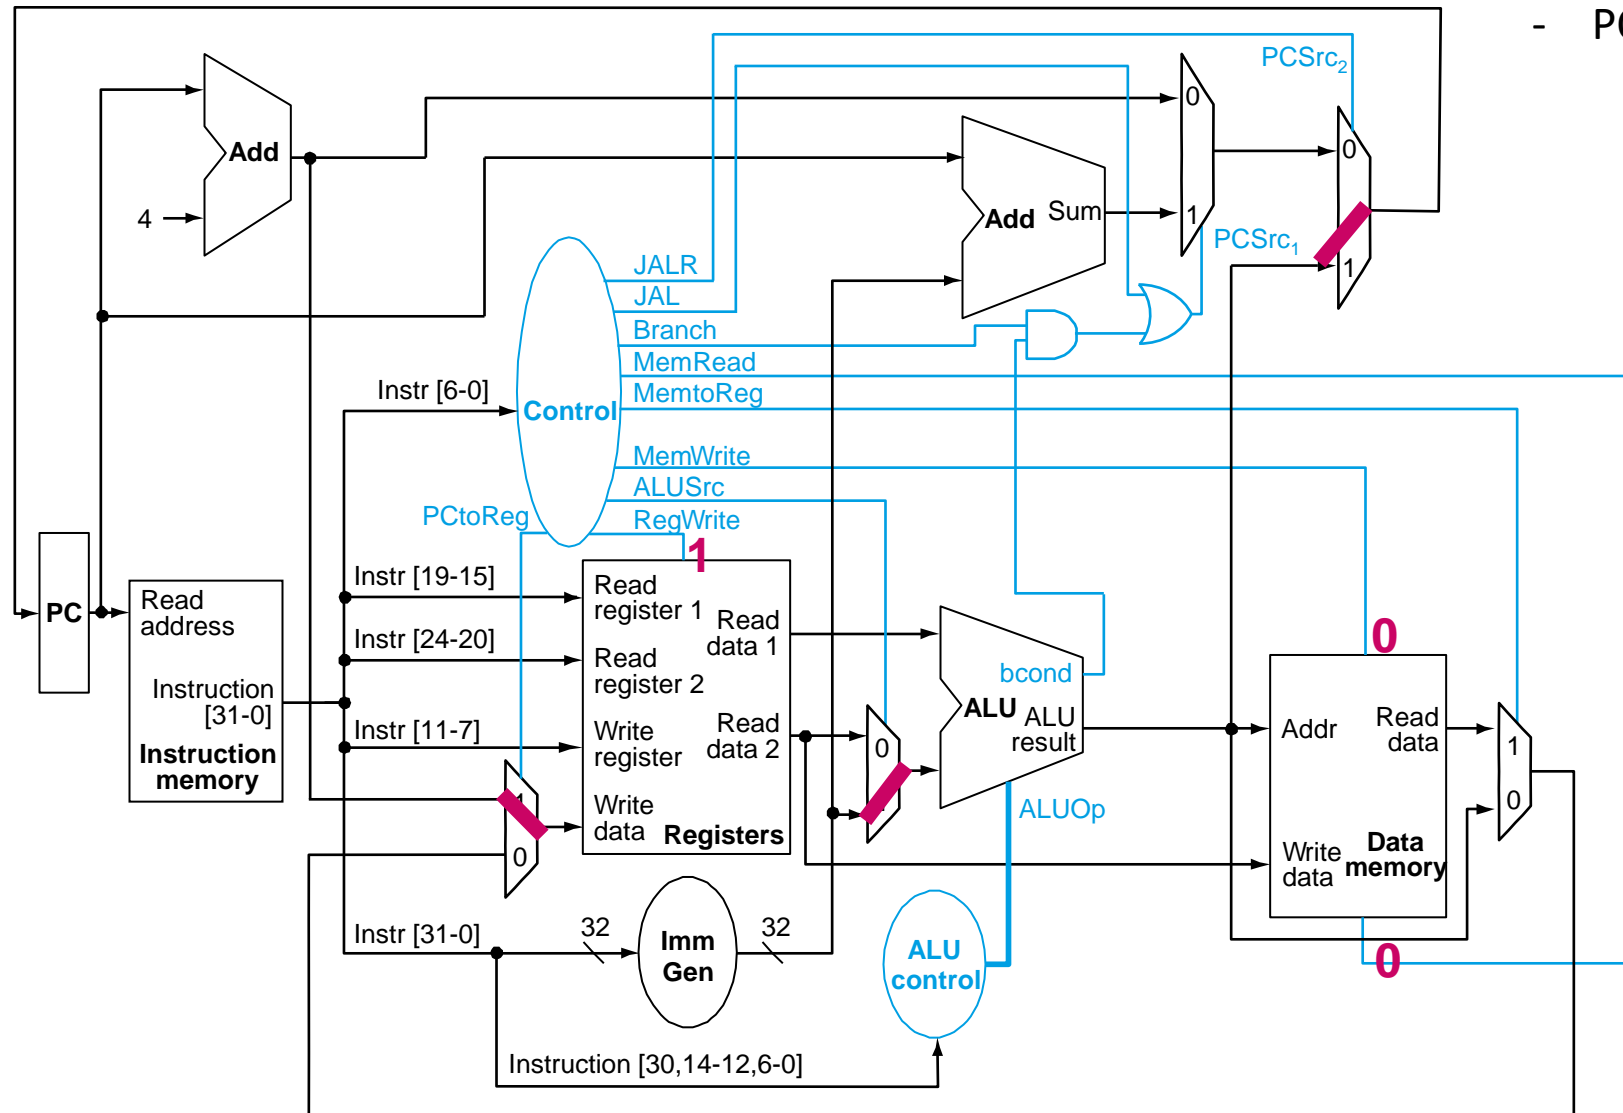
JAL

- $\text{target} = \text{PC} + \text{sign-extend}(\text{imm}_{21})$
- $\text{GPR}[\text{rd}] \leftarrow \text{PC} + 4$
- $\text{PC} \leftarrow \text{target}$



JALR

- $\text{target} = \text{GPR}[\text{rs1}] + \text{sign-extend}(\text{imm}_{12})$
- $\text{target} \&= 0\text{x}\text{FFFFFFFE}$ //ignored for simplicity
- $\text{GPR}[\text{rd}] \leftarrow \text{PC} + 4$
- $\text{PC} \leftarrow \text{target}$



Now you know how to design a single-cycle RISC-V CPU !!

- However, this is a very slow CPU
- We will cover multi-cycle CPU implementations (faster!) in the next lecture

Question?

Announcements

- Reading: P&H (RISC-V ed.) Appendix C