

CSED311 Computer Architecture – Lecture 17


Cache Coherence

Eunhyeok Park

Department of Computer Science and Engineering
POSTECH


Disclaimer: Slides developed in part by Profs. Austin, Brehob, Falsafi, Hill, Hoe, Lipasti, Martin, Roth, Shen, Smith, Sohi, Tyson, Vijaykumar, Mutlu, and Wenisch, Gwangsun Kim @ Carnegie Mellon University, University of Michigan, Purdue University, University of Pennsylvania, University of Wisconsin, and POSTECH

Shared Memory Multiprocessor Abstraction




- What is the key advantage of shared-memory multiprocessor (MP) system?
“Single Physical Address Space”
- Latencies from P_n to Memory
UMA (uniform memory access) vs. NUMA (non-uniform memory access)

Implementation #1: Bus Interconnect



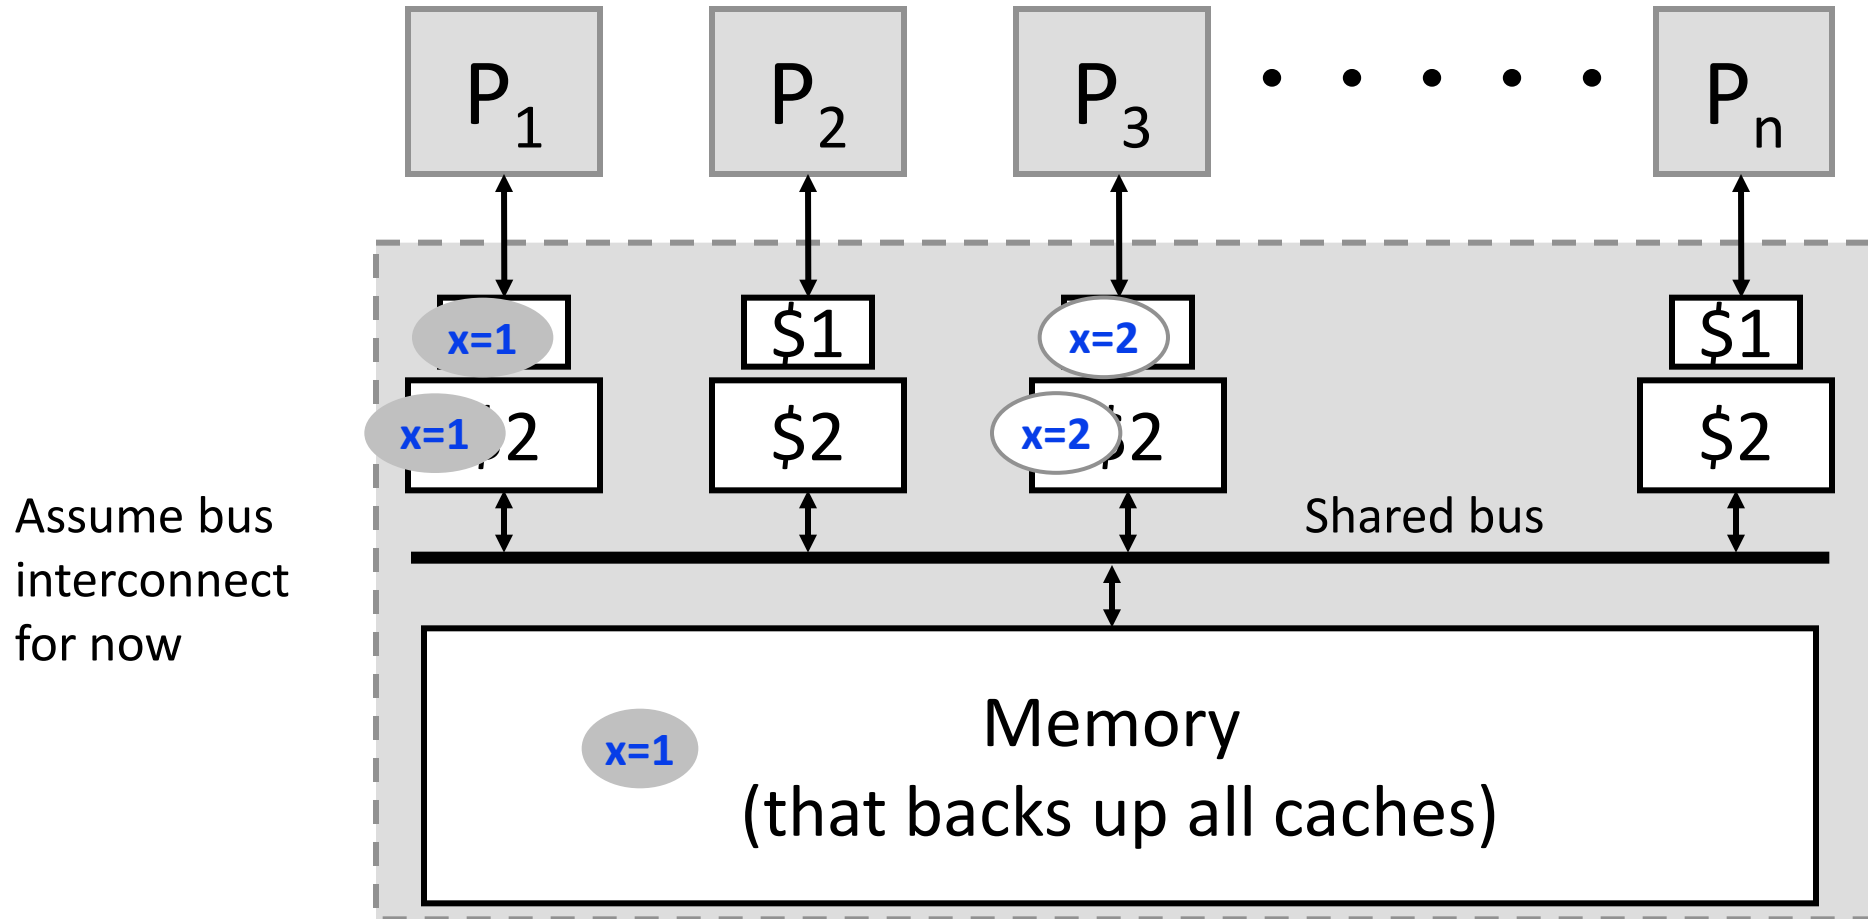
- Bus-based systems
 - Limited scalability: 2–8 processors
 - The bus can be a bandwidth bottleneck
 - Typically, processors split from memories (UMA)
 - Sometimes multiple processors on single chip (CMP)
 - Symmetric multiprocessors (SMPs)
 - Used to be common until ~2006 (e.g., Intel Front-Side Bus)

Impl. #2: Point-to-Point Interconnect



- General point-to-point network-based systems
 - Typically, processor/memory/router blocks (NUMA)
 - **Glueless MP**: no need for additional “glue” (or router) chips
 - Better scalability than bus (can scale to even 1000s of processors)
 - **Massively parallel processors (MPPs)**
 - Used by most dual-socket or quad-socket servers today
 - Intel Ultra-Path Interconnect, AMD Infinity fabric, etc.

MPs with Private Caches: Cache Coherence?



The goal of “cache coherence” is to make all the processors believe they are connected to the same memory directly (warning: not a formal definition)

Extreme Solutions to Cache Coherence

#1. No caching of shared variables

#2. Allow multiple copies, but make sure they all have the same value at all times

- Updates to one copy must be visible to all copies wherever they may be (memory and all of the caches)
- Thus, can have **multiple readers and writers** at once

#3. Allow only one copy of a memory location at a time

- If location **x** is cached in one cache then it is not valid in memory or another cache
- Another processor must have a way to find out who has location **x** and take over ownership before reading or writing
- Thus, can only have **one reader/writer** per location

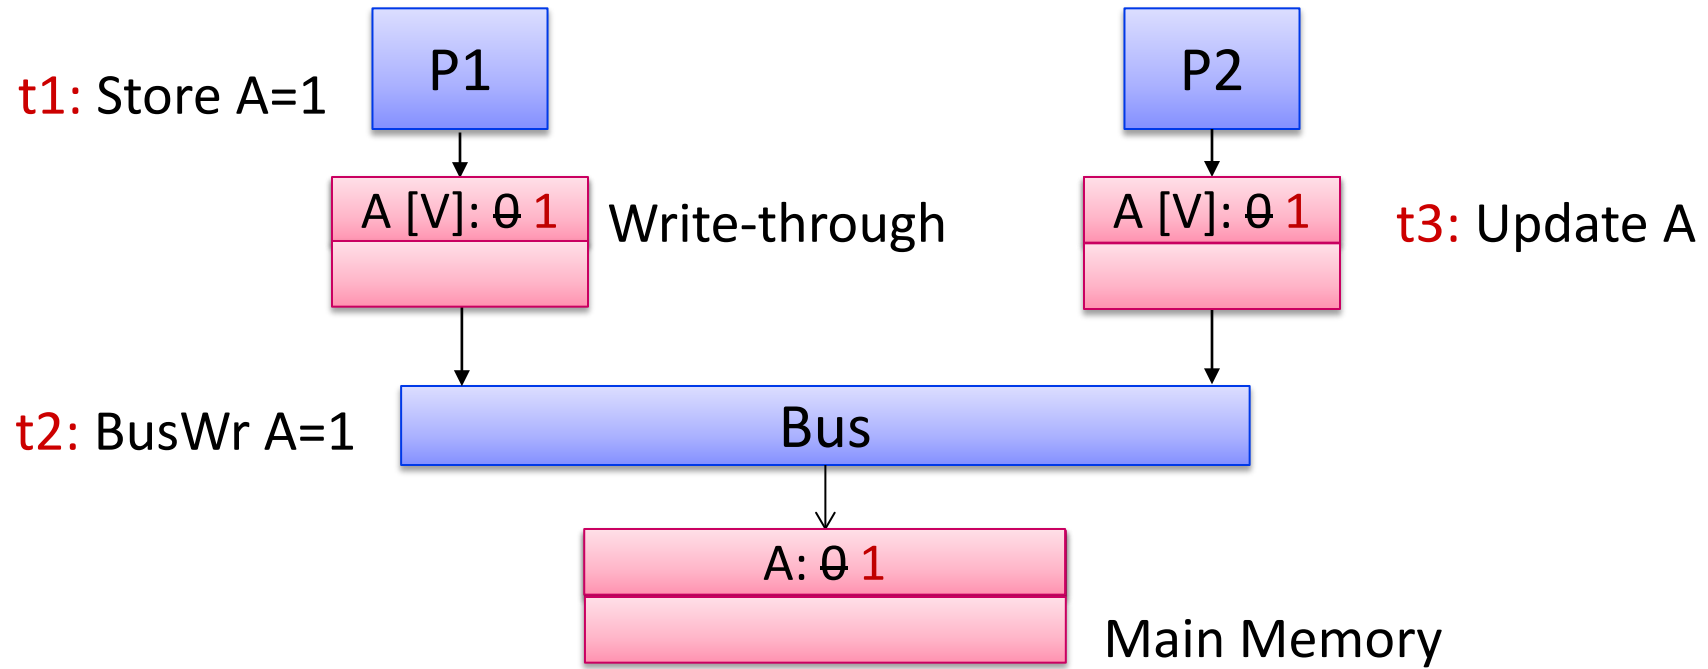
What are the drawbacks of the approaches?

A cache coherence protocol is the “rule” between caches to enforce a particular policy

#2. “Multiple Identical Copies” Protocol

- A cache line can be either **Valid** or **Invalid**
- Based on a write-through (and write-no-allocate) scheme
 - A cache issues a read transaction on a read
 - A cache issues a write transaction to memory whenever its cached line is modified
 - A cache does not need to write back when a line is displaced
- Allows **multiple readers**, but must write through to bus
- All caches “snoop” the bus for other’s write transactions
 - Check if the write is to a currently cached location
 - **Tremendous bus traffic and tag BW** (typically 15% of cache accesses are stores)
 - **If a write goes to a cached location**
 - (1) overwrite the old value with the new snoop value (=write-update protocol)**
 - OR (2) delete the old value (=valid-invalid protocol)**

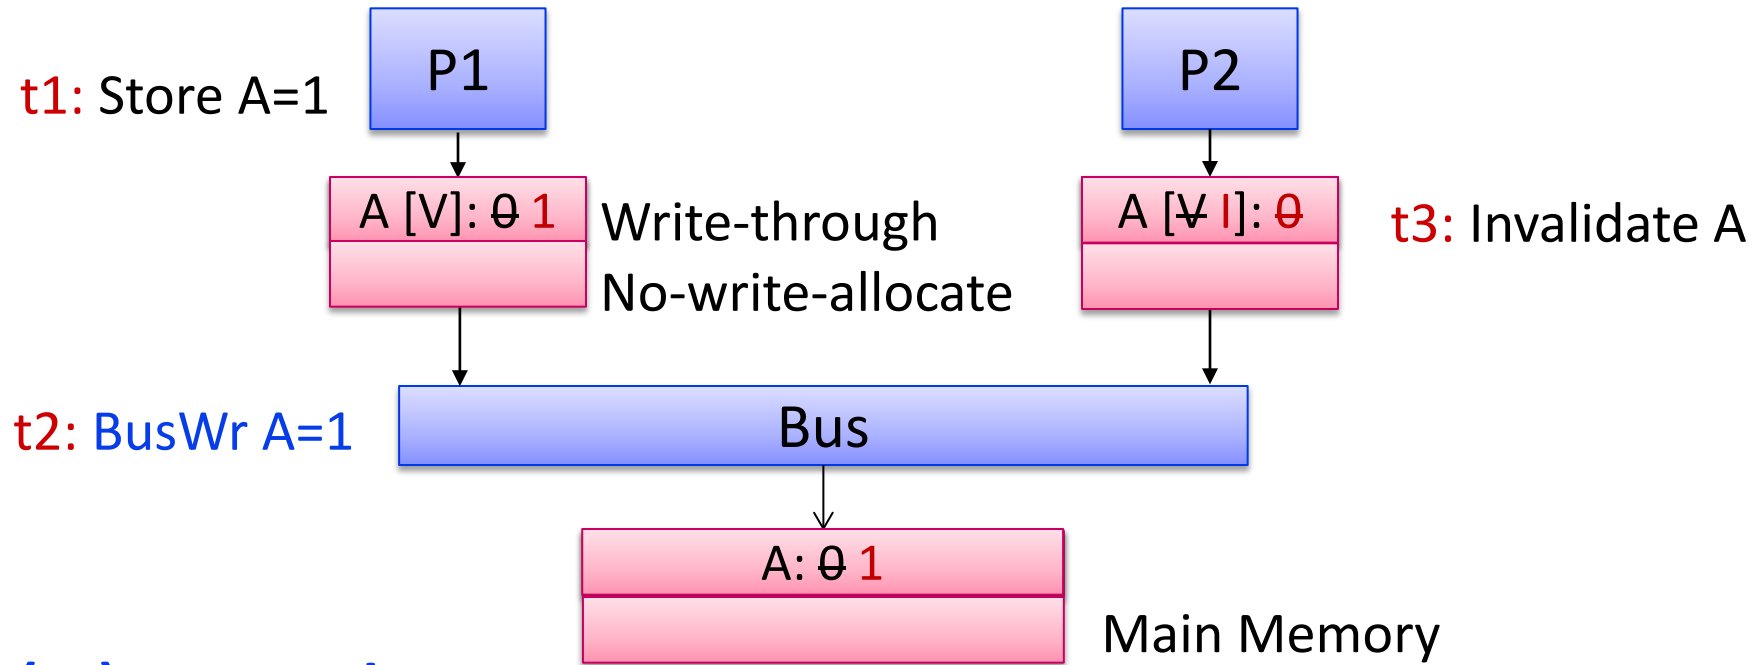
#2-1: Write-Update Protocol (for Write-Through \$)



Write-Update protocol

- “Update” all values of A in caches and memory
→ support multiple copies (for both reads & writes)

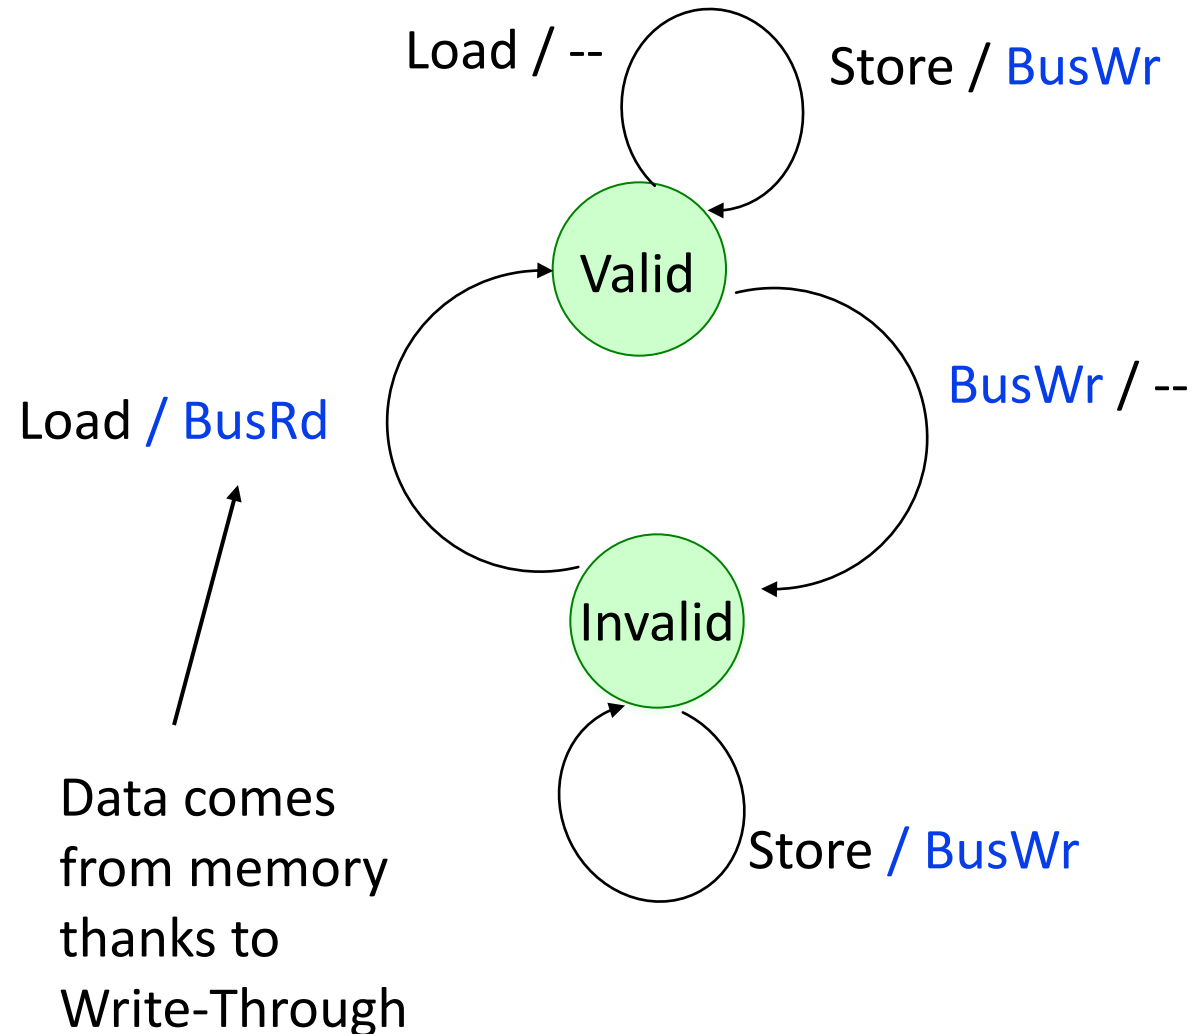
#2-2: Valid-Invalid Protocol (for Write-Through \$)



Valid-Invalid (VI) protocol

- Whenever a cache write is done, invalidate other copies
→ support multiple copies for reads but not writes
- Both write-update and valid-invalid protocol can be implemented with a simple state machine for each cache line

Valid-Invalid Snooping Protocol




- Write-through, no-write-allocate cache
- Actions: Ld, St, BusRd, BusWr
- 1 bit of storage overhead per cache line (valid bit)
- Support multiple readers

#3. “One Copy at All Time” Protocol, Using Write-Back Cache

- Write-through schemes consume too much bandwidth
- How about using write-back caches?
 - Cache issues a read transaction on a read or write miss
 - Cache issues a write-back to memory only when a line is displaced
- A cache line can be either **Modified** or **Invalid**
- All caches “snoop” the bus for other’s *read* transactions
 - If a cache observes a request to a currently cached line, then respond with a value in its cache and mark its own copy **Invalid**
 - Alternatively, a cache can also ask the requestor to retry later and, in the meantime, write back its copy to memory

Why don’t caches need to snoop for *write-back* transactions?

Modified-Invalid Snooping Protocol



- **Write-back**, write-allocate
- **Actions:**
Ld, St, BusRd,
- 1 bit of storage overhead per cache line
- **One reader, one writer**

How about “Multiple Copies” Protocol Using “Write-Back” Cache?

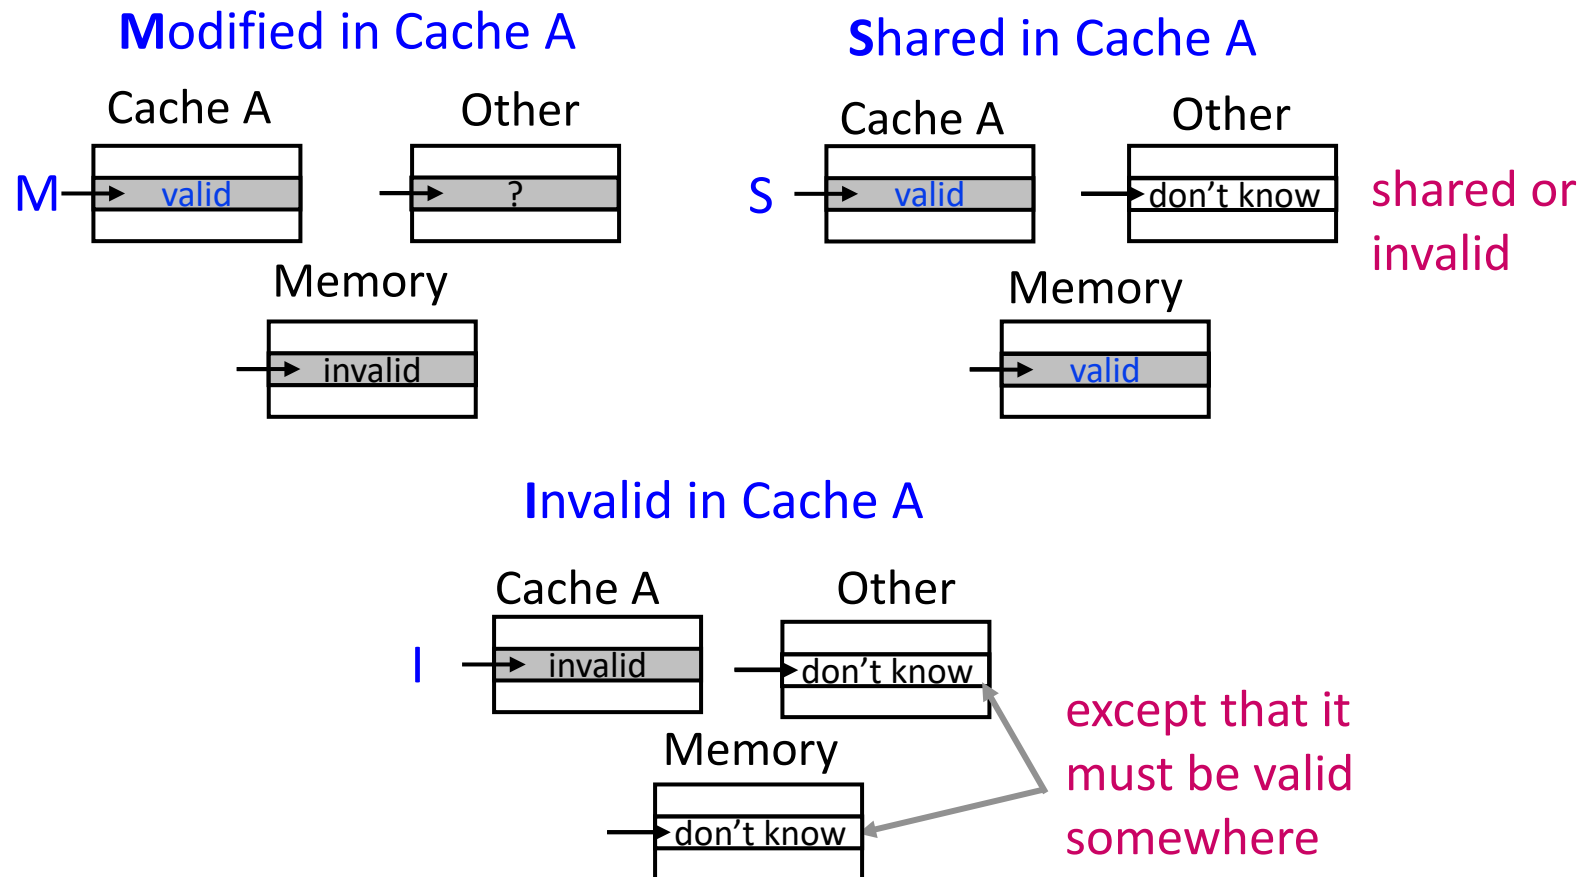
- Key idea: differentiate between “mutually exclusive” and “shared” states
 - **Mutual exclusion** – when “owner” has the only replica of a cache block, it may update it freely
 - **Sharing** – multiple readers are ok, but they cannot write without gaining ownership
- (1) Must find which cache (if any) is an owner on read misses
- (2) Must eventually update memory so that writes are not lost

CC Protocol for Bus-based Systems

- Bus is a broadcast medium, bus “snooping” allows every cache to see what everyone else wants to do
- A cache can even intervene in another cache’s bus transaction (e.g., a cache might ask another cache to “retry” the transaction later or respond in place of the memory)
- Besides the usual status bits, additional information (i.e., **cache coherence states**) is recorded with each cache line
- Example:
 - **Invalid**: cache line does not have valid data
 - **Shared**: valid line, but other caches may have copies
(presumably all identical and unchanged from memory)
 - **Modified**: cache line has been written to since it was brought in

MSI Cache Coherence States

- Given the state of an address in one cache, what are the possible states of the same address elsewhere (if they exist)?



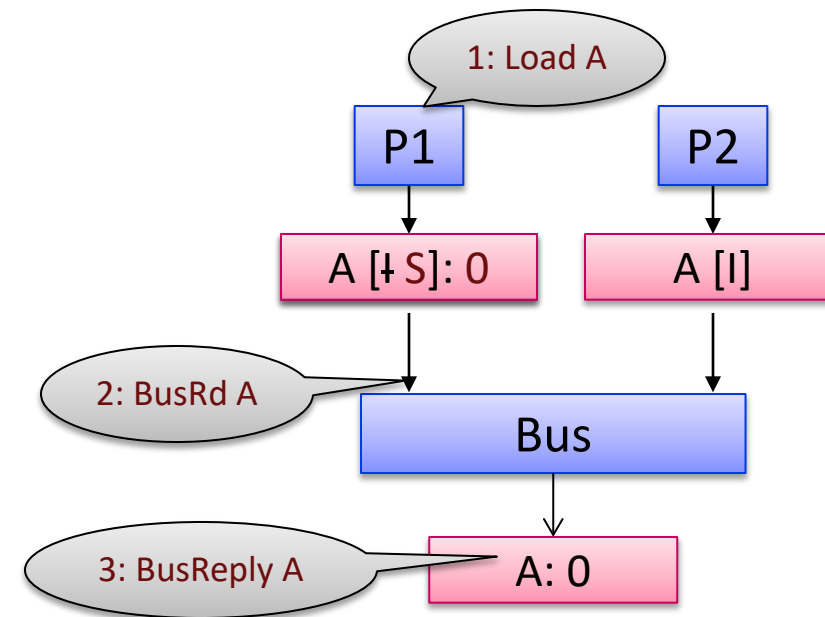

MSI Protocol

- An efficient policy for **single-writer/multi-reader** usage
 - Allow multiple read-only copies (all identical) (**Shared**)
 - Allow only a single writable copy (**Modified**)
 - Reduces the number of bus transactions
- Based on a write-back scheme
 - On a read miss, issue a read transaction for a read-only copy
 - On a write miss, issue a “read-with-intent-to-modify” for an exclusive copy
 - On a write hit to a read-only copy, issue an “invalidate” transaction for other caches
 - When displacing a **Modified** line, write the dirty value back to memory
- All caches “snoop” the bus for other caches’ read, “read-with-intent-to-modify” and invalidate transactions

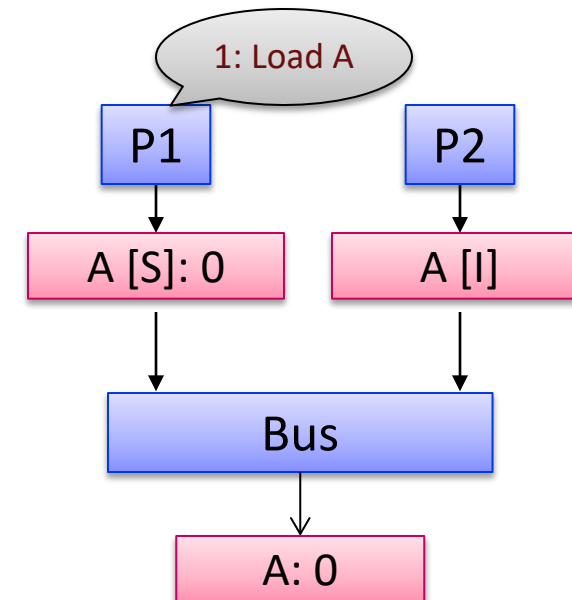

Modified-Shared-Invalid (MSI) Protocol

- Three states tracked per-block at each cache
 - **Invalid** – cache does not have a copy
 - **Shared** – cache has a read-only copy; clean (i.e., memory is up-to-date)
 - Some other caches might have this copy as well
 - **Modified** – cache has a writable, “dirty” copy
 - Dirty = memory is out of date
 - No other cache has this copy
- Three actions **from processor**
 - **Load, Store, Evict**
- Five messages **from bus**
 - **BusRd, BusRdX, BusInv, BusWB, BusReply**

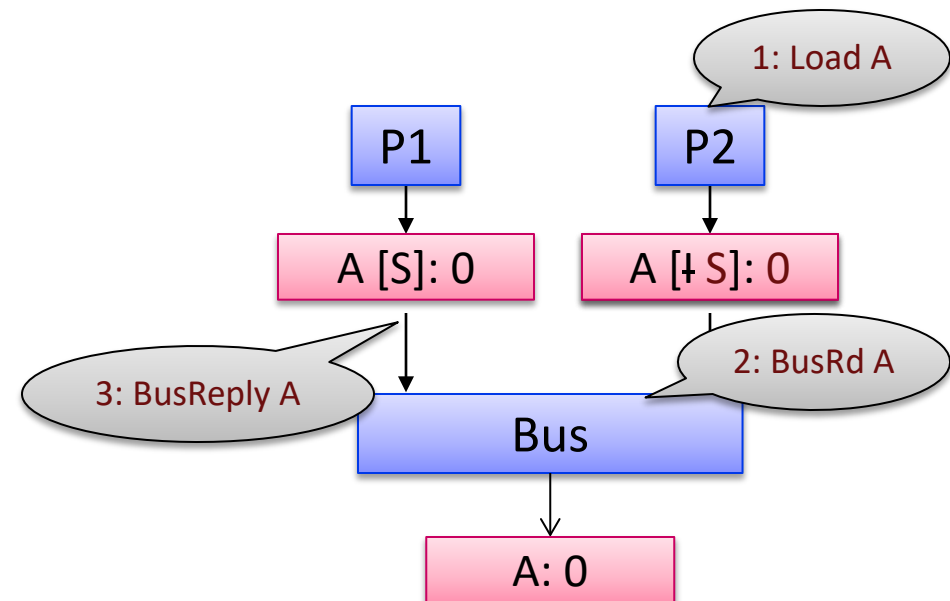

Modified-Shared-Invalid (MSI) Protocol



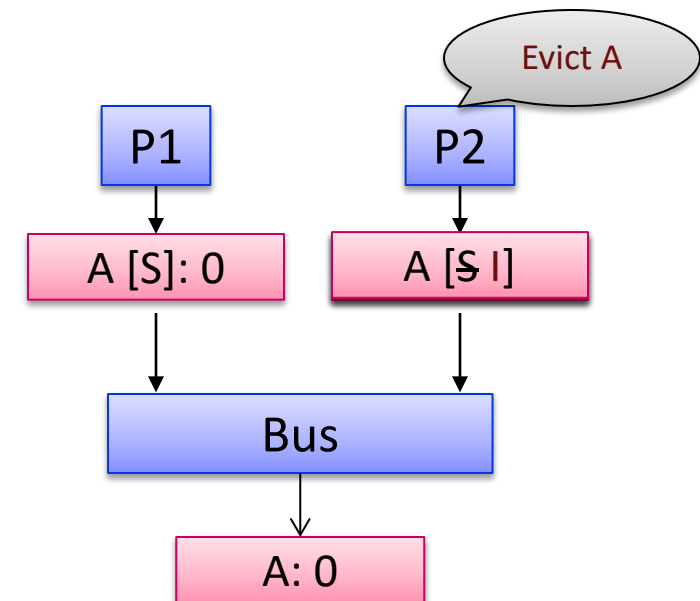

Modified-Shared-Invalid (MSI) Protocol



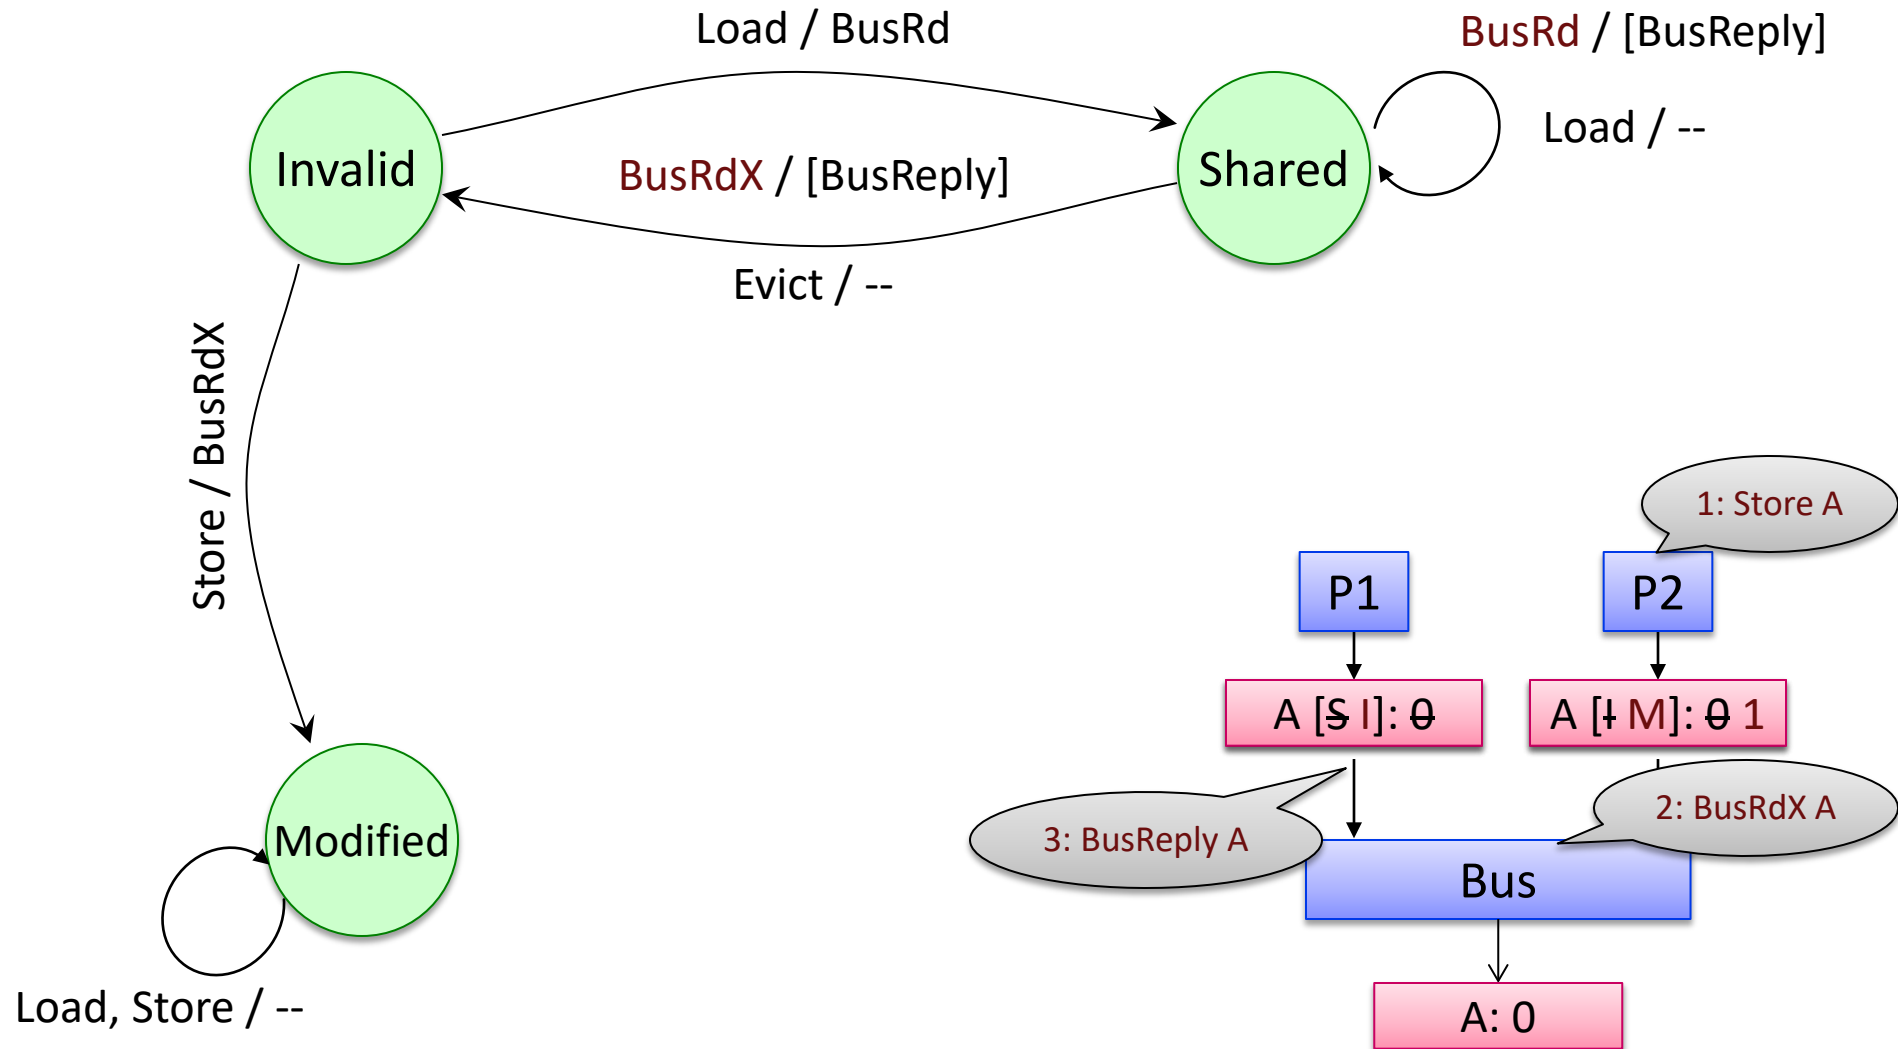
Modified-Shared-Invalid (MSI) Protocol



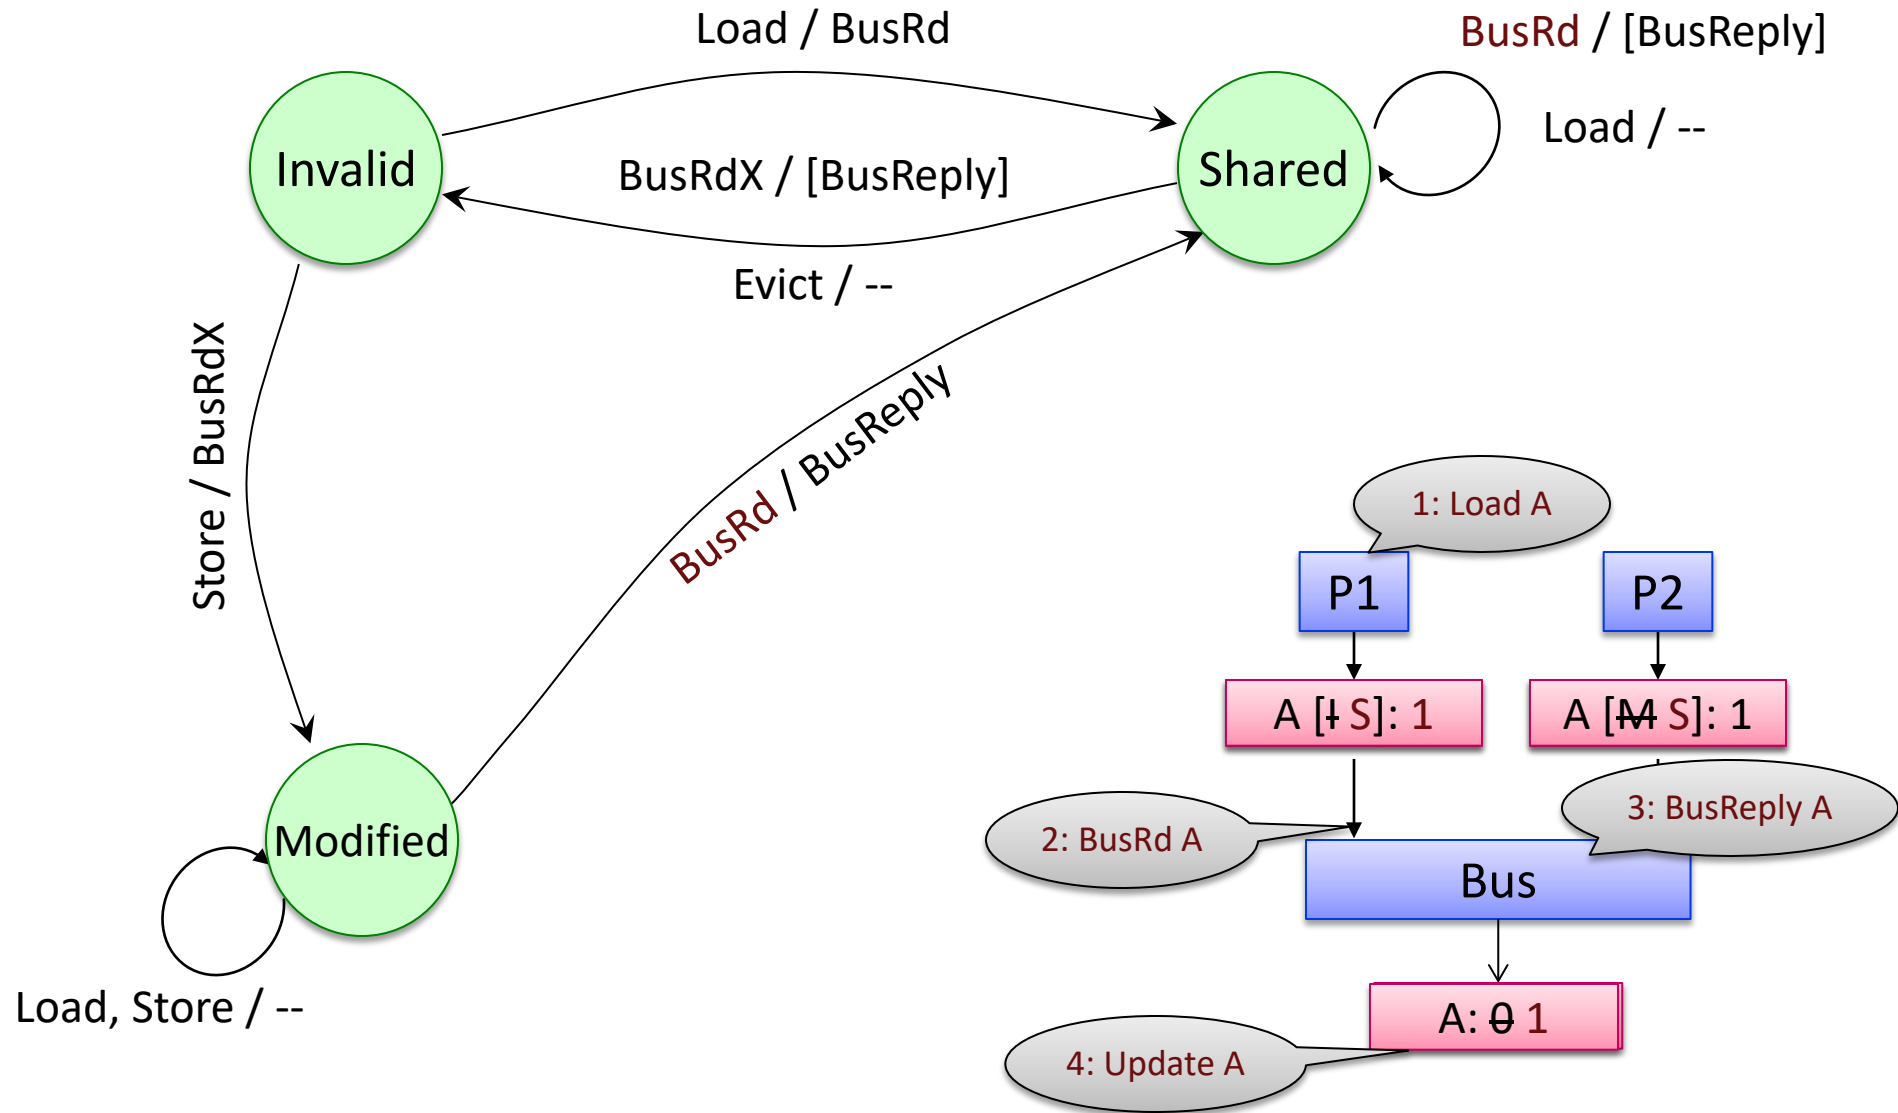
Modified-Shared-Invalid (MSI) Protocol



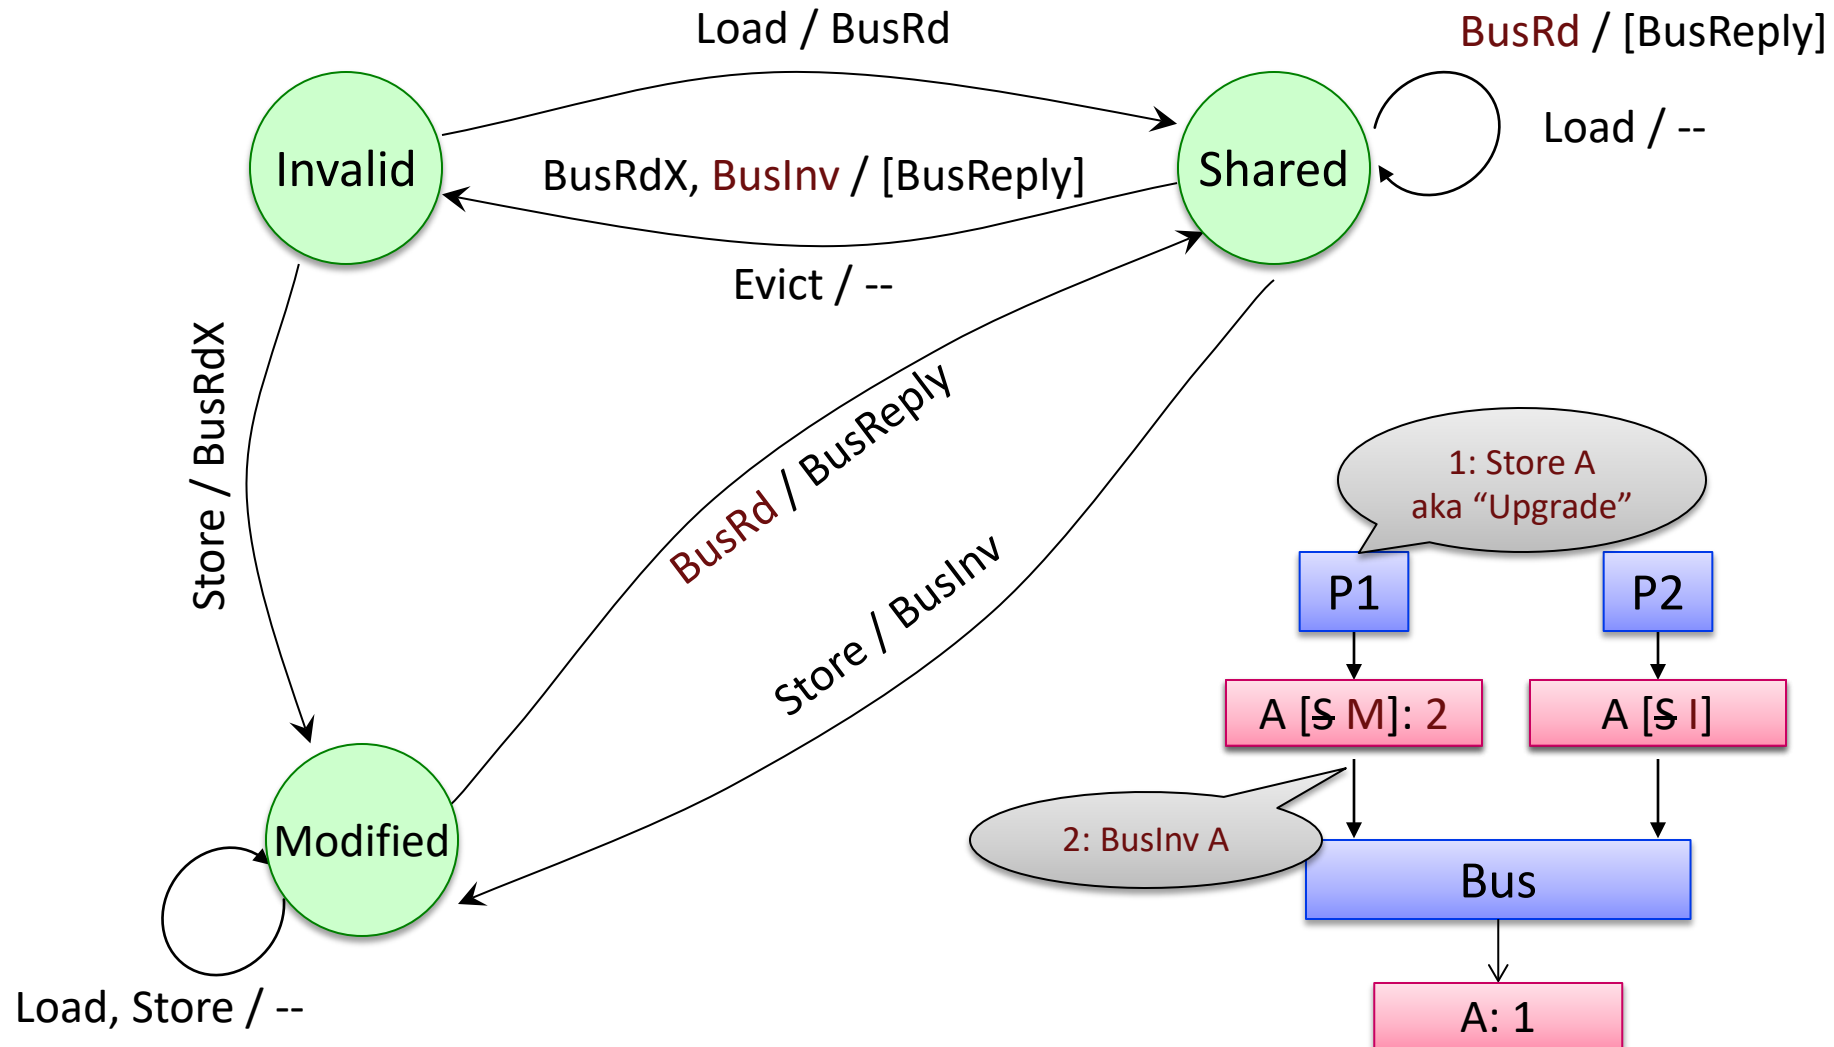
Modified-Shared-Invalid (MSI) Protocol



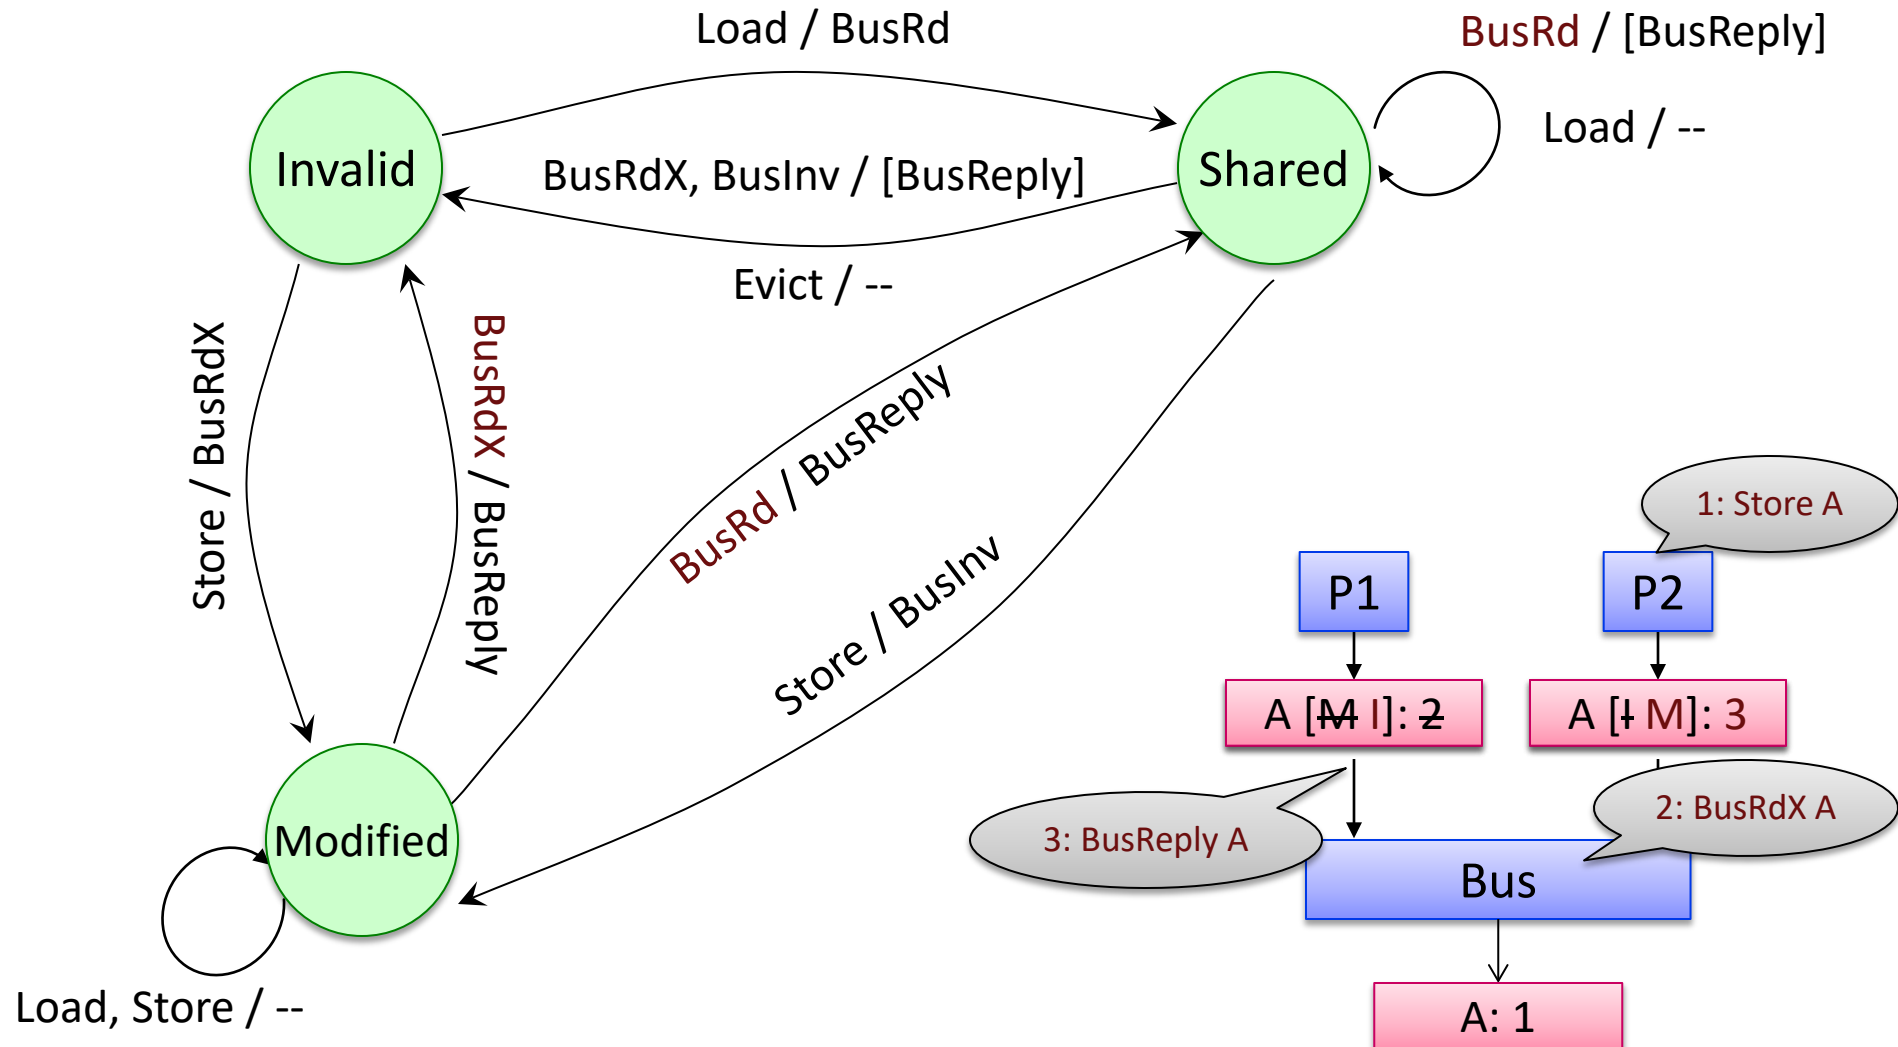
Modified-Shared-Invalid (MSI) Protocol



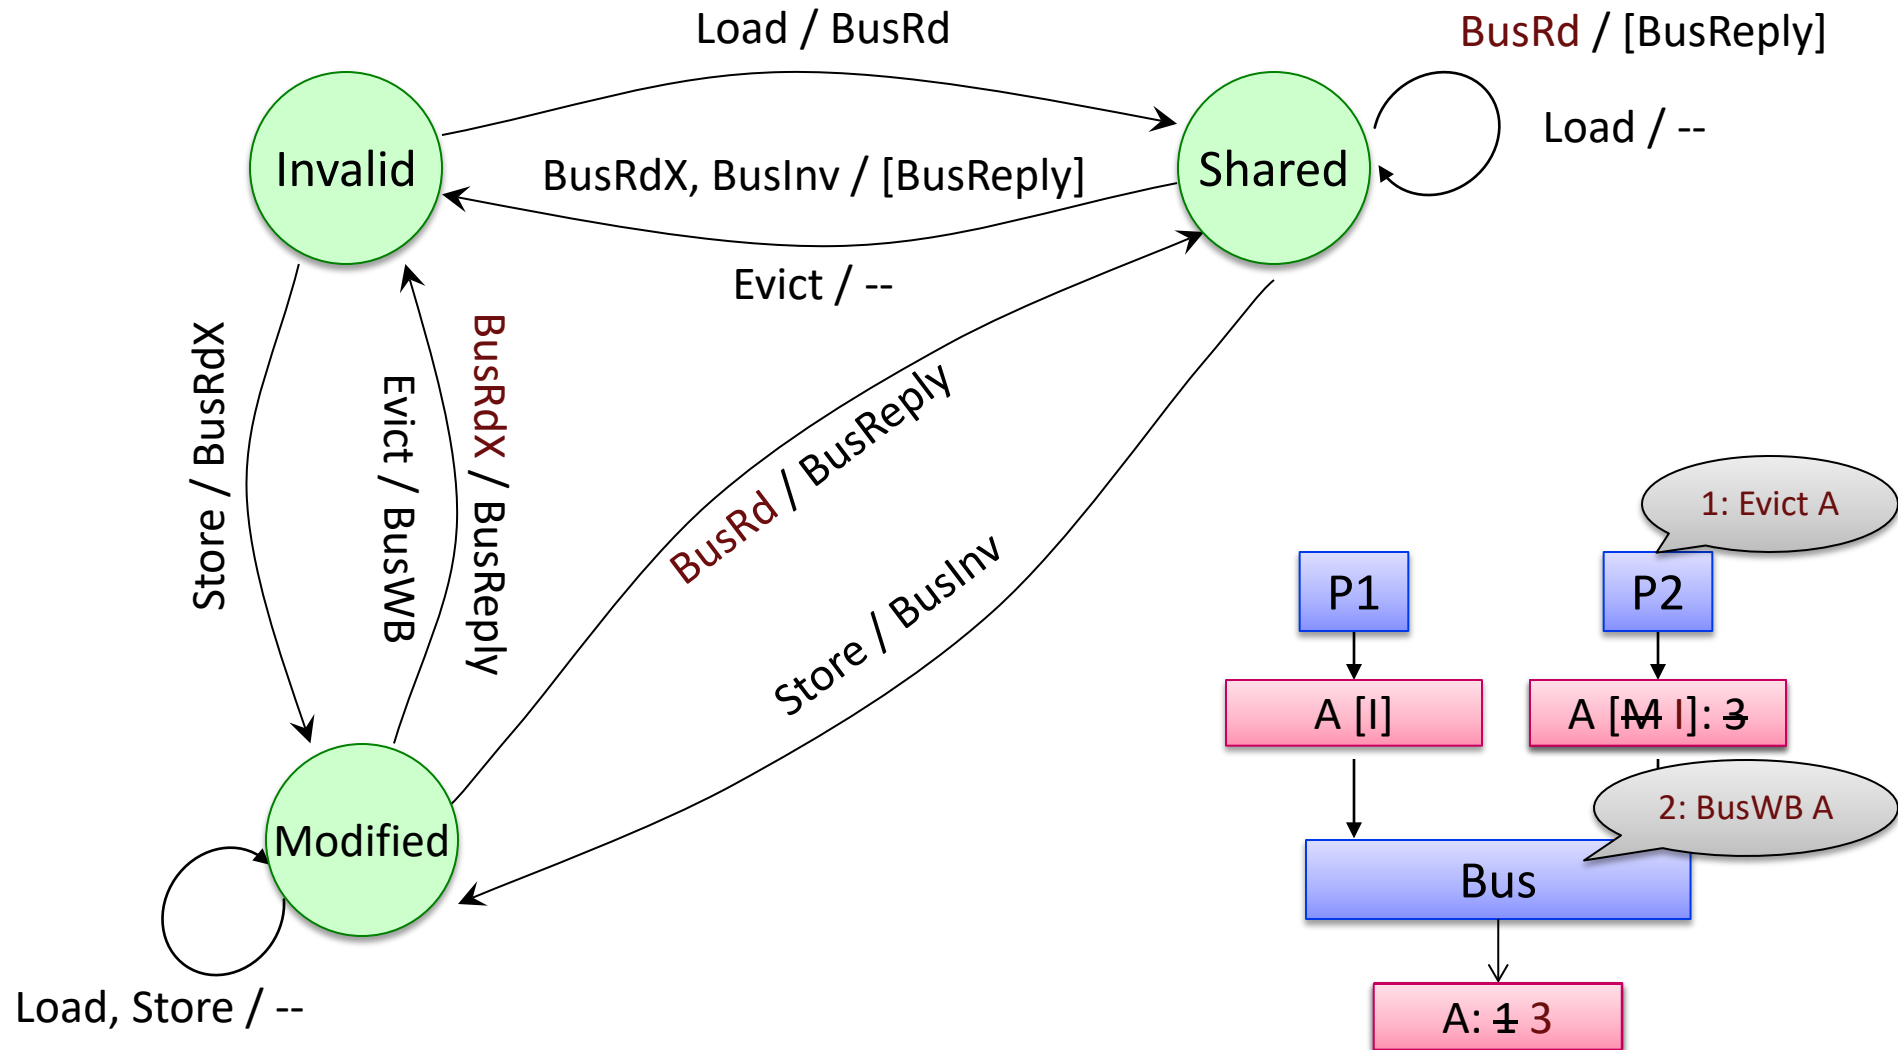
Modified-Shared-Invalid (MSI) Protocol



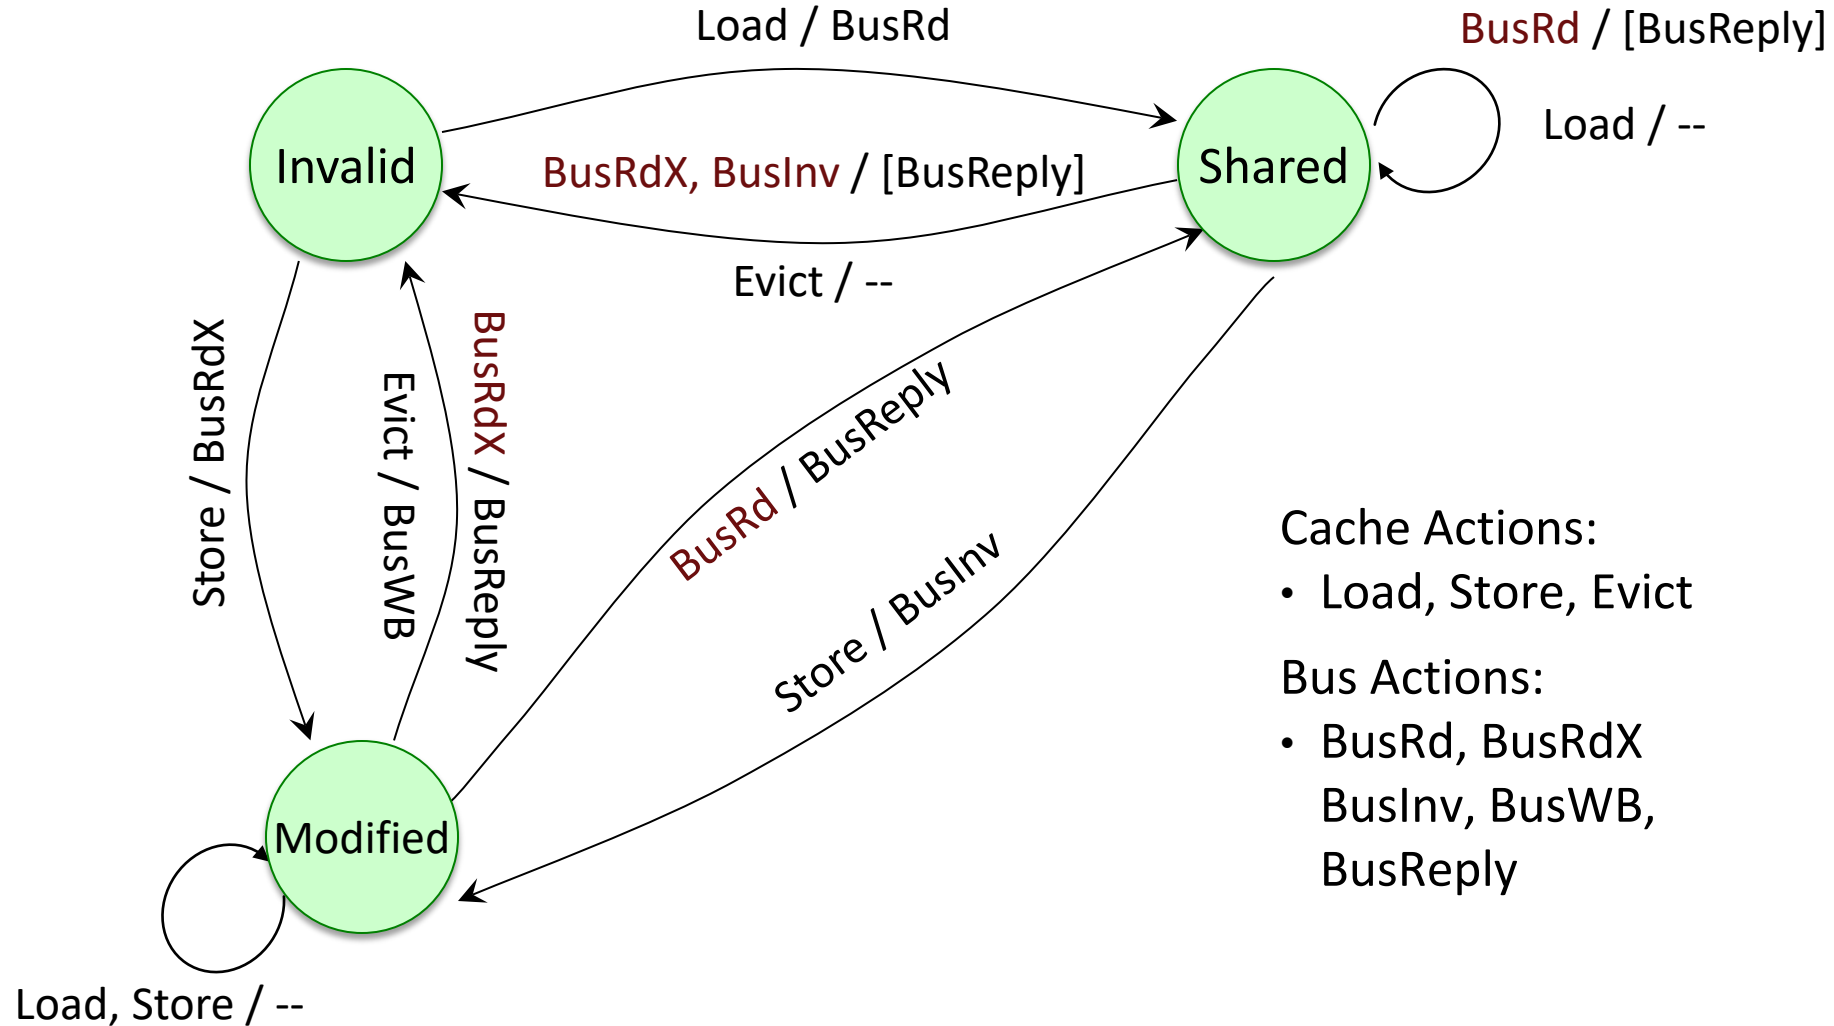
Modified-Shared-Invalid (MSI) Protocol



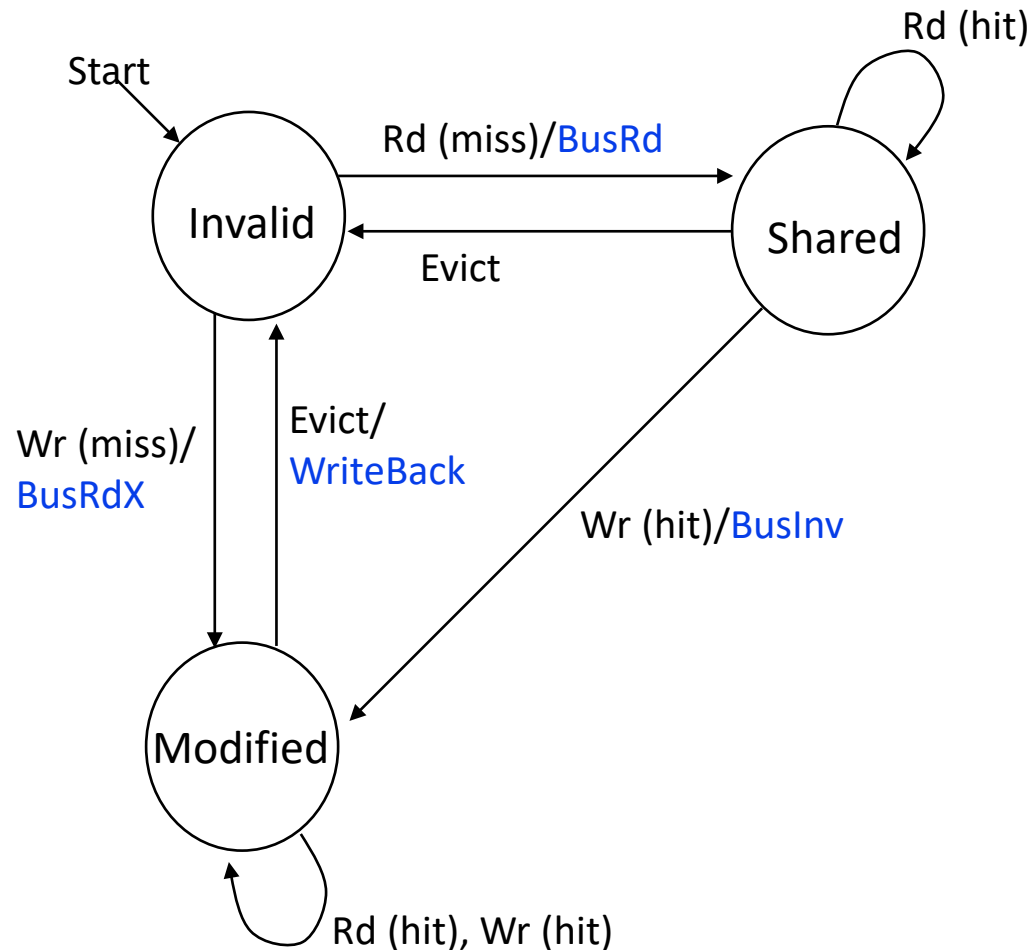
Modified-Shared-Invalid (MSI) Protocol



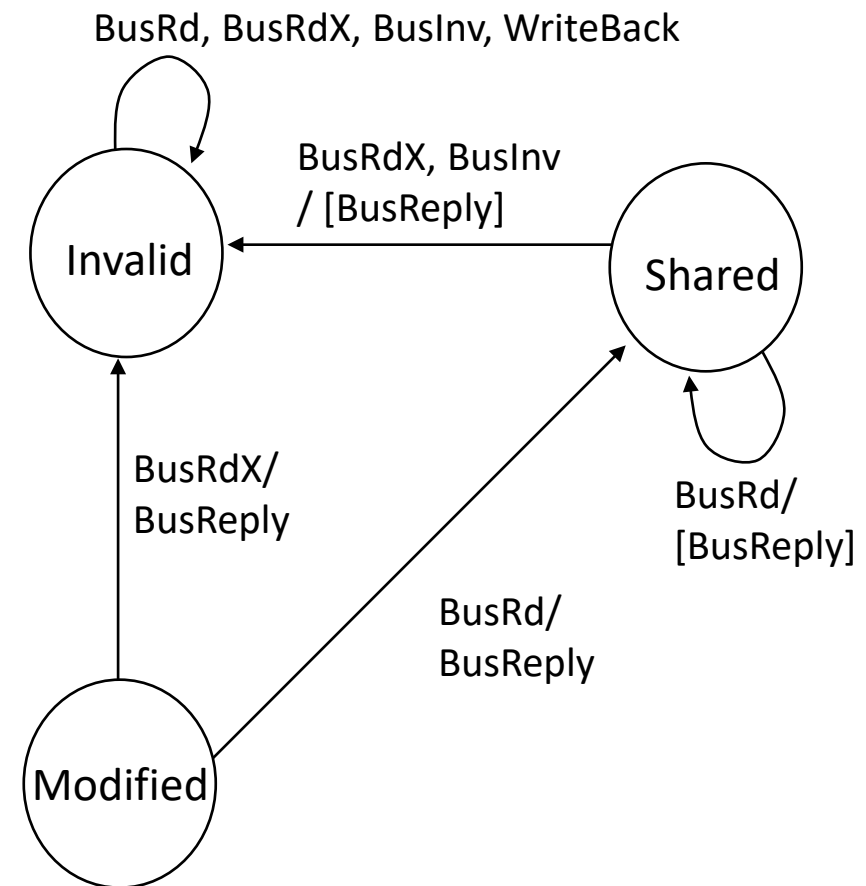
MSI Protocol Summary



MSI State Transition Diagram



a. State transitions using signals from the CPU

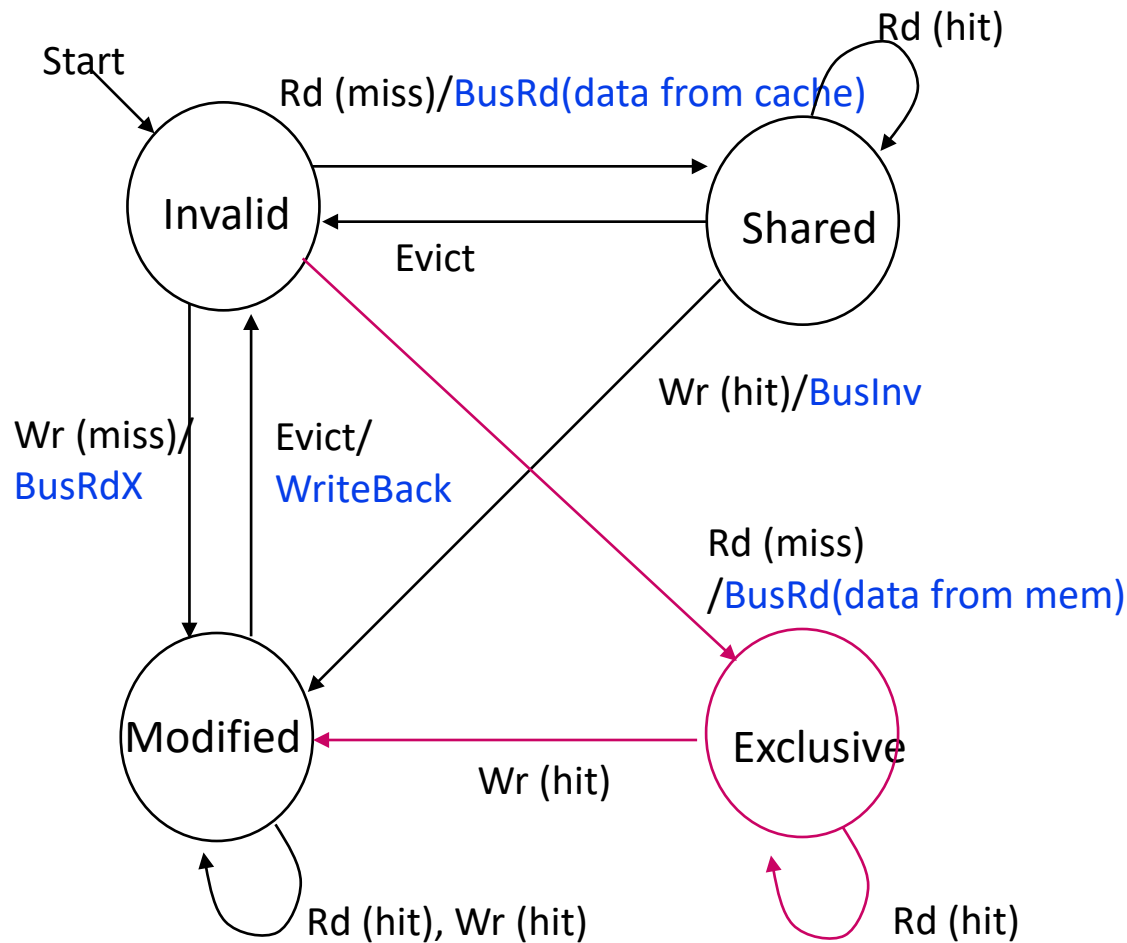


b. State transitions using signals from the BUS

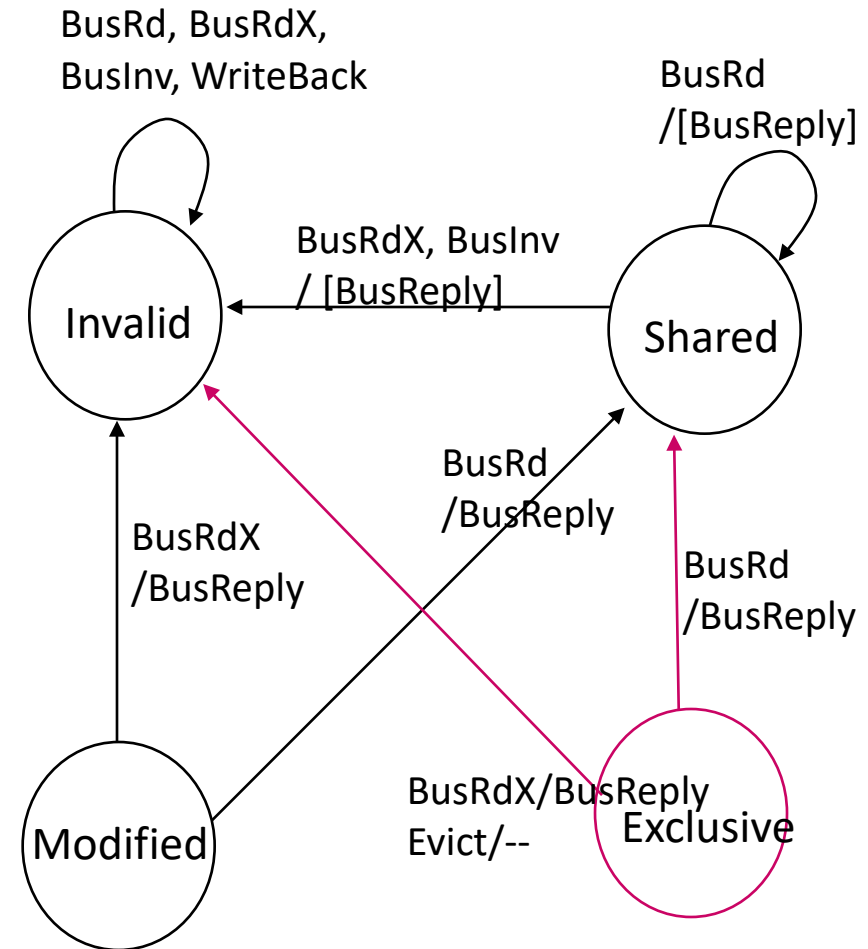
MESI Protocol

- MSI suffers from frequent read-upgrade sequences
 - Leads to two bus transactions, even for private blocks (=one sharer)
 - BusRd: gets the block in S
 - BusInv (or BusRdX, or BusUpgr): convert S to M
 - Must broadcast useless invalidation messages
 - Uniprocessors don't have this problem
- Solution: add an “Exclusive” state
 - Exclusive – only one copy; writable; **clean**
 - Can detect exclusivity when memory provides reply to a read
 - No need for a BusWB from Exclusive
 - For stores, silently transition to “Modified” to indicate data is dirty

Modified-Exclusive-Shared-Invalid (MESI) Protocol



a. State transitions using signals from the CPU



b. State transitions using signals from the BUS

Implementation of Snoopy Bus

- **Every bus snoop results in a cache lookup**

- Needs a dual ported cache or at least an extra tag lookup port
- In an inclusive L1 and L2 arrangement, snoop (i.e., messages from the bus) only goes to L2 and does not contend with processor for L1 bandwidth
(True benefits of “inclusive” cache!)

- **Bus is not very scalable**

- Physical limits (such as number of drops and physical extent of the bus) force bigger busses to clock slower
- Bus bandwidth is divided when you add more processors

- **Implementing snoopy (or any coherence) protocols can get complicated**

- MSI and MESI state transitions are not really atomic (need many *transient* states)
- CPU and bus transactions are not atomic
- CC issues can become intertwined with memory consistency

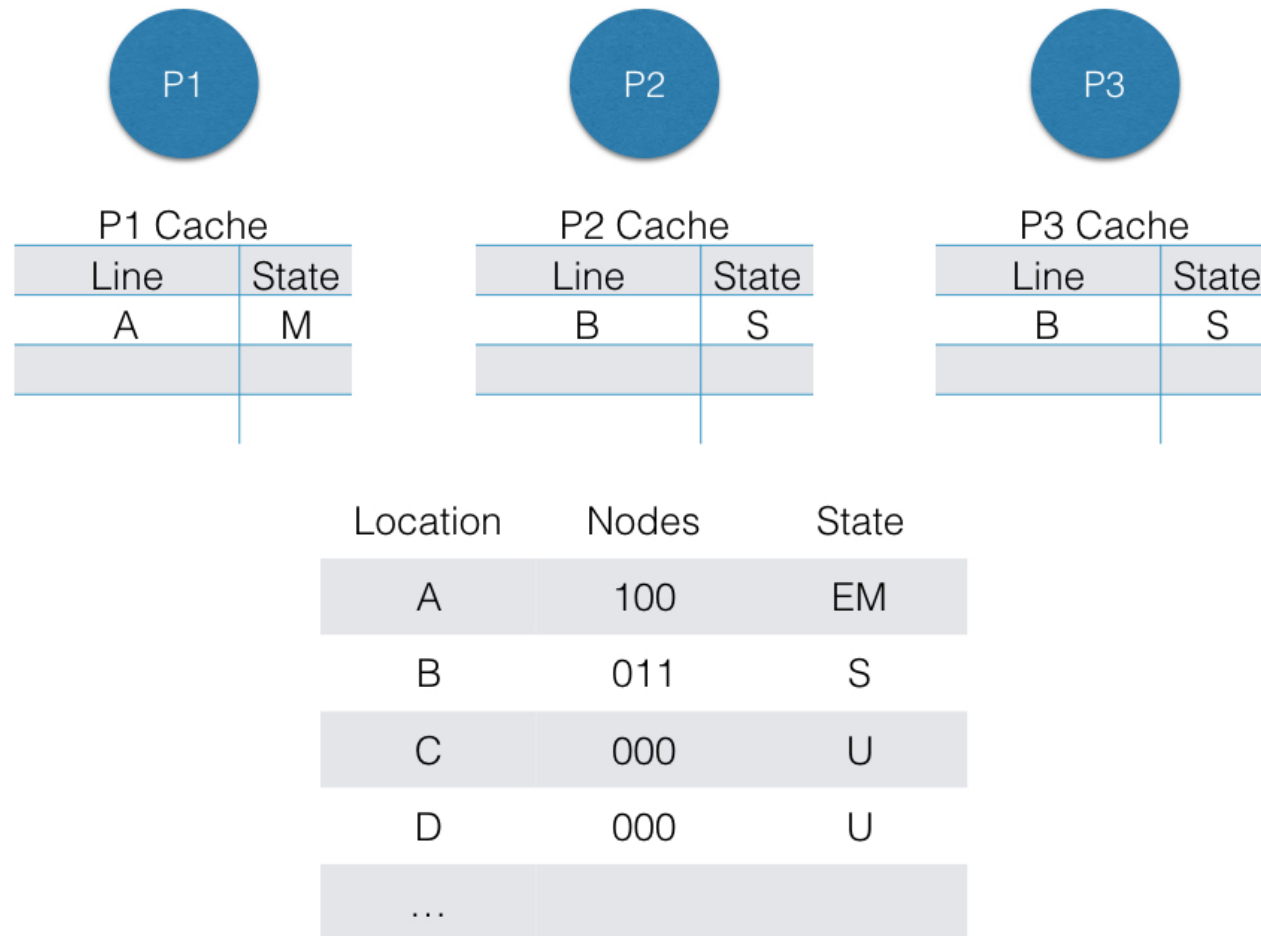
Scalable Cache Coherence

- **Scalable cache coherence**: two part solution
- Part I: **Bus bandwidth**
 - Replace non-scalable bandwidth substrate (bus) with scalable bandwidth one (point-to-point network, e.g., mesh)
- Part II: **Processor snooping bandwidth**
 - Observation: most snoops result in no action
 - Replace non-scalable broadcast protocol (spam everyone) with scalable **Directory protocol** (only spam processors that care)

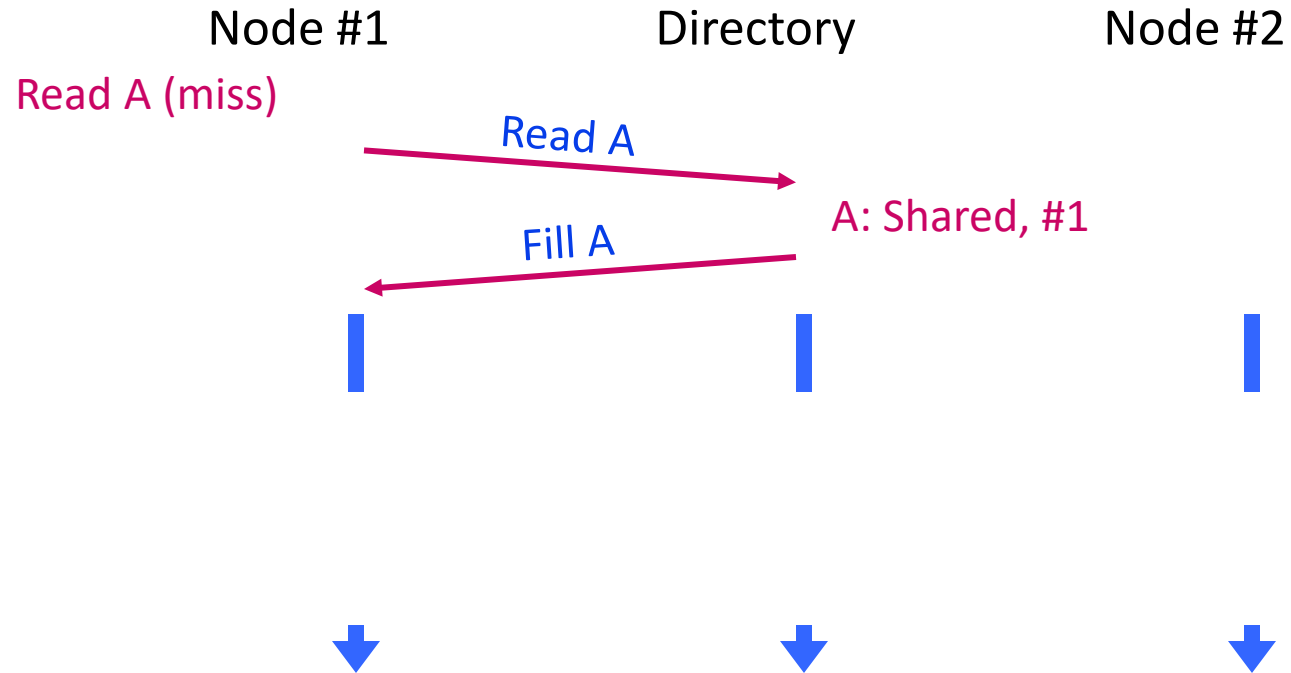
Directory Coherence Protocols

- Idea: A logically-central directory keeps track of where the copies of each cache block reside. Caches consult this directory to ensure coherence.
- **Directories** for non-broadcast coherence protocol
 - Extend memory to track caching information
 - For each physical cache line, the directory records:
 - **Owner**: which processor has a dirty (M state) or exclusive (E state) copy
 - **Sharers**: which processors have clean (S state) copies – 1 bit for each processor
 - Processor sends coherence event to the directory. For example, (for the corresponding directory entry)
 - On a read: set the cache's sharer bit and arrange the supply of data
 - On a write: invalidate all caches that have the block and reset their bits
- Advantage: scalable (less coherence traffic)
- Disadvantages: additional storage overhead, increased latency (\$ → dir. → mem.)

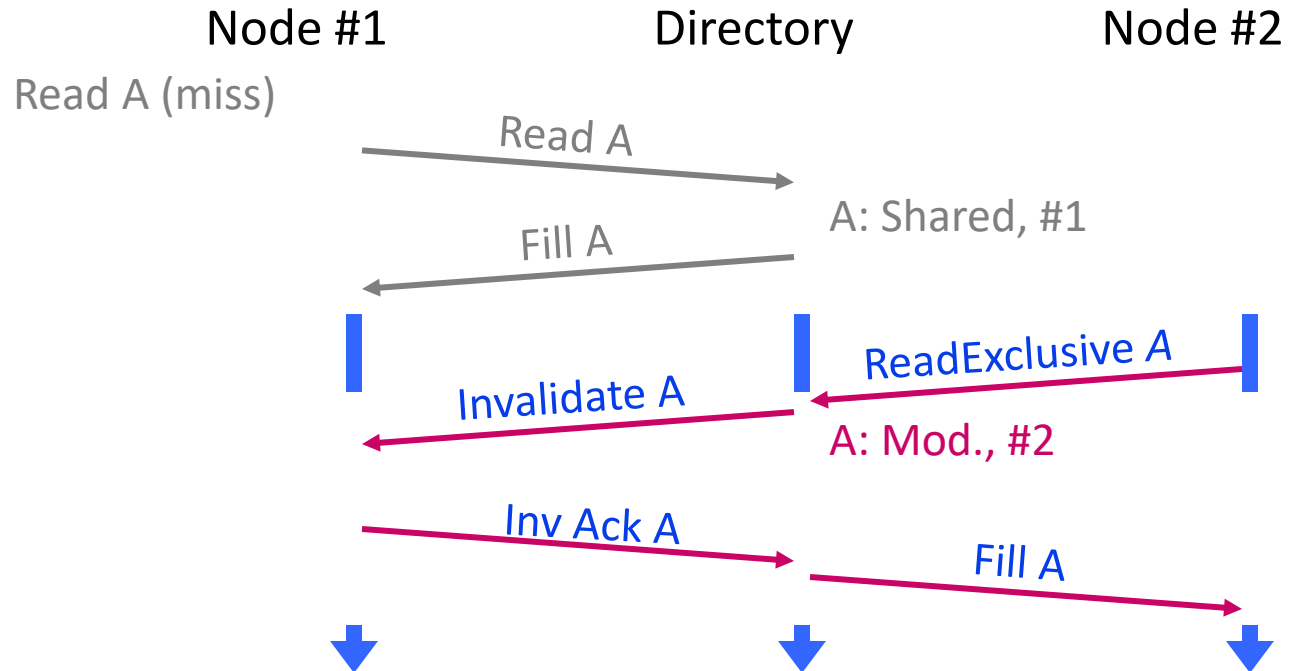
Directory-based Cache Coherence Example



Basic Operation: Read



Basic Operation: Write




Centralized Directory

- **Single directory** contains a copy of cache tags from all nodes
- **Advantages:**
 - Send Invalidate/Update only to nodes that have copies
 - Central serialization point: easy to get memory consistency (just like a bus...)
- **Problems:**
 - Less scalable than a distributed directory (imagine traffic from 1000's of nodes...)
 - Directory size/organization changes with the number of nodes

Distributed Directory

- **Distribute directory** among memory modules
 - Memory block = Coherence block (usually = cache line)
 - “**Home node**” → Node with directory entry
 - Usually also dedicated main memory storage for cache line
 - **Scalable** – Directory grows with memory capacity
 - But the directory can **no longer serialize** accesses across all addresses
 - **Memory consistency** becomes responsibility of CPU interface



Directory for 1000 nodes?

■ Scalability Problem

- What if I have 2, 4, 8, 16, 32, ..., 1000s CPUs?

(1000s-CPU systems have been implemented)

Snoop-based CC for 100~1000 CPUs (or cores)?

■ Directory-based CC for many nodes

(1) Bit Vector-based Directory CC: N-bit directory info (=bit vector) added to cache state bits for N-CPU system.

(e.g., [10010110] → CPU 1, 4, 6, 7 are sharing the block)

(2) For too many nodes, Pointer-based Directory CC can be used.

(e.g., [1001 0110] → CPU 9, CPU 6 are sharing the block)


Which one is better?

Designing a Directory Protocol: Nomenclature

- Local Node (L)
 - Node initiating the transaction we care about
- Home Node (H)
 - The node where the memory location and the directory entry of an address reside
- Remote Node (R)
 - Any other node that participates in the transaction

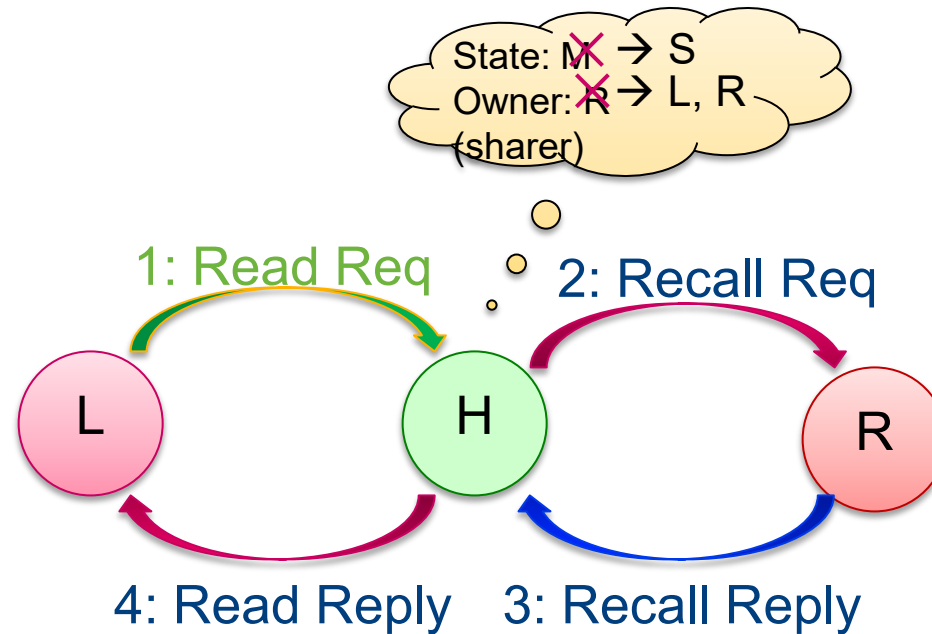
Read Transaction (data in home memory)

- L has a cache miss on a load instruction



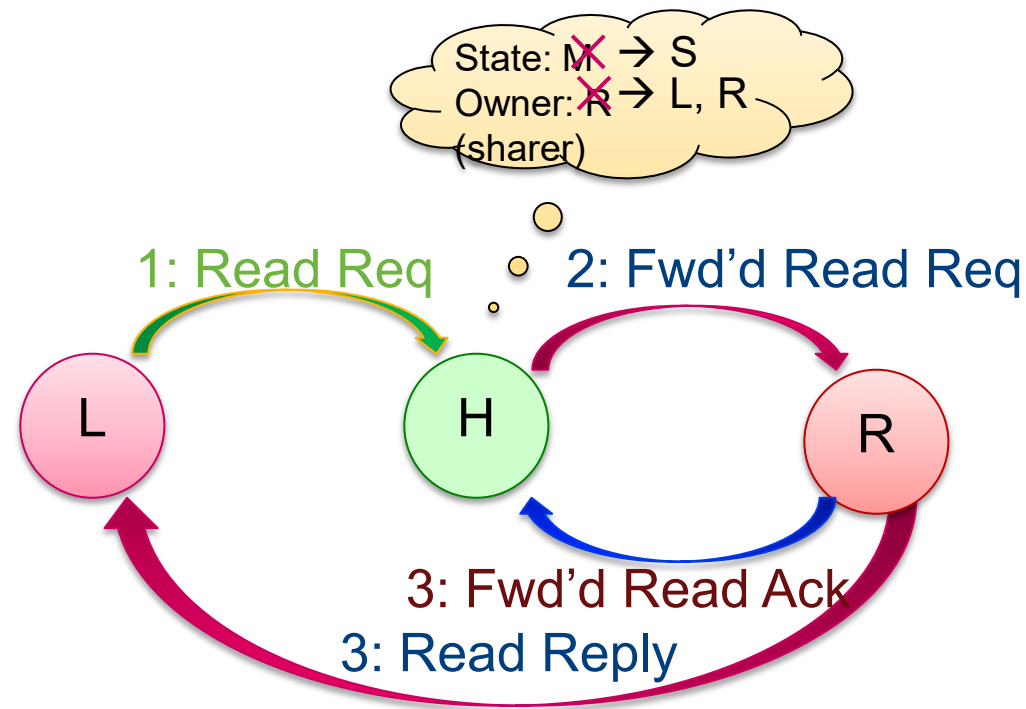
4-hop Read Transaction (data in a remote cache)

- L has a cache miss on a load instruction
 - Block was previously in modified state at R



3-hop Read Transaction (data in a remote cache)


- L has a cache miss on a load instruction
 - Block was previously in modified state at R



Performance Issues (General Idea)


■ Deadlock

- No one is making a progress
e.g., Everyone stopped to wait for a forever-busy resource




■ Livelock

- Some are making progress, but to reach the old same status
e.g., Everyone tries to get a resource at every odd second, while the resource becomes available at every even second



■ Starvation

- Only some are making progress (= unfair progress)
e.g., only one guy keeps winning in getting the resource



Performance Issues in CC

■ Deadlock example

Different addresses A & B happen to map to ~~a same~~ different cache lines.

P_A with a modified copy of A placed a read-miss req for B on the bus.

P_B with a modified copy of B placed a read-miss req for A on the bus.

What if P_A and P_B want to serve their older requests first?

■ Livelock example

If all CPUs try to write to A simultaneously, they must invalidate other copies of A in CPUs.

What if As in all caches continue to get invalidated before any write is completed?

■ Starvation example

What if a specific CPU continues to become successful in writing to A, while other CPUs are failing to write to A?

Extremely difficult to design a high-performance, scalable, and functionally correct cache-coherence protocol.

Shared Memory Summary (1/2)

- Shared-memory multiprocessors
 - + Simple software: easy data sharing, easy programming
 - Complex hardware: must provide illusion of global address space
- Two basic implementations
 - **Symmetric (UMA) multi-processors (SMPs)**
 - Underlying communication network: bus (ordered)
 - + Low-latency, simple protocols that rely on global order
 - Low-bandwidth, poor scalability
 - **Scalable (NUMA) multi-processors (MPPs)**
 - Underlying communication network: point-to-point (unordered)
 - + Scalable bandwidth
 - Higher-latency, complex protocols

Shared Memory Summary (2/2)

- Two aspects to global memory space illusion
 - **Coherence**: consistent view of individual cache lines
 - Absolute coherence not needed, relative coherence OK
 - VI and MSI protocols, cache-to-cache transfer optimization
 - Implementation?
 - SMP: snooping, MPP: directories
 - **Consistency** (next lecture): consistent global view of all memory locations
 - Programmers intuitively expect sequential consistency (SC)
 - Global interleaving of individual processor access streams
 - Not naturally provided by coherence, needs extra stuff
 - Relaxed ordering: consistency only for synchronization points

Question?

Announcements

- Textbook reading: P&H Ch 5.10
- Paper reading: “The Directory-Based Cache Coherence Protocol for the DASH Multiprocessor”,
Lenoski, Laudon, and Gharachorloo, ISCA’90